

CS 6410: Compilers

Fall 2023

Lecture 2

Tamara Bonaci

t.bonaci@northeastern.edu



Credits For Course Material

- **Big thank you to UW CSE faculty member, Hal Perkins**
- Some direct ancestors of this course:
 - UW CSE 401 (Chambers, Snyder, Notkin, Perkins, Ringenburg, Henry, ...)
 - UW CSE PMP 582/501 (Perkins)
 - Cornell CS 412-3 (Teitelbaum, Perkins)
 - Rice CS 412 (Cooper, Kennedy, Torczon)
 - Many books (Appel; Cooper/Torczon; Aho, [[Lam,] Sethi,] Ullman [Dragon Book], Fischer, [Cytron ,] LeBlanc; Muchnick, ...)

Agenda


- Review – introduction to compilers
 - Front end
 - Back end
- Quick review of basic concepts of formal grammars
- Regular expressions
- Lexical specification of programming languages
- Using finite automata to recognize regular expressions
- Scanners and Tokens

CS 6410, Compilers – Fall 2023


Introduction To Compilers - Review

A Structure of a Compiler

- At a high level, a compiler has two pieces:
 - **Front end – analysis**
 - Read source program, and discover its structure and meaning
 - **Back end – synthesis**
 - Generate equivalent target language program




Compiler: Front End



- **Front end is usually split into two parts:**
 1. **Scanner** – responsible for converting character stream to token stream: keywords, operators, variables, constants
 - Also: strips out white space, comments
 2. **Parser** - reads token stream; generates IR
 - Either here or shortly after, perform semantics analysis to check for things like type errors

Compiler: Front End



- Front end is usually split into two parts:
 - **Scanner** – responsible for converting character stream to token stream
 - **Parser** – reads token stream; generates IR
- Both of these can be generated automatically
 - Use a formal grammar to specify the source language
 - Tools read the grammar and generate scanner & parser (lex/yacc or flex/bison for C/C++, JFlex/CUP for Java)

Parser Output - IR

- Given a token stream from a scanner, the parser must produce output that captures the meaning of the program
- Most common output from a parser is an **Abstract Syntax Tree (AST)**
 - Represents the essential meaning of program without syntactic noise
 - Nodes are operations, children are operands
- **Many different forms**
 - Engineering tradeoffs have changed over time
 - Tradeoffs (and IRs) can also vary between different phases of a single compiler

Static Semantic Analysis (SSA)

- During or after parsing, **check that the program is legal and collect info for the back end**
 - Type checking
 - Check language requirements like proper declarations
 - Preliminary resource allocation
- Collect other information needed by back end analysis and code generation
- **Key data structure: Symbol Table(s)**
 - Maps names -> meaning/types/details

Back End Structure

- Typically split into two major parts
 - **“Optimization” – code improvement**
 - Examples: common sub-expression elimination, constant folding, code motion (move invariant computations outside of loops)
 - Optimization phases often interleaved with analysis
 - **Target Code Generation (machine specific)**
 - Instruction selection & scheduling, register allocation
- Usually walks the AST to generate lower-level intermediate code before optimization

Compilers and Interpreters

- Programs can be compiled or interpreted (or sometimes both)
- **Compiler**
 - A program that translates a program from one language (the *source*) to another (the *target*)
 - ***Languages are sometimes even the same(!)***
- **Interpreter**
 - A program that reads a source program and produces the results of executing that program on some input

Compiler

- Read and analyze **entire program**
- Translate to semantically equivalent program in another language
 - Presumably easier or more efficient to execute
- **Offline process**
- Tradeoff: compile time overhead
 - (preprocessing) vs execution performance

Typically Implemented with Compilers

- FORTRAN, C, C++, COBOL, and many other programming languages
- (La)TeX, SQL (databases), VHDL, many others
- Particularly appropriate if significant optimization wanted/needed

Interpreter

- Typically implemented as an “**execution engine**”
- Program analysis interleaved with execution:

```
running = true;
while (running) {
  analyze next statement;
  execute that statement; }
```
- Usually requires repeated analysis of individual statements (particularly in loops and functions)
 - But hybrid approaches can avoid some of this overhead
- But: immediate execution, good debugging/interaction...

Often Implemented with Interpreters

- Javascript, PERL, Python, Ruby, awk, sed
- Shells (bash),
- Scheme/Lisp/ML/OCaml,
- postscript/pdf,
- machine simulators
- **Particularly efficient if interpreter overhead is low relative to execution cost of individual statements**
 - But even if not (machine simulators), flexibility, immediacy, or portability may be worth it

Hybrid Approaches

- Compiler generates byte code intermediate language, e.g., compile Java source to Java Virtual Machine .class files, then
- Interpret byte codes directly, or
- Compile some or all byte codes to native code
 - **Variation: Just-In-Time compiler (JIT)** – detect hot spots & compile on the fly to native code
- Also widely use for Javascript, many functional and other languages (Haskell, ML, Racket, Ruby), C# and Microsoft Common Language Runtime

Programming Language Specs

- Since the 1960s, the syntax of every significant programming language has been specified by a **formal grammar**
 - First done in 1959 with BNF (Backus-Naur Form), used to specify ALGOL 60 syntax
 - Borrowed from the linguistics community (Chomsky)

CS 6410, Compilers – Fall 2023

Formal Languages and Automata


Formal Languages & Automata Theory

(One slide review)

- **Alphabet**: a finite set of symbols and characters
- **String**: a finite, possibly empty sequence of symbols from an alphabet
- **Language**: a set of strings (possibly empty or infinite)
- Finite specifications of (possibly infinite) languages
 - **Automaton** – a **recognizer**; a machine that accepts all strings in a language (and rejects all other strings)
 - **Grammar** – a generator; a system for producing all strings in the language (and no other strings)
- A particular language may be specified by many different grammars and automata
- A grammar or automaton specifies only one language

Chomsky's Language Hierarchy

- **Regular (Type-3)** languages are specified by regular expressions/grammars and finite automata (FSAs)
 - Specs and implementation of scanners
- **Context-free (Type-2)** languages are specified by context-free grammars and pushdown automata (PDAs)
 - Specs and implementation of parsers
- **Context-sensitive (Type-1)** languages ... aren't too important (at least for us)
- **Recursively-enumerable (Type-0)** languages are specified by general grammars and Turing machines



Backus-Naur Form (BNF)

- **Backus-Naur Form (BNF):** a syntax for describing language grammars in terms of transformation *rules*, of the form:
$$\langle \text{symbol} \rangle ::= \langle \text{expression} \rangle \mid \langle \text{expression} \rangle \dots \mid \langle \text{expression} \rangle$$
 - **Terminal:** a fundamental symbol of the language
 - **Non-terminal:** a high-level symbol describing language syntax, which can be transformed into other non-terminal or terminal symbol(s) based on the rules of the grammar.
 - Developed by two Turing-award-winning computer scientists in 1960 to describe their new ALGOL programming language

An Example BNF Grammar

`<s> ::= <n> <v>`

`<n> ::= Tamara | Emily | Daniel |`

`<v> ::= laughed | cooked | slept`

- Some sentences that could be generated from this grammar:

Tamara cooked

Emily laughed

Daniel slept

Another Example BNF Grammar

`<s> ::= <np> <v>`

`<np> ::= <pn> | <dp> <n>`

`<pn> ::= Tamara | Emily | Daniel |`

`<dp> ::= a | the`

`<n> ::= ball | hamster | carrot | computer`

`<v> ::= laughed | cooked | slept`

- Some sentences that could be generated from this grammar:

the hamster cooked

Tamara slept

A computer laughed

Yet Another Example BNF Grammar

`<s> ::= <np> <v>`

`<np> ::= <pn> | <dp> <adj> <n>`

`<pn> ::= Tamara | Daniel | Emily | Jessica`

`<dp> ::= a | the`

`<adj> ::= silly | invisible | loud | romantic`

`<n> ::= ball | hamster | carrot | computer`

`<v> ::= cried | slept | belched`

- Some sentences that could be generated from this grammar:

the invisible carrot cried

Jessica belched

a computer slept

a romantic ball belched

BNF Grammar and Recursion

`<s> ::= <np> <v>`

`<np> ::= <pn> | <dp> <adjp> <n>`

`<pn> ::= Tamara | Daniel | Emily | Jessica`

`<dp> ::= a | the`

`<adjp> ::= <adj> <adjp> | <adj>`

`<adj> ::= silly | invisible | loud | romantic`

`<n> ::= ball | hamster | carrot | computer`

`<v> ::= cried | slept | belched`

- Grammar rules can be defined *recursively*, so that the expansion of a symbol can contain that same symbol.
 - There must also be expressions that expand the symbol into something non-recursive, so that the recursion eventually ends.

Example:

Grammar for a Tiny Programming Language

$program ::= statement \mid program \ statement$

$statement ::= assignStmt \mid ifStmt$

$assignStmt ::= id = expr ;$

$ifStmt ::= \text{if } (expr) \ statement$

$expr ::= id \mid int \mid expr + expr$

$id ::= a \mid b \mid c \mid i \mid j \mid k \mid n \mid x \mid y \mid z$

$int ::= 0 \mid 1 \mid 2 \mid 3 \mid 4 \mid 5 \mid 6 \mid 7 \mid 8 \mid 9$

Exercise 1: Derive a Simple Program

```
program ::= statement | program statement  
statement ::= assignStmt | ifStmt  
assignStmt ::= id = expr ;  
ifStmt ::= if ( expr ) statement  
expr ::= id | int | expr + expr  
id ::= a | b | c | i | j | k | n | x | y | z  
int ::= 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9
```

Exercise 2: Derive Another Simple Program

```
program ::= statement | program statement  
statement ::= assignStmt | ifStmt  
assignStmt ::= id = expr ;  
ifStmt ::= if ( expr ) statement  
expr ::= id | int | expr + expr  
id ::= a | b | c | i | j | k | n | x | y | z  
int ::= 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9
```

Productions

- The rules of a grammar are called **productions**
- **Rules contain:**
 - Non-terminal symbols: grammar variables (program, statement, id, etc.)
 - Terminal symbols: concrete syntax that appears in programs (a, b, c, 0, 1, if, =, (,), ...)
- Meaning of
nonterminal ::= <sequence of terminals and non-terminals>
 - In a derivation, an instance of nonterminal can be replaced by the sequence of terminals and non-terminals on the right of the production
- Often there are several productions for a non-terminal – can choose any in different parts of derivation

Alternative Notations

- There are several notations for productions in common use; all mean the same thing

$ifStmt ::= \text{if (} expr \text{) statement}$

$ifStmt \rightarrow \text{if (} expr \text{) statement}$

$\langle ifStmt \rangle ::= \text{if (} \langle expr \rangle \text{) } \langle statement \rangle$

CS 6410, Compilers – Fall 2023


Parsing

Parsing

- **Parsing:** reconstruct the derivation (syntactic structure) of a program
- In principle, a single recognizer could work directly from a concrete, character-by-character grammar
- In practice this is never done

Parsing & Scanning

- In real compilers, the recognizer is split into two phases
 - **Scanner**: translate input characters to tokens
 - Also, report lexical errors like illegal characters and illegal symbols
 - **Parser**: read token stream and reconstruct the derivation



Why Separate the Scanner and Parser?

- **Simplicity & Separation of Concerns**
 - Scanner hides details from parser (comments, whitespace, input files, etc.)
 - Parser is easier to build; has simpler input stream (tokens) and simpler interface for input
- **Efficiency**
 - Scanner recognizes regular expressions – proper subset of context free grammars
 - (But still often consumes a surprising amount of the compiler's total execution time)

But ...

- Not always possible to separate cleanly
- Example: C/C++/Java *type vs identifier*
 - Parser would like to know which names are types and which are identifiers, but...
 - Scanner doesn't know how things are declared
- So, we hack around it somehow...
 - Either use simpler grammar and disambiguate later, or communicate between scanner & parser
 - Engineering issue: try to keep interfaces as simple & clean as possible

Typical Tokens in Programming Languages

- Operators & Punctuation
 - + - * / () { } [] ; : :: < <= == = != ! ...
 - Each of these is a distinct lexical class
- Keywords
 - if while for goto return switch void ...
 - Each of these is also a distinct lexical class (*not* a string)
- Identifiers
 - A single ID lexical class, but parameterized by actual id
- Integer constants
 - A single INT lexical class, but parameterized by int value
- Other constants, etc.

Principle of Longest Match

- In most languages, the scanner should pick the longest possible string to make up the next token if there is a choice
- Example

return maybe != iffy;
should be recognized as 5 tokens

RETURN	ID(maybe)	NEQ	ID(iffy)	SCOLON
--------	-----------	-----	----------	--------

i.e., != is one token, not two; “iffy” is an ID, not IF followed by ID(fy)

Lexical Complications

- Most modern languages are free-form
 - Layout doesn't matter
 - Whitespace separates tokens
- Alternatives
 - Fortran – line oriented
 - Haskell, Python – indentation and layout can imply grouping
- And other confusions
 - In C++ or Java, is >> a shift operator or the end of two nested templates or generic classes?

CS 6410, Compilers – Fall 2023

Regular Expressions

Regular Expressions

- The lexical grammar (structure) of most programming languages can be specified with regular expressions
 - (Sometimes a little cheating is needed)
- Tokens can be recognized by a deterministic finite automaton
 - Can be either table-driven or built by hand based on lexical grammar

Regular Expressions

- Defined over some alphabet Σ
 - For programming languages, alphabet is usually ASCII or Unicode
- If re is a regular expression, $L(re)$ is the language (set of strings) generated by re

Fundamental REs

re	$L(re)$	Notes
a	$\{ a \}$	Singleton set, for each a in Σ
ε	$\{ \varepsilon \}$	Empty string
\emptyset	$\{ \}$	Empty language

Operations on REs

re	$L(re)$	Notes
rs	$L(r)L(s)$	Concatenation
$r s$	$L(r) \cup L(s)$	Combination (union)
r^*	$L(r)^*$	0 or more occurrences (Kleene closure)

- **Precedence:** $*$ (highest), concatenation, $|$ (lowest)
- Parentheses can be used to group REs as needed

Examples

<i>re</i>	Meaning
+	single + character
!	single ! character
=	single = character
!=	2 character sequence "!="
xyzzzy	5 character sequence "xyzzzy"
$(1 0)^*$	0 or more binary digits
$(1 0)(1 0)^*$	1 or more binary digits
$0 1(0 1)^*$	sequence of binary digits with no leading 0's, except for 0 itself

Abbreviations

- The basic operations generate all possible regular expressions, but there are common abbreviations used for convenience. Some examples:

Abbr.	Meaning	Notes
r^+	(rr^*)	1 or more occurrences
$r?$	$(r \mid \epsilon)$	0 or 1 occurrence
$[a-z]$	$(a \mid b \mid \dots \mid z)$	1 character in given range
$[abxyz]$	$(a \mid b \mid x \mid y \mid z)$	1 of the given characters

More Examples

<i>re</i>	Meaning
[abc]+	
[abc]*	
[0-9]+	
[1-9][0-9]*	
[a-zA-Z][a-zA-Z0-9_]*	

Abbreviations

- Many systems allow abbreviations to make writing and reading definitions or specifications easier

name ::= *re*

- Restriction: abbreviations may not be circular (recursive) either directly or indirectly (else would be non-regular)

Example

- Possible syntax for numeric constants

digit ::= [0-9]

digits ::= *digit*+

number ::= *digits* (. *digits*)?

([eE] (+ | -)? *digits*) ?

- How would you describe this set in English?
- What are some examples of legal constants (strings) generated by *number* ?
 - What are the differences between these and numeric constants in YFPL? (Your Favorite Programming Language)

CS 6410, Compilers – Fall 2023

Recognizing Res and Finite Automata


Recognizing REs

- Finite automata can be used to recognize strings generated by regular expressions
- Can build by hand or automatically
 - Reasonably straightforward, and can be done systematically
 - Tools like Lex, Flex, JFlex et seq do this automatically, given a set of REs

Finite State Automaton

- A finite set of states
 - One marked as initial state
 - One or more marked as accepting (final) states
 - States sometimes labeled or numbered
- A set of transitions from state to state
 - Each labeled with symbol from Σ , or ϵ
 - Common to allow multiple labels (symbols) on one edge to simplify diagrams
- Operate by reading input symbols (usually characters)
 - Transition can be taken if labeled with current symbol
 - ϵ -transition can be taken at any time
- Accept when final state reached & no more input
 - Slightly different in a scanner where the FSA is a subroutine that accepts the longest input string matching a token regular expression, starting at the current location in the input
- Reject if no transition possible, or no more input and not in final state (DFA)
 - Some versions require an explicit “error” state and transitions to it on all “no legal transition possible” input. OK to omit that for CSE 401

Example: FSA for “cat”

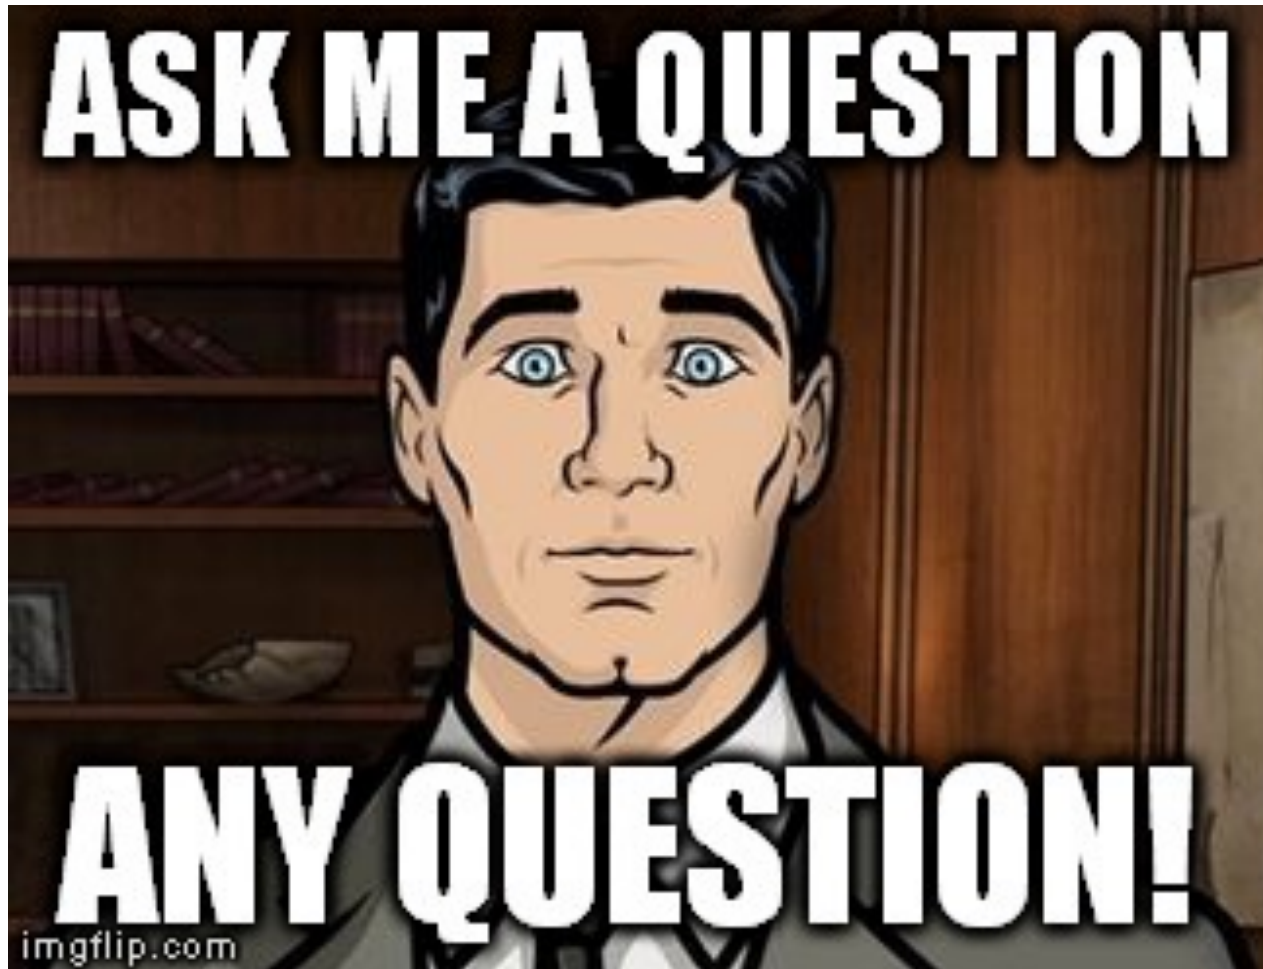


DFA vs NFA

- **Deterministic Finite Automata (DFA)**
 - No choice of which transition to take under any condition
 - No ϵ transitions (arcs)
- **Non-deterministic Finite Automata (NFA)**
 - Choice of transition in at least one case
 - Accept if some way to reach a final state on given input
 - Reject if no possible way to final state
 - i.e., may need to guess right path or backtrack

Coming Attractions

- **First homework:** paper exercises on regular expressions, automata, etc.
- **Then:** first part of the compiler assignment – the scanner
- **Next topic:** more scanning



[Meme credit: imgflip.com]