

CS 6410: Compilers

Fall 2023

Tamara Bonaci
t.bonaci@northeastern.edu

Thank you to UW faculty Hal Perkins. Today lecture notes are a modified version of his lecture notes.

Credits For Course Material

- Big thank you to UW CSE faculty member, Hal Perkins
- Some direct ancestors of this course:
 - UW CSE 401 (Chambers, Snyder, Notkin, Perkins, Ringenburt, Henry, ...)
 - UW CSE PMP 582/501 (Perkins)
 - Cornell CS 412-3 (Teitelbaum, Perkins)
 - Rice CS 412 (Cooper, Kennedy, Torczon)
 - Many books (Appel; Cooper/Torczon; Aho, [[Lam,] Sethi,] Ullman [Dragon Book], Fischer, [Cytron ,] LeBlanc; Muchnick, ...)

Agenda

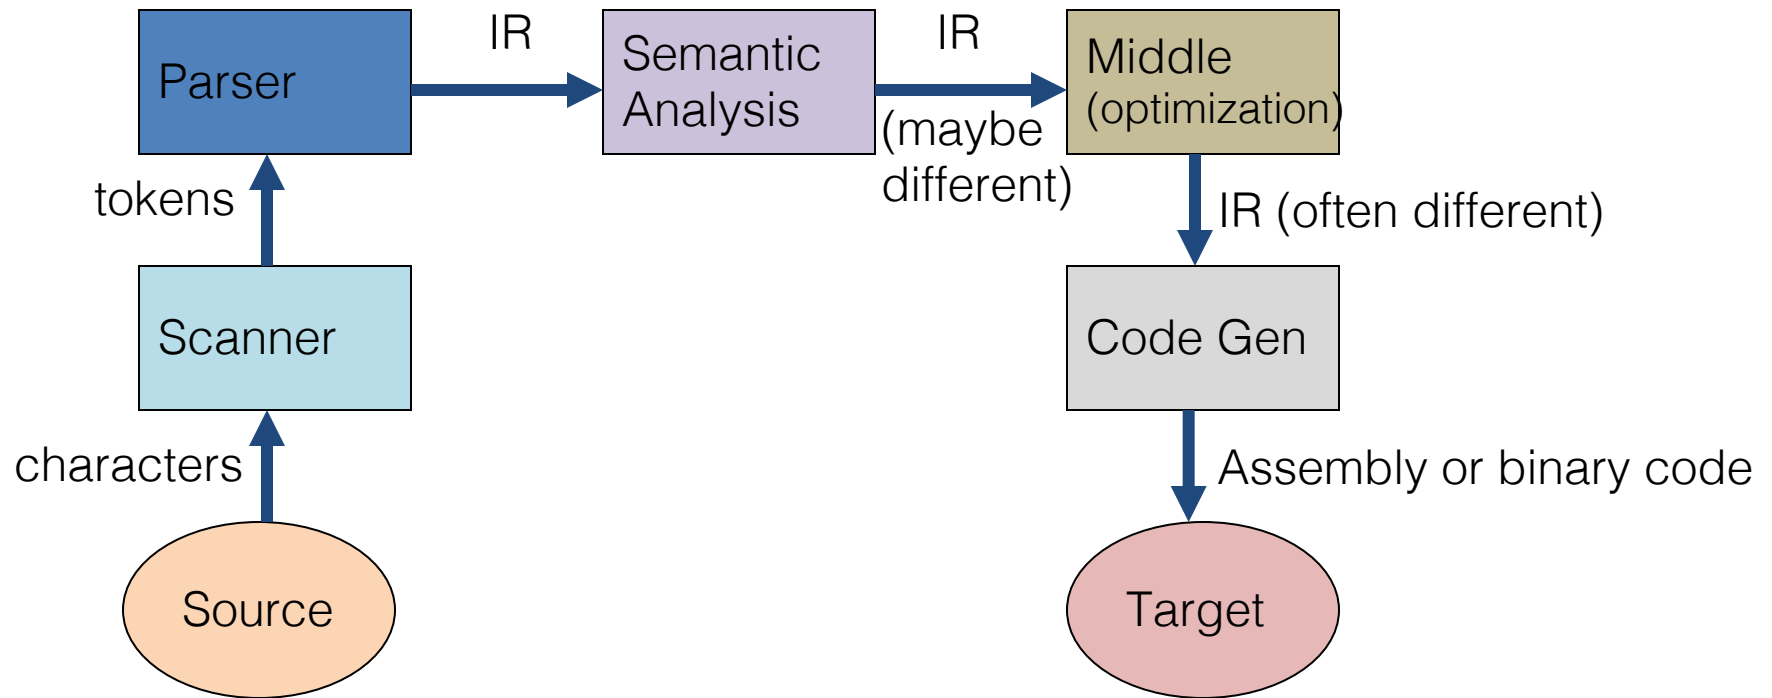
- Static semantics
- Attribute grammars
- Symbol tables
- Types & type checking
- Intermediate representations

Reading:

- Cooper & Torczon – chapter 3 and 5
- The Dragon book, chapters 4 and 6.1, 6.2

Semantics

Review: Compiler Structure



Review: Semantic Analysis

- Main tasks:
 - Extract types and other information from the program
 - Check language rules that go beyond the context-free grammar
 - **Resolve names** – connect declarations and uses
 - “Understand” the program – last phase of front end ...
 - ... so, program is “correct” for hand-off to back end
- Key data structure: **Symbol tables**
 - For each identifier in the program, record its attributes (kind, type, etc.)
 - Later: assign storage locations (stack frame offsets) for variables, add other annotations

Review: A Sampling of Semantic Checks (0)

- Appearance of a name: *id*
 - Check: *id* has been declared and is in scope
 - Compute: Inferred type of *id* is its declared type
- Constant: *v*
 - Compute: Inferred type and value are explicit

Review: A Sampling of Semantic Checks (1)

- Binary operator: $\text{exp}_1 \text{ op } \text{exp}_2$
 - **Check:** exp_1 and exp_2 have compatible types
 - Identical, or
 - Well-defined conversion to appropriate types
 - **Compute:** Inferred type is a function of the operator and operand types

Review: A Sampling of Semantic Checks (2)

- Assignment: $\text{exp}_1 = \text{exp}_2$
 - **Check:** exp_1 is assignable (not a constant or expression)
 - **Check:** exp_1 and exp_2 have (assignment-)compatible types
 - Identical, or
 - exp_2 can be converted to exp_1 (e.g., char to int), or
 - Type of exp_2 is a subclass of type of exp_1 (can be decided at compile time)
 - **Compute:** Inferred type is type of exp_1

Review: A Sampling of Semantic Checks (3)

- Cast: $(exp_1) exp_2$
 - Check: exp_1 is a type
 - Check: exp_2 either
 - Has same type as exp_1
 - Can be converted to type exp_1 (e.g., double to int)
 - **Downcast**: is a superclass of exp_1 (usually requires a runtime check to verify; at compile time we can at least decide if it could be true)
 - **Upcast (Trivial)**: is the same or a subclass of exp_1
 - Compute: Inferred type is exp_1

Review: A Sampling of Semantic Checks (4)

- Field reference: `exp.f`
 - Check: `exp` is a reference type (not value type)
 - Check: The class of `exp` has a field named `f`
 - Compute: Inferred type is declared type of `f`

Review: A Sampling of Semantic Checks (5)

- **Method call:** $\text{exp.m}(e_1, e_2, \dots, e_n)$
 - **Check:** exp is a reference type (class instance)
 - **Check:** The class of exp has a method named m
 - **Check:** The method exp.m has n parameters
 - Or, if overloading allowed, at least one version of m exists with n parameters
 - **Check:** Each argument has a type that can be assigned to the associated parameter
 - Same “assignment compatible” check for assignment
 - Overloading: need to find a “best match” among available methods if more than one is compatible – or reject if result is ambiguous (e.g., C++, others)
 - **Compute:** Inferred type is given by method declaration (or could be void)

Review: A Sampling of Semantic Checks (6)

- Return statement: `return exp;` or: `return;`
- Check:
 - If the method is not void: The expression can be assigned to a variable with the declared return type of the method – exactly the same test as for assignment statement
 - If the method is void: There is no expression

Attribute Grammars

Attribute Grammars

- A systematic way to think about semantic analysis
- Formalize properties checked and computed during semantic analysis and relate them to grammar productions in the CFG (or AST)
- Sometimes used directly, but even when not, AGs are a useful way to organize the analysis and think about it

Attribute Grammars

- **Idea:** associate attributes with each node in the (abstract) syntax tree
- **Examples of attributes**
 - Type information
 - Storage location
 - Assignable (e.g., expression vs variable – lvalue vs rvalue in C/C++ terms)
 - Value (for constant expressions)
- **Notation:** $X.a$ if a is an attribute of node X

Inherited and Synthesized Attributes

Given a production $X ::= Y_1 Y_2 \dots Y_n$

- A *synthesized* attribute $X.a$ is a function of some combination of the attributes of the Y_i 's (bottom up)
- An *inherited* attribute $Y_i.b$ is a function of some combination of attributes $X.a$ and other $Y_j.c$ (top down)
 - Often restricted a bit: only Y 's to the left can be used (has implications for evaluation)

Attribute Equations

- For each kind of node we give a set of equations relating attribute values of the node and its neighbors (usually children)
 - Example: $\text{plus.val} = \text{exp}_1.\text{val} + \text{exp}_2.\text{val}$
- Attribution (evaluation) means implicitly finding a solution that satisfies all of the equations in the tree
 - This is an example of a constraint language

Informal Example of Attribute Rules (1)

- Suppose we have the following grammar for a trivial language

```
program ::= decl stmt
decl ::= int id;
stmt ::= exp = exp ;
exp ::= id | exp + exp | 1
```
- What attributes would we create to check types and assignability?

Informal Example of Attribute Rules (2)

- Attributes of nodes
 - env (environment, e.g., symbol table)
 - Synthesized by decl, inherited by stmt
 - Each entry maps a name to its type and kind
 - type (expression type)
 - synthesized
 - kind (variable [var or lvalue] vs value [val or rvalue])
 - synthesized

Attributes for Declarations

$\text{decl} ::= \text{int id};$
 $\text{decl.env} = \{\text{id} \rightarrow (\text{int}, \text{var})\}$

Attributes for Program

program ::= decl stmt
 stmt.env = decl.env

Attributes for Constants

$\text{exp} ::= 1$

$\text{exp.kind} = \text{val}$

$\text{exp.type} = \text{int}$

Attributes for Identifier Expressions

$\text{exp} ::= \text{id}$

$(\text{type}, \text{kind}) = \text{exp.env.lookup}(\text{id})$

$\text{exp.type} = \text{type}$ (i.e., id type)

$\text{exp.kind} = \text{kind}$ (i.e., id kind)

Attributes for Addition

$\text{exp} ::= \text{exp}_1 + \text{exp}_2$

$\text{exp}_1.\text{env} = \text{exp}.\text{env}$

$\text{exp}_2.\text{env} = \text{exp}.\text{env}$

error if $\text{exp}_1.\text{type} \neq \text{exp}_2.\text{type}$

(or error if not compatible, depending on language rules)

$\text{exp}.\text{type} = \text{exp}_1.\text{type}$ (or $\text{exp}_2.\text{type}$)

$\text{exp}.\text{kind} = \text{val}$

Attribute Rules for Assignment

$\text{stmt} ::= \text{exp}_1 = \text{exp}_2;$

$\text{exp}_1.\text{env} = \text{stmt}.\text{env}$

$\text{exp}_2.\text{env} = \text{stmt}.\text{env}$

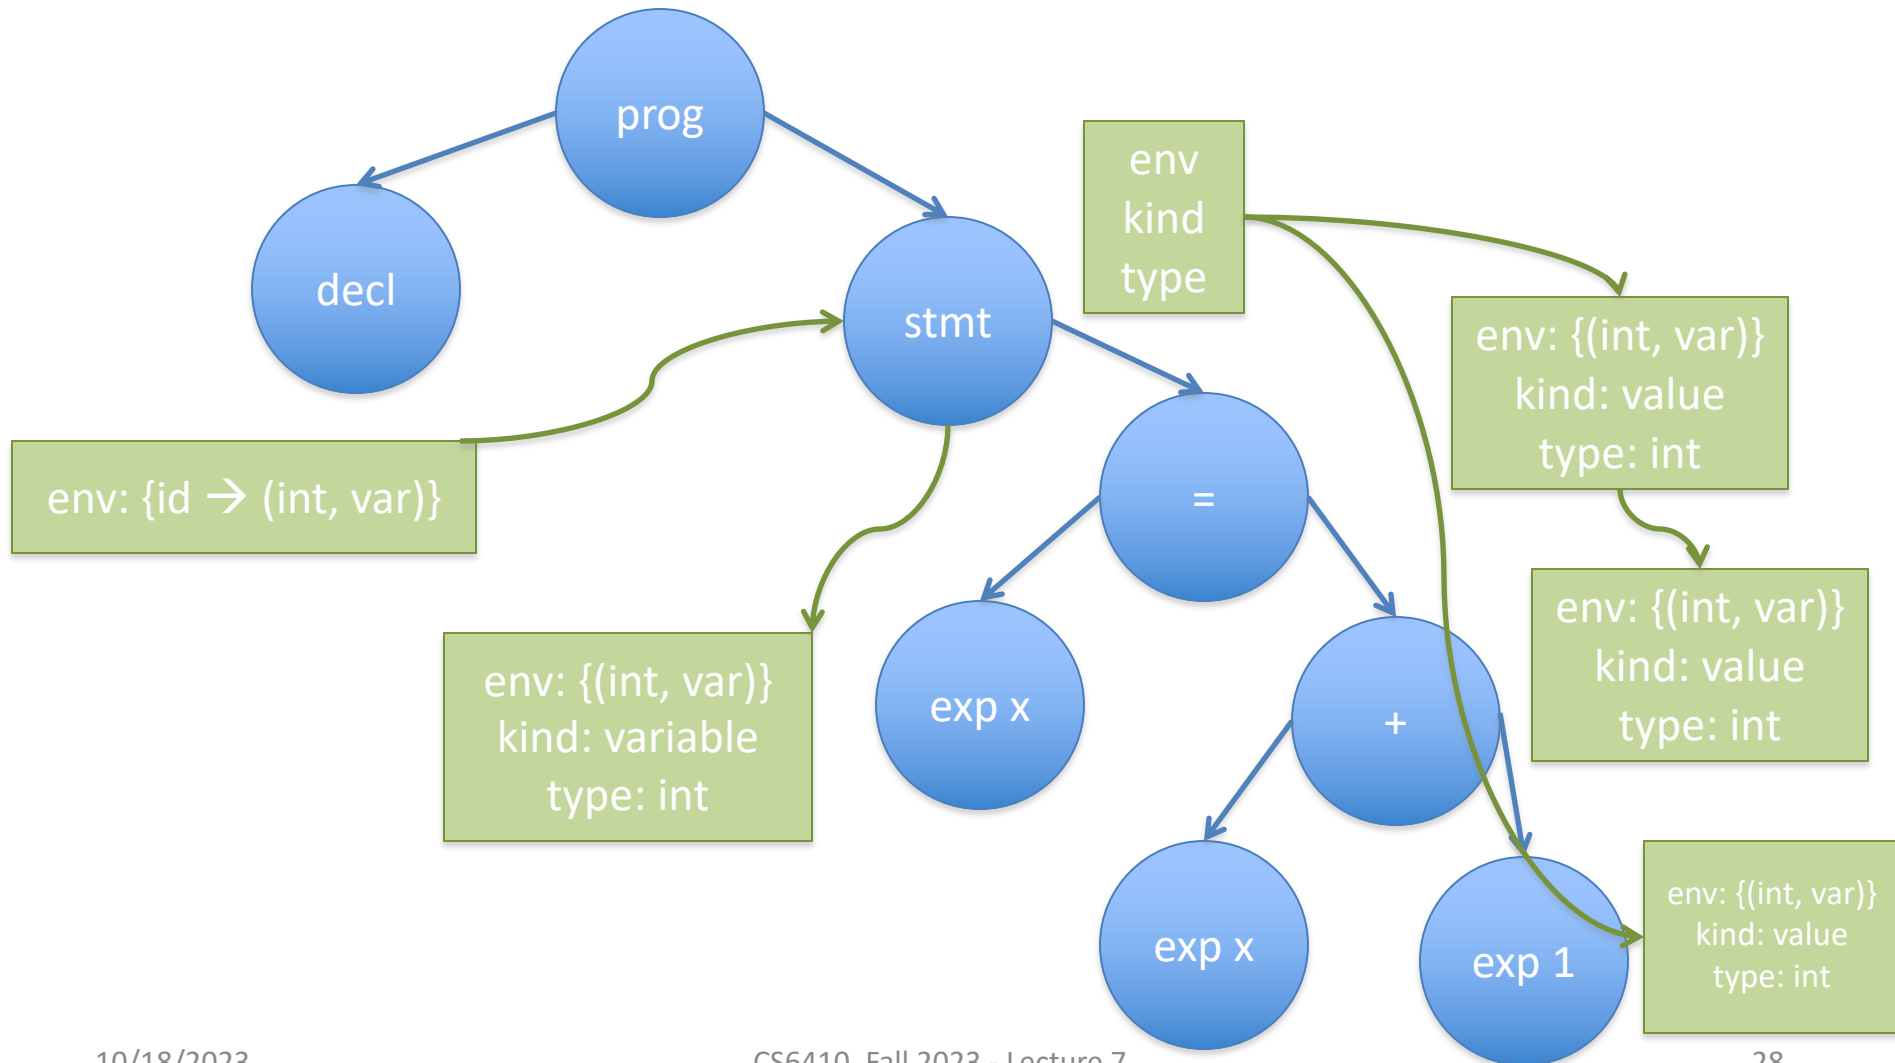
Error if $\text{exp}_2.\text{type}$ is not assignment compatible
with $\text{exp}_1.\text{type}$

Error if $\text{exp}_1.\text{kind}$ is not var (can't be val)

Example

```
int x; x = x + 1;
```

Example: `int x; x = x + 1;`



Extensions

- This can be extended to handle sequences of declarations and statements
 - Sequences of declarations builds up larger environments
 - Each declaration synthesizes a new environment from previous one, plus the new binding
 - Full environment is passed down to statements and expressions

Observations

- These are equational computations
 - Think functional programming, no side effects
- Solver can be automated, provided the attribute equations are non-circular
- But implementation problems
 - Non-local computation
 - Can't afford to literally pass around copies of large, aggregate structures like environments

In Practice

- Attribute grammars give us a good way of thinking about how to structure semantic checks
- Symbol tables will hold environment information
- Add fields to AST nodes to refer to appropriate attributes (symbol table entries for identifiers, types for expressions, etc.)
 - Put in appropriate places in AST class inheritance tree and exploit inheritance. Most statements don't need types, for example, but all expressions do.

Symbol Tables

Symbol Tables

- Map identifiers to
<type, kind, location, other properties>
- Operations
 - Lookup(id) => information
 - Enter(id, information)
 - Open/close scopes
- Build & use during semantics pass
 - Build first from declarations
 - Then use to check semantic rules
- Use (and augment) in later compiler phases

Aside: Implementing Symbol Tables

- Big topic in classical (i.e., ancient) compiler courses: implementing a hashed symbol table
- **These days:** use the collection classes that are provided with the standard language libraries (Java, C#, C++, ML, Haskell, etc.)
 - Then tune & optimize if it really matters
 - In production compilers, it really matters
 - Up to a point...
- **Java:**
 - Map (HashMap) will handle most cases
 - List (ArrayList) for ordered lists (parameters, etc.)

Symbol Tables for MiniJava

- We'll outline a scheme that does what we need, but feel free to modify/adapt as needed
- Mix of global and local tables

Symbol Tables for MiniJava: Global

- Global – Per Program Information
 - Single global table to map class names to per-class symbol tables
 - Created in a pass over class definitions in AST
 - Used in remaining parts of compiler to check class types and their field/method names and extract information about them

Symbol Tables for MiniJava: Class

- One symbol table for each class
 - One entry per method/field declared in the class
 - **Contents:** type information, public/private, parameter types (for methods), storage locations (later), etc
- Reached from global table of class names
- In Java, we actually need multiple symbol tables (or more complex symbol table) per class
 - The same identifier can be used for both a method name and a field name in a single class

Symbol Tables for MiniJava: Global/Class

- All global tables persist throughout the compilation
 - And beyond in a real compiler...
 - Symbolic information in Java .class or MSIL files, link-time optimization information in gcc)
 - Debug information in .o and .exe files
 - Some or all information in library files (.a, .so)
 - Type information for garbage collector

Symbol Tables for MiniJava: Methods

- One local symbol table for each method
 - One entry for each local variable or parameter
 - **Contents:** type info, storage locations (later), etc
 - Needed for project only while compiling the method; can discard when done in a single pass compiler
 - But if type checking and code gen, etc. are done in separate passes, this table needs to persist until we're done with it
 - And beyond: often need type info for runtime debugging, memory management/garbage collection, etc
 - Even for our project, the MiniJava compiler will likely have multiple passes

Beyond MiniJava

- What we aren't dealing with: nested scopes
 - Inner classes
 - Nested scopes in methods – reuse of identifiers in parallel or inner scopes; nested functions (ML, ...)
 - Lambdas and function closures
- **Basic idea:** new symbol table for inner scopes, linked to surrounding scope's table (i.e., stack of symbol tables, top = current innermost scope)
 - Look for identifier in inner scope; if not found look in surrounding scope (recursively)
 - Pop symbol table when we exit a scope
- Also ignoring static fields/methods, accessibility (public, protected, private), package scopes, ...

Engineering Issues (1)

- In multipass compilers, inner scope symbol tables need to persist for use in later passes
 - So really can't delete symbol tables on scope exit
 - Retain and add a pointer to the parent scope (effectively a reverse tree of scope symbol tables with root = global table)
 - Keep a pointer to current innermost scope (leaf) and start looking for symbols there

Engineering Issues (2)

- In practice, want to retain $O(1)$ lookup or something close to it
 - Would like to avoid $O(\text{depth of scope nesting})$, although some compilers assume this will be small enough not to matter
 - When it matters, use hash tables with additional information (linked lists of various sorts) to get the scope nesting right
 - Scope entry/exit operators

Error Recovery

- What to do when an undeclared identifier is encountered?
 - Only complain once (Why?)
 - Can forge a symbol table entry for id once you've complained so it will be found in the future
 - Assign the forged entry a type of “unknown”
 - “Unknown” is the type of all malformed expressions and is compatible with all other types
 - Allows you to only complain once! (How?)

“Predefined” Things

- Many languages have some “predefined” items (constants, functions, classes, namespaces, standard libraries, ...)
- Include initialization code or declarations to manually create symbol table entries for these when the compiler starts up
 - Rest of compiler generally doesn’t need to know the difference between “predeclared” items and ones found in the program
 - Can put “standard prelude” information in a file or data resource and use that to initialize
 - Tradeoffs?

Types and Type Checking

Types

- Classical roles of types in programming languages
 - Run-time safety
 - Compile-time error detection
 - Improved expressiveness (method or operator overloading, for example)
 - Provide information to optimizer
 - In strongly typed languages, allows compiler to make assumptions about possible values

Type Checking Terminology

Static vs. dynamic typing

- **Static:** checking done prior to execution (e.g. compile-time)
- **Dynamic:** checking during execution

Strong vs. weak typing

- **Strong:** guarantees no illegal operations performed
- **Weak:** can't make guarantees

Caveats:

- Hybrids common
- Inconsistent usage common
- “untyped,” “typeless” could mean dynamic or weak

	static	dynamic
strong	Java, SML	Scheme, Ruby
weak	C	PERL

Type Systems

- Base Types
 - Fundamental, atomic types
 - Typical examples: int, double, char, bool
- Compound/Constructed Types
 - Built up from other types (recursively)
 - Constructors include records/structs/classes, arrays, pointers, enumerations, functions, modules, ...
 - Most language provide a small collection of these

How to Represent Types in a Compiler?

Create a shallow class hierarchy

- Example:

```
abstract class Type { ... }    // or  
interface  
    class BaseType extends Type { ... }  
    class ClassType extends Type { ... }
```

- Should not need too many of these

Types vs ASTs

- Types nodes are not AST nodes!
- **AST** = abstract representation of source program (including source program type info)
- **Types** = abstract representation of type semantics for type checking, inference, etc.
 - Can include information not explicitly represented in the source code, or may describe types in ways more convenient for processing
- Be sure you have a separate “type” class hierarchy in your compiler distinct from the AST

Base Types

- For each base type create exactly one object to represent it (singleton pattern!)
 - Symbol table entries and AST nodes reference these objects to represent entry/node types
 - Usually created at compiler startup
- Useful to create a type “void” object to tag functions that do not return a value
- Also useful to create a type “unknown” object for errors
 - (“void” and “unknown” types reduce the need for special case code in various places in the type checker; don’t have to return “null” for “no type” or “not declared” cases)

Compound Types

- Basic idea: use a appropriate “type constructor” object that refers to the component types
 - Limited number of these – correspond directly to type constructors in the language (pointer, array, record/struct/class, function)
 - So a compound type is represented as a graph
- Some examples...

Class Types

- Type for: class Id { fields and methods }

```
class ClassType extends Type {  
    Type baseClassType;           // ref to base class  
    Map fields;                   // type info for  
    fields  
    Map methods;                  // type info for  
    methods  
}
```

(MiniJava note: May not want to represent class types exactly like this, depending on how class symbol tables are represented; e.g., the class symbol table(s) might be a sufficient representation of a class type.)

Array Types

- For regular Java this is simple: only possibility is # of dimensions and element type (which can be another array type or anything else)

```
class ArrayType extends Type {  
    int nDims;  
    Type elementType;  
}
```

Methods/Functions

- Type of a method is its result type, plus an ordered list of parameter types

```
class MethodType extends Type {  
    Type resultType;           // type or  
    "void"  
    List parameterTypes;  
}
```

- Sometimes called the method “signature”

Type Equivalence

- For base types this is simple: types are the same if they are identical
 - Can use pointer comparison in the type checker if you have a singleton object for each base type
- Normally there are well defined rules for coercions between arithmetic types
 - Compiler inserts these automatically where required by the language spec or when written explicitly by programmer (casts) – often involves inserting cast or conversion nodes in AST

Type Equivalence for Compound Types

- Two basic strategies
 - *Structural equivalence*: two types are the same if they are the same kind of type and their component types are equivalent, recursively
 - *Name equivalence*: two types are the same only if they have the same name, even if their structures match
- Different language design philosophies
 - e.g., are Complex and Point the same?
 - e.g., are Point (Cartesian) and Point (Polar) the same?

Structural Equivalence

- Structural equivalence says two types are equal iff they have same structure
 - Atomic types are tautologically the same structure and equal if they are the same type
 - For type constructors: equal if the same constructor and, recursively, type (constructor) components are equal
- Ex: atomic types, array types, ML record types
- Implement with recursive implementation of equals, or by canonicalization of types when types created, then use pointer/ref. equality

Name Equivalence

- Name equivalence says that two types are equal iff they came from the same textual occurrence of a type constructor
 - Ex: class types, C struct types (struct tag name), datatypes in ML
 - special case: type synonyms (e.g. typedef in C) do not define new types
- Implement with pointer equality assuming appropriate representation of type info

Type Equivalence and Inheritance

- Suppose we have

```
class Base { ... }
```

```
class Derived extends Base { ... }
```

- A variable declared with type Base has a *compile-time type* or *static type* of Base
- During execution, that variable may refer to an object of class Base or any of its subclasses like Derived (or can be null), often called the the *runtime type* or *dynamic type*
 - Since subclass is guaranteed to have all fields/methods of base class, type checker only needs to deal with declared compile-time types of variables and, in fact, can't track all possible runtime types

Type Casts

- In most languages, one can explicitly cast an object of one type to another
 - Sometimes cast means a conversion (e.g., casts between numeric types)
 - Sometimes cast means a change of static type without doing any computation (casts between pointer types or pointer and numeric types in C)
 - For objects can be a upcast (free and always safe) or downcast (requires runtime check to be safe)

Type Conversions and Coercions

- In full Java, we can explicitly convert a value of type double to one of type int
 - can represent as unary operator
 - typecheck, codegen normally
- In full Java, can implicitly coerce an value of type int to one of type double
 - compiler must insert unary conversion operators, based on result of type checking

C and Java: type casts

- In C/C++: safety/correctness of casts not checked
 - Allows writing low-level code that's type-unsafe
 - C++ has more elaborate casts, and at least one of them does imply runtime checks
- In Java: downcasts from superclass to subclass need runtime check to preserve type safety
 - static typechecker allows the cast
 - codegen introduces runtime check
 - (same code needed to handle “instanceof”)
 - Java's main need for dynamic type checking

Various Notions of Type Compatibility

- There are usually several relations on types that we need to analyze in a compiler:
 - “is the same as”
 - “is assignable to”
 - “is same or a subclass of”
 - “is convertible to”
- Exact meanings and checks needed depend on the language specifications
- Be sure to check for the right one(s)

Useful Compiler Functions

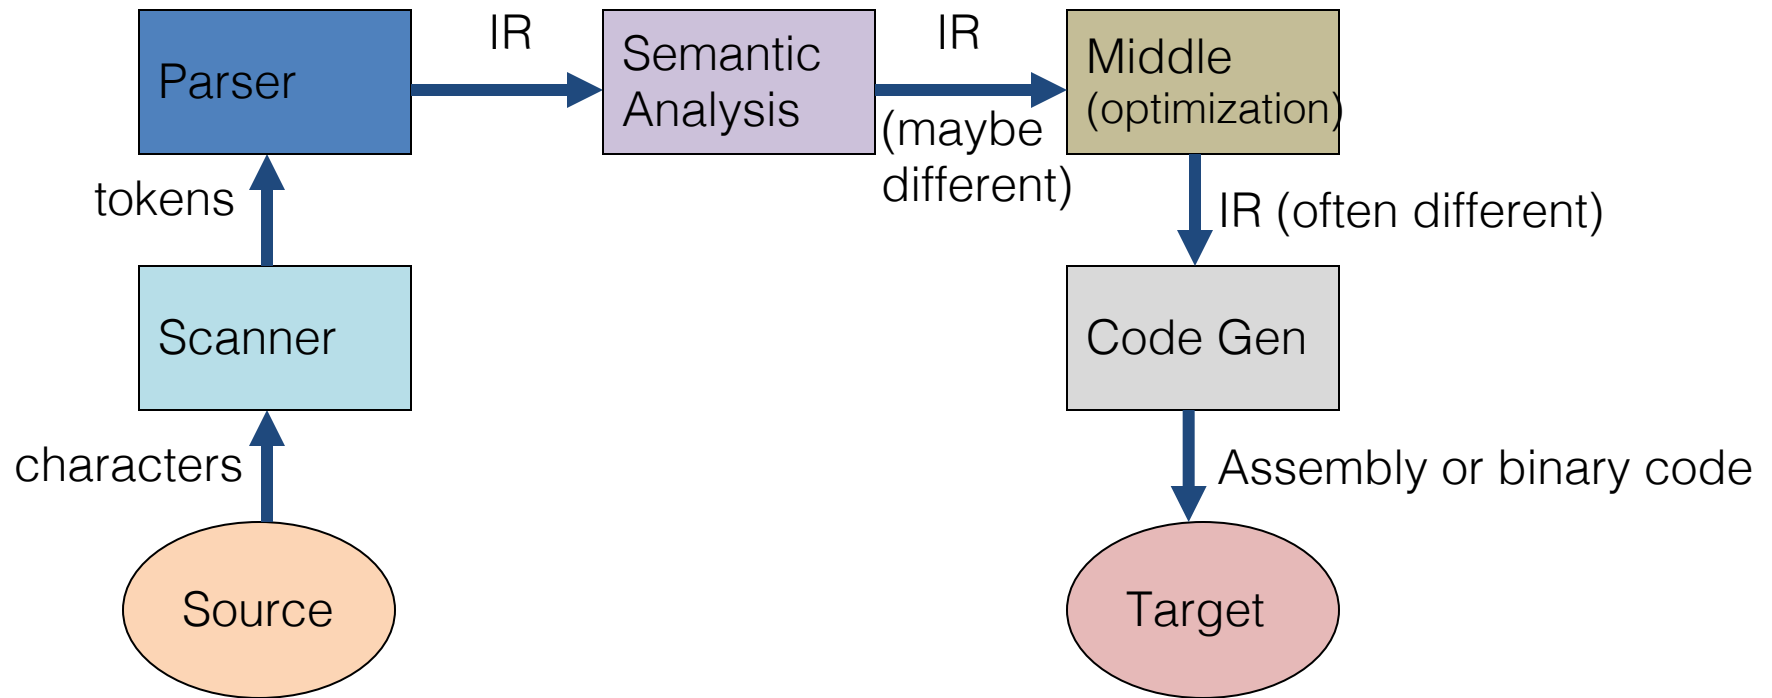
- Create a handful of methods to decide different kinds of type compatibility:
 - Types are identical
 - Type t_1 is assignment compatible with t_2
 - Parameter list is compatible with types of expressions in the method call
- Usual modularity reasons: isolates these decisions in one place and hides the actual type representation from the rest of the compiler
- Probably belongs in the same package with the type representation classes

Implementing Type Checking for MiniJava

- Create multiple visitors for the AST
- First pass/passes: gather information
 - Collect global type information for classes
 - Could do this in one pass, or might want to do one pass to collect class information, then a second one to collect per-class information about fields, methods – you decide
- Next set of passes: go through method bodies to check types, other semantic constraints

Intermediate Representations

Review: Compiler Structure



Intermediate Representations

- In most compilers:
 - The parser builds an **intermediate representation** (IR, typically an Abstract Syntax Tree)
 - Rest of the compiler transforms the IR to optimize it
 - Typically will transform initial IR to one or more different IRs along the way
 - IR eventually translate to final target code

IR Design

- Decisions affect speed and efficiency of the rest of the compiler
 - General rule: compile time is important, but performance of generated code often more important
 - Typical case for production code: compile a few times, run many times
 - Although the reverse is true during development
 - So make choices that improve compile time as long as they don't compromise the result

IR Design

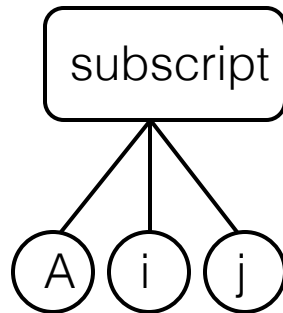
- Desirable properties
 - Easy to generate
 - Easy to manipulate
 - Expressive
 - Appropriate level of abstraction
- Different tradeoffs depending on compiler goals
- Different tradeoffs in different parts of the same compiler
 - So often different IRs in different parts

IR Design Taxonomy

- Structure
 - Graphical (trees, graphs, etc.)
 - Linear (code for some abstract machine)
 - Hybrids are common (e.g., control-flow graphs whose nodes are basic blocks of linear code)
- Abstraction Level
 - High-level, near to source language
 - Low-level, closer to machine (exposes more details to compiler)

Examples: Array Reference

source: $A[i,j]$



$t1 \leftarrow A[i,j]$

```
loadl 1 => r1
sub rj,r1 => r2
loadl 10 => r3
mult r2,r3 => r4
sub ri,r1 => r5
add r4,r5 => r6
loadl @A => r7
add r7,r6 => r8
load r8 => r9
```

Levels of Abstraction

- Key design decision: how much detail to expose
 - Affects possibility and profitability of various optimizations
 - Depends on compiler phase: some semantic analysis & optimizations are easier with high-level IRs close to the source code. Low-level usually preferred for other optimizations, register allocation, code generation, etc.
 - Structural (graphical) IRs are typically fairly high-level
 - Linear IRs are typically low-level

Graphical IRs

- IR represented as a graph (or a tree)
- Nodes and edges typically reflect some structure of the program
 - E.g., source code, control flow, data dependence
- May be large (especially syntax trees)
- High-level examples:
 - Syntax trees
 - DAGs
 - Generally used in early phases of compilers
- Other examples:
 - Control flow graphs,
 - Data dependency graphs
 - Often used in optimization and code generation

Concrete Syntax Trees

- The full grammar is needed to guide the parser, but contains many extraneous details
 - Chain productions
 - Rules that control precedence and associativity
- Typically the full syntax tree (parse tree) does not need to be used explicitly, but sometimes we want it (structured source code editors or transformations, ...)

Example

- Concrete syntax for $x = 2^*(n+m)$

assign ::= *id* = *expr* ;

expr ::= *expr* + *term* | *expr* - *term* | *term*

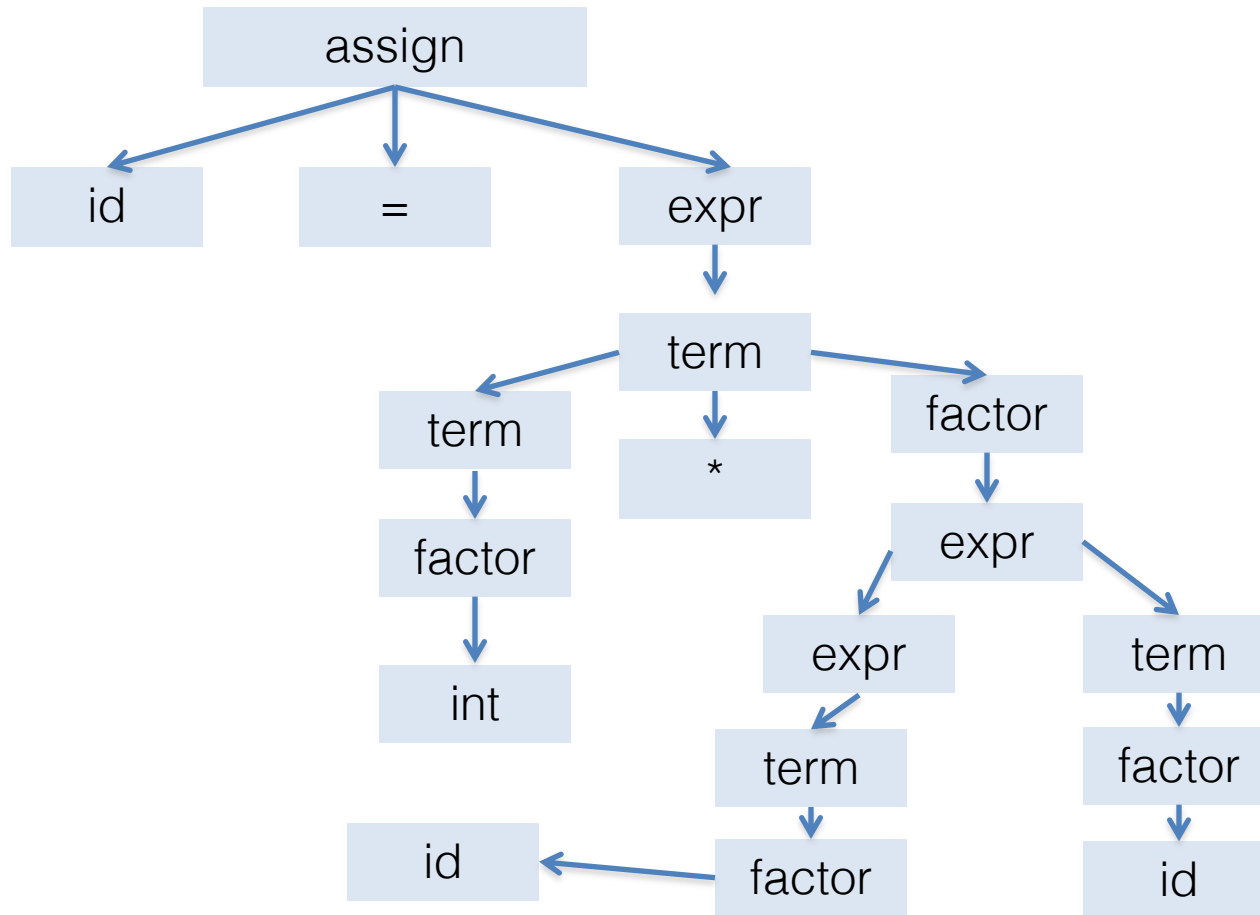
term ::= *term* * *factor* | *term* / *factor* | *factor*

factor ::= *int* | *id* | (*expr*)

$assign ::= id = expr;$ $expr ::= expr + term \mid expr - term \mid term$ $term ::= term * factor \mid term / factor \mid factor$ $factor ::= int \mid id \mid (expr)$

Example

- Concrete syntax for $x = 2^*(n+m)$



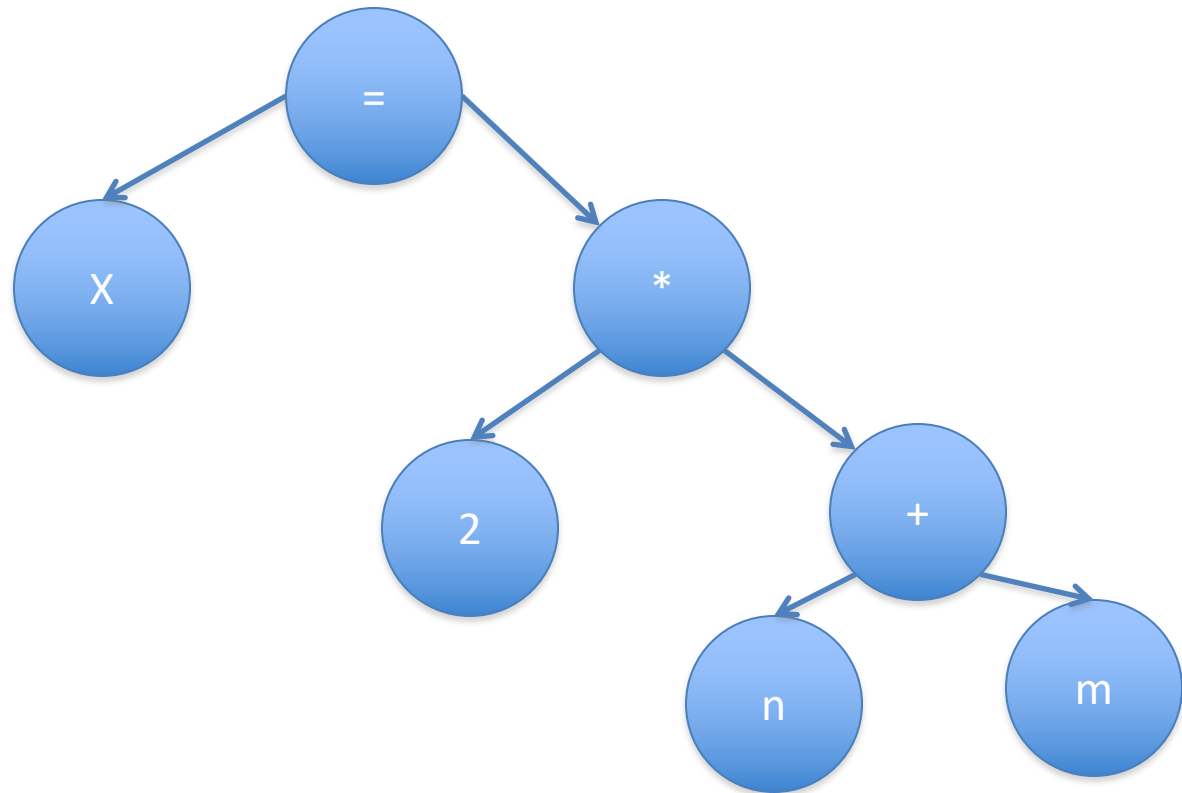
Abstract Syntax Trees

- Common output from parser:
 - Used for static semantics (type checking, etc)
 - Sometimes used high-level optimizations
- Focus on essential structural information
- Can be represented:
 - Explicitly as a tree or
 - In a linear form
- Example: LISP/Scheme S-expressions are essentially ASTs

$assign ::= id = expr;$
 $expr ::= expr + term \mid expr - term \mid term$
 $term ::= term * factor \mid term / factor \mid factor$
 $factor ::= int \mid id \mid (expr)$

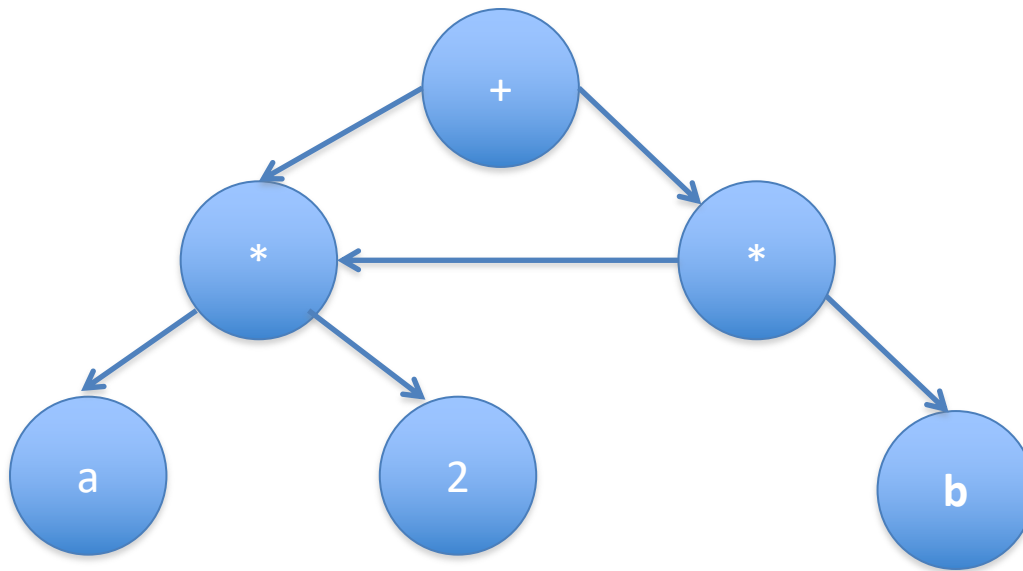
Example

- Abstract syntax for $x = 2^*(n+m)$



DAGs (Directed Acyclic Graphs)

- Variation on ASTs with **shared substructures**
- **Pro:** saves space, exposes redundant sub-expressions
- **Con:** less flexibility if part needs to be changed



Linear IRs

- Pseudo-code for some abstract machine
- Level of abstraction varies
- Simple, compact data structures
 - Commonly used: arrays, linked lists
- Examples:
 - 3-address code,
 - Stack machine code

```
t1 ← 2
t2 ← b
t3 ← t1 *
t2
t4 ← a
t5 ← t4 -
t3
```

- Fairly compact
- Compiler can control reuse of names – clever choice can reveal optimizations
- ILOC & similar code

```
push 2
push b
multiply
push a
subtract
```

- Each instruction consumes top of stack & pushes result
- Very compact
- Easy to create and interpret
- Java bytecode, MSIL

Abstraction Levels in Linear IR

- Linear IRs can be:
 - High-level abstraction (close to the source language)
 - Medium-level abstraction
 - Very low-level abstraction
- Examples: Linear IRs for C array reference $a[i][j+2]$
 - High-level: $t1 \leftarrow a[i,j+2]$

More IRs for $a[i][j+2]$

- Medium-level

$t1 \leftarrow j + 2$

$t2 \leftarrow i * 20$

$t3 \leftarrow t1 + t2$

$t4 \leftarrow 4 * t3$

$t5 \leftarrow \text{addr } a$

$t6 \leftarrow t5 + t4$

$t7 \leftarrow *t6$

- Low-level

$r1 \leftarrow [\text{fp}-4]$

$r2 \leftarrow r1 + 2$

$r3 \leftarrow [\text{fp}-8]$

$r4 \leftarrow r3 * 20$

$r5 \leftarrow r4 + r2$

$r6 \leftarrow 4 * r5$

$r7 \leftarrow \text{fp} - 216$

$f1 \leftarrow [r7+r6]$

Abstraction Level Tradeoffs

- High-level abstraction:
 - Good for some source-level optimizations and semantic checking
 - Can't optimize things that are hidden – like address arithmetic for array subscripting
- Medium-level abstraction:
 - More details, but keeps more higher-level semantic information
 - great for machine-independent optimizations
 - Many (all?) optimizing compilers work at this level
- Low-level abstraction:
 - Need for good code generation and resource utilization in back end
 - Loses semantic knowledge (e.g., variables, data aggregates, source relationships are usually missing)
- Many compilers use all 3 in different phases

Three-Address Code (TAC)

- Usual form: $x \leftarrow y \text{ op } z$
 - One operator
 - Maximum of 3 names
 - (Copes with: nullary $x \leftarrow y$ and unary $x \leftarrow \text{op } y$)
- Eg: $x = 2 * (m + n)$ becomes
 - $t1 \leftarrow m + n; \quad t2 \leftarrow 2 * t1; \quad x \leftarrow t2$
 - You may prefer: `add t1, m, n; mul t2, 2, t1; mov x, t2`
- “Expression temps”:
 - Invent as many new temp names as needed
 - Don’t correspond to any user variables; de-anonymize expressions
- Store in a quad(ruple)
 - $\langle \text{lhs}, \text{rhs1}, \text{op}, \text{rhs2} \rangle$

Three Address Code

- Advantages
 - Resembles code for actual machines
 - Explicitly names intermediate results
 - Compact
 - Often easy to rearrange
- Various representations
 - Quadruples, triples, SSA (Static Single Assignment)
 - We will see much more of this...

Stack Machine Code Example

Hypothetical code for $x = 2 * (m + n)$

```
pushaddr  x
pushconst 2
pushval   n
pushval   m
add
mult
store
```

m
n
2
@x
?

m + n
2
@x
?

$2*(m+n)$
@x
?

?

Compact: common opcodes just 1 byte wide; instructions have 0 or 1 operand

Stack Machine Code

- Originally used for stack-based computers (famous example: B5000, ~1961)
- Now also used for virtual machines:
 - UCSD Pascal – pcode
 - Forth
 - Java bytecode in a .class files (generated by Java compiler)
 - MSIL in a .dll or .exe assembly (generated by C#/F#/VB compiler)
- Advantages
 - Compact; mostly 0-address opcodes (fast download over network)
 - Easy to generate; easy to write a FrontEnd compiler, leaving the 'heavy lifting' and optimizations to the JIT
 - Simple to interpret or compile to machine code
- Disadvantages
 - Inconvenient/difficult to optimize directly
 - Does not match up with modern chip architectures

Hybrid IRs

- Combination of structural and linear
- Level of abstraction varies
- Most common example: control-flow graph (CFG)

Control Flow Graph (CFG)

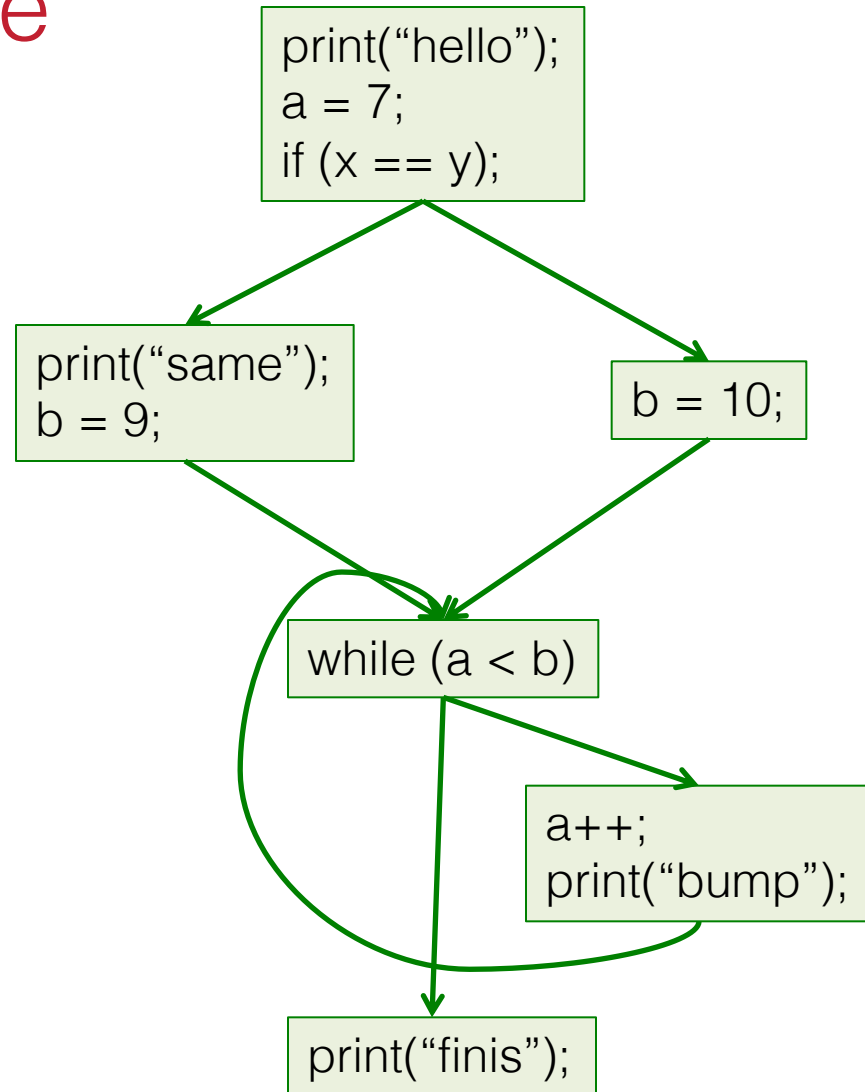
- **Nodes:** basic blocks
- **Edges:** possible flow of control from one block to another, (i.e., possible execution orderings)
 - Edge from A to B if B could execute immediately after A in some possible execution
- Required for much of the analysis done during optimization phases

Basic Blocks

- Fundamental concept in analysis/optimization
- A *basic block* is:
 - A sequence of code
 - One entry, one exit
 - Always executes as a single unit (“straight-line code”) – so it can be treated as an indivisible block
 - We’ll ignore exceptions, at least for now
- Usually represented as some sort of a list although Trees/DAGs are possible

CFG Example

```
print("hello");  
a=7;  
if (x == y) {  
    print("same");  
    b = 9;  
} else {  
    b = 10;  
}  
while (a < b) {  
    a++;  
    print("bump");  
}  
print("finis");
```



Basic Blocks: Start with Tuples

```
1 i = 1
2 j = 1
3 t1 = 10 * i
4 t2 = t1 + j
5 t3 = 8 * t2
6 t4 = t3 - 88
7 a[t4] = 0
8 j = j + 1
9 if j <= 10 goto #3

10 i = i + 1
11 if i <= 10 goto #2
12 i = 1
13 t5 = i - 1
14 t6 = 88 * t5
15 a[t6] = 1
16 i = i + 1
17 if i <= 10 goto #13
```

Typical "tuple stew" - IR generated by traversing an AST

Partition into **Basic Blocks**:

- Sequence of consecutive instructions
- No jumps into the middle of a BB
- No jumps out of the middles of a BB
- "I've started, so I'll finish"
- (Ignore exceptions)

Basic Blocks: Leaders

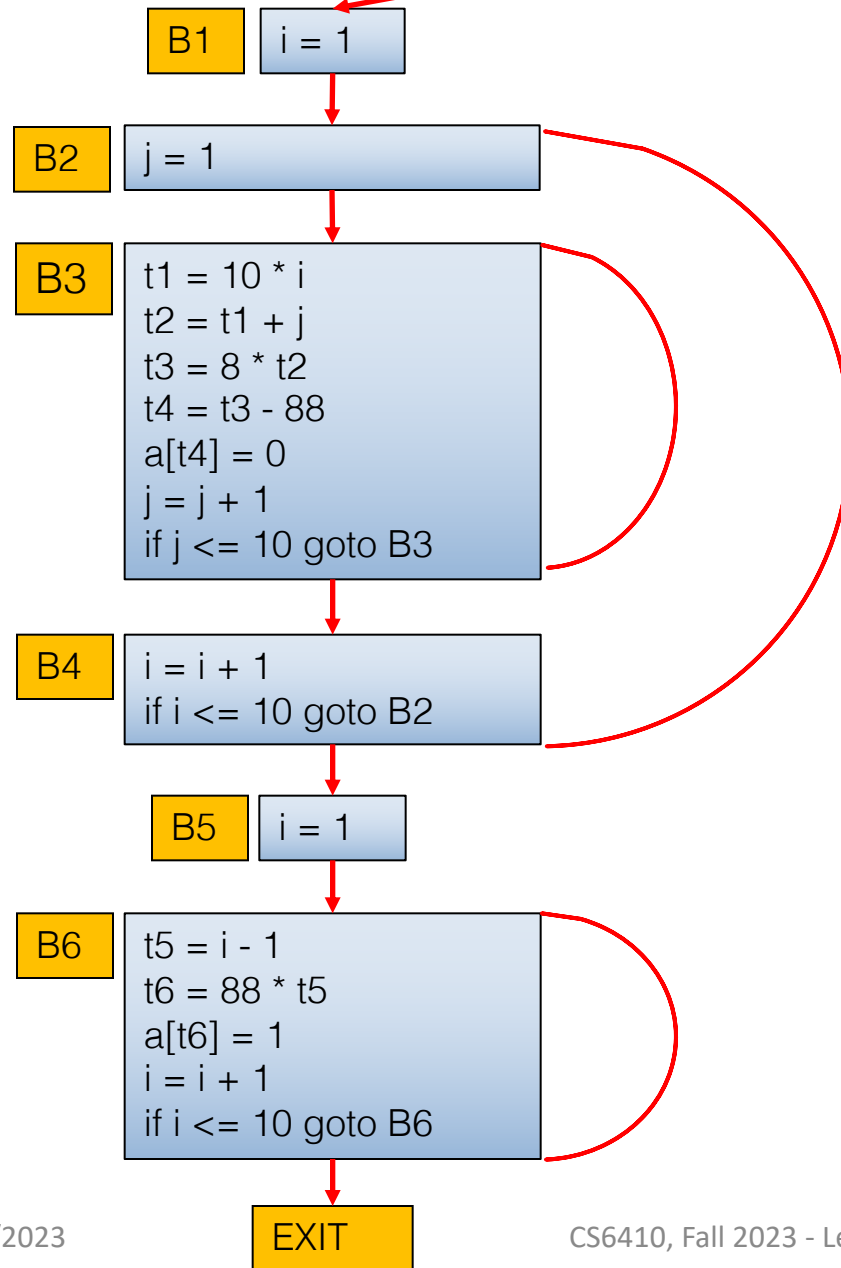
1 i = 1	10 i = i + 1
2 j = 1	11 if i <= 10 goto #2
3 t1 = 10 * i	12 i = 1
4 t2 = t1 + j	13 t5 = i - 1
5 t3 = 8 * t2	14 t6 = 88 * t5
6 t4 = t3 - 88	15 a[t6] = 1
7 a[t4] = 0	16 i = i + 1
8 j = j + 1	17 if i <= 10 goto #13
9 if j <= 10 goto #3	

Identify Leaders (first instruction in a basic block):

- First instruction is a leader
- Any target of a branch/jump/goto
- Any instruction immediately after a branch/jump/goto

Leaders in **red**. Why is each leader a leader?

Basic Blocks: Flowgraph



Control Flow Graph ("CFG", again!)

- 3 loops total
- 2 of the loops are nested

Most of the executions likely spent in loop bodies; that's where to focus efforts at optimization

Identifying Basic Blocks: Recap

- Perform linear scan of instruction stream
- A basic blocks begins at each instruction that is:
 - The beginning of a method
 - The target of a branch
 - Immediately follows a branch or return

Dependency Graphs

- Often used in conjunction with another IR
- **Data dependency**: edges between nodes that reference common data
- **Examples**
 - **RAW – read after write**: Block A defines x then B reads it
 - **WAR – “anti-dependence”**: Block A reads x then B writes it
 - **WAW**: Blocks A and B both write x – order of blocks must reflect original program semantics
- These restrict reorderings the compiler can do

What IR to Use?

- Common choice: all(!)
 - AST used in early stages of the compiler
 - Closer to source code
 - Good for semantic analysis
 - Facilitates some higher-level optimizations
 - Lower to linear IR for optimization and codegen
 - Closer to machine code
 - Use to build control-flow graph
 - Exposes machine-related optimizations
 - Hybrid (graph + linear IR = CFG) for dataflow & opt

Coming Attractions

- To get a running compiler we need:
 - Execution model for language constructs
 - x86-64 assembly language for compiler writers
 - Code generation and runtime bootstrap details
- We'll also spend considerable time on compiler optimization
 - Intermediate reps., graphs, SSA, dataflow
 - Optimization analysis and transformations
- Immediate problem is to keep lectures from getting too far ahead of the project - maybe hold off on runtime details?
 - Thoughts? Suggestions? Opinions?



[Meme credit: imgflip.com]