

# CS 6410: Compilers

Fall 2023

Tamara Bonaci  
[t.bonaci@northeastern.edu](mailto:t.bonaci@northeastern.edu)

Thank you to UW faculty Hal Perkins. Today lecture notes are a modified version of his lecture notes.

## Credits For Course Material

- Big thank you to UW CSE faculty members, Hal Perkins and Justin Hsia
- Some direct ancestors of this course:
  - UW CSE 401 (Chambers, Snyder, Notkin, Perkins, Ringenburt, Henry, ...)
  - UW CSE PMP 582/501 (Perkins)
  - Cornell CS 412-3 (Teitelbaum, Perkins)
  - Rice CS 412 (Cooper, Kennedy, Torczon)
  - Many books (Appel; Cooper/Torczon; Aho, [[Lam,] Sethi,] Ullman [Dragon Book], Fischer, [Cytron ,] LeBlanc; Muchnick, ...)

# Agenda

- Overview of x86-64 architecture (core part only)
- Mapping source code to x86-64<sup>1</sup>
  - Basic statements and expressions
  - Object representation and layout
  - Field access
  - What is `this`?
  - Object creation – `new`
  - Method calls
    - Dynamic dispatch
    - Method tables
    - `super`
  - Runtime type information

<sup>1</sup> Mapping for other common architectures is similar

Reading:

Cooper and Torczon, chapters 4.1-4.4, 5.5, 6.2-6.5 and 7.1-7.4

Dragon book, chapters 6.3-6.5, 7.1-7.7, 8.1-8.4

# Review: x86-64 Assembly “Data Types”

- Integral data of 1, 2, 4, or 8 bytes
  - Data values
  - Addresses
- Floating point data of 4, 8, 10 or 2x8 or 4x4 or 8x2
  - Different registers for those (e.g. xmm1, ymm2)
  - Come from extensions to x86 (SSE, AVX, ...)
- No aggregate types such as arrays or structures
  - Just contiguously allocated bytes in memory
- Two common syntaxes
  - “AT&T”: used by our course, slides, textbook, gnu tools, ...
  - “Intel”: used by Intel documentation, Intel tools, ...
  - Must know which you’re reading

# Review: Intel vs. GNU Assembler

- Main differences between Intel docs and gcc assembler

	Intel/Microsoft	AT&T/GNU as
Operand order: op a,b	a = a op b (dst first)	b = b op a (dst last)
Memory address	[baseregister+offset]	offset(baseregister)
Instruction mnemonics	mov, add, push, ...	movq, addq, pushq [operand size is added to end]
Register names	rax, rbx, rbp, rsp, ...	%rax, %rbx, %rbp, %rsp, ...
Constants	17, 42	\$17, \$42
Comments	; to end of line	# to end of line or /* ... */

- Intel docs also include many complex, historical instructions and artifacts not commonly used by modern compilers – we won't use them either

# Review: x86-64 Memory Model

- 8-bit bytes, byte addressable
- 16-, 32-, 64-bit words, double words and quad words (Intel terminology)
  - That's why the 'q' in 64-bit instructions like `movq`, `addq`, etc.
- Data should normally be aligned on “natural” boundaries for performance, although unaligned accesses are generally supported – but with a big performance penalty on many machines
- Little-endian – address of a multi-byte integer is address of low-order byte

## Review: x86-64 registers

- 16 64-bit general registers
  - %rax, %rbx, %rcx, %rdx, %rsi, %rdi, %rbp, %rsp, %r8-%r15
- Registers can be used as 64-bit integers or pointers, or 32-bit ints
  - Also possible to reference low-order 16- and 8-bit chunks – we won't for the most part
- To simplify our project we'll use only 64-bit data (ints, pointers, even booleans!)

# Review: Instruction Format

- Typical data manipulation instruction

label: opcode src,dst # comment

- Meaning is  
dst  $\leftarrow$  dst op src
- Normally, one operand is a register, the other is a register, memory location, or integer constant
  - Can't have both operands in memory – can't encode two memory addresses in a single instruction (e.g., cmp, mov)
- Language is free-form, comments and labels may appear on lines by themselves (and can have multiple labels per line of code)



# Review: Operand types

- **Immediate:** Constant integer data
  - Examples: `0x400`, `-533`
  - Encoded with 1, 2, 4, or 8 bytes depending on the instruction
- **Register:** 1 of 16 integer registers
  - Examples: `rax`, `r13`
  - But `rsp`, `rbp` reserved for special use
  - Others have special uses for particular instructions
- **Memory:** Consecutive bytes of memory at a computed address
  - Simplest example: `DWORD PTR [rax]`
  - Various other “address modes”

rax
rcx
rdx
rbx
rsi
rdi
rsp
rbp
rN

# Review: Three Basic Kinds of Instructions

## 1) Transfer data between memory and register

- **Load** data from memory into register
  - $\text{reg} = \text{Mem}[\text{address}]$
- **Store** register data into memory
  - $\text{Mem}[\text{address}] = \text{reg}$

Remember:  
Memory is indexed  
just like an array of  
bytes!

## 2) Perform arithmetic operation on register or memory data

- $c = a + b; \quad z = x \ll y; \quad i = h \& g;$

## 3) Control flow: what instruction to execute next

- Unconditional jumps to/from procedures
- Conditional branches

# Review: x86-64 - Instructions Overview

- Data transfer instruction (`mov`)
- Arithmetic operations
- Memory addressing modes
- Address computation instruction (`lea`)

# Review: x86-64 Memory Stack

- Register `%rsp` points to the “top” of stack
  - Dedicated for this use; don’t use otherwise
  - Points to the last 64-bit quadword pushed onto the stack (not next “free” quadword)
  - Should always be quadword (8-byte) aligned
    - It will start out this way, and will stay aligned unless your code does something bad
    - Software generally requires 16-byte alignment when function is called
  - Stack grows down

# Review: Stack Instructions

## pushq src

$\%rsp \leftarrow \%rsp - 8;$

$\text{memory}[\%rsp] \leftarrow \text{src}$

(e.g., push src onto the stack)

## popq dst

$\text{dst} \leftarrow \text{memory}[\%rsp];$

$\%rsp \leftarrow \%rsp + 8$

(e.g., pop top of stack into dst and logically remove it from the stack)

## Review: Stack Frames

- When a method is called, a **stack frame** is traditionally allocated on logical “top” of the stack to hold its local variables
- Frame is popped on method return
- By convention, %rbp (base pointer) points to a known offset into the stack frame
  - Local variables referenced relative to %rbp
  - Base pointer common in 32-bit x86 code; less so in x86-64 code where push/pop used less & stack frame has a fixed size so locals can be referenced from %rsp easily
  - We will use %rbp in our project – simplifies addressing of local variables and compiler bookkeeping

## Review: Operand Address Modes (1)

- These should cover most of what we'll need

```
movq    $17,%rax           # store 17 in %rax
movq    %rcx,%rax          # copy %rcx to %rax
movq    -16(%rbp),%rax      # copy memory to %rax
movq    %rax,-24(%rbp)      # copy %rax to memory
```

- References to object fields work similarly – put the object's memory address in a register and use that address plus an offset
- Remember: can't have two memory addresses in a single instruction

## Review:Operand Address Modes (2)

- A memory address can combine the contents of two registers (with one optionally multiplied by 2, 4, or 8) plus a constant:
- $$\text{basereg} + \text{indexreg} * \text{scale} + \text{constant}$$
- Main use of general form is for array subscripting or small computations - if the compiler is clever
  - **Example:** suppose we have an array of 8-byte ints with address of the array A in %rcx and subscript i in %rax. Code to store %rbx in A[i]

```
movq    %rbx, (%rcx,%rax,8)
```



## qword ptr – Intel assembler

- Obscure, but sometimes necessary...
- If the assembler can't figure out the size of the operands to move, you can explicitly tell it to move 64bits with the qualifier "qword ptr"

```
mov    qword ptr [rax+16], [rbp-8]
```

- Similarly, for dword ptr, etc
- Use this if the assembler complains; otherwise ignore
- Not an issue in GNU assembler – operand sizes encoded in opcode mnemonics

# Basic Data Movement and Arithmetic Instructions

`movq src, dst`

$\text{dst} \leftarrow \text{src}$

`addq src, dst`

$\text{dst} \leftarrow \text{dst} + \text{src}$

`subq src, dst`

$\text{dst} \leftarrow \text{dst} - \text{src}$

`incq dst`

$\text{dst} \leftarrow \text{dst} + 1$

`decq dst`

$\text{dst} \leftarrow \text{dst} - 1$

`negq dst`

$\text{dst} \leftarrow -\text{dst}$

(2's complement  
arithmetic negation)

# Integer Multiply and Divide

`imulq src, dst`

$\text{dst} \leftarrow \text{dst} * \text{src}$

dst must be a register

`cqto`

$\%rdx:\%rax \leftarrow$  128-bit  
sign extended copy of  
 $\%rax$

(why??? To prep  
numerator for `idivq`!)

`idivq src`

Divide  $\%rdx:\%rax$  by `src`  
( $\%rdx:\%rax$  holds sign-  
extended 128-bit value;  
cannot use other  
registers for division)

$\%rax \leftarrow$  quotient

$\%rdx \leftarrow$  remainder

(no division in MiniJava!)

# Bitwise Operations

`andq src, dst`

$\text{dst} \leftarrow \text{dst} \& \text{src}$

`orq src, dst`

$\text{dst} \leftarrow \text{dst} | \text{src}$

`xorq src, dst`

$\text{dst} \leftarrow \text{dst} \wedge \text{src}$

`notq dst`

$\text{dst} \leftarrow \sim \text{dst}$

(logical or 1's  
complement)

## Shifts and Rotates

`shlq count, dst`  
dst shifted left count bits

`shrq count, dst`  
dst  $\leftarrow$  dst shifted right  
count bits (0 fill)

`sarq count, dst`  
dst  $\leftarrow$  dst shifted right  
count bits (sign bit fill)

`rolq count, dst`  
dst  $\leftarrow$  dst rotated left  
count bits

`rorq count, dst`  
dst  $\leftarrow$  dst rotated right  
count bits

## Uses for Shifts and Rotates

- Can often be used to optimize multiplication and division by small constants (mul/div by powers of 2)
  - If you're interested, look at “Hacker's Delight” by Henry Warren, A-W, 2<sup>nd</sup> ed, 2012
    - Lots of very cool bit fiddling and other algorithms
  - But be careful – be sure semantics are OK
    - Example: right shift is not the same as integer divide for negative numbers – shift truncates towards  $-\infty$
- There are additional instructions that shift and rotate double words, use a calculated shift amount instead of a constant, etc.

# Address Computation Instruction

- `lea dst, src`
  - `lea` stands for **load effective address**
  - `src` is address expression (in brackets)
  - `dst` is a register
  - Sets `dst` to the address computed by the `src` expression (**does not go to memory! – it just does math**)
  - Example: `lea rax, QWORD PTR[rdx+rcx*4+1]`
- Uses:
  - Computing addresses without a memory reference
    - e.g. translation of `p = &x[i];`
  - Computing arithmetic expressions of the form `x+k*i+d`
    - Though `k` can only be 1, 2, 4, or 8

## Load Effective Address

- The unary & operator in C/C++

`leaq src, dst`      #  $\text{dst} \leftarrow \text{address of src}$

- `dst` must be a register
- Address of `src` includes any address arithmetic or indexing
- Useful to capture addresses for pointers, reference parameters, etc.
- Also useful for computing arithmetic expressions that match  $r1 + \text{scale} * r2 + \text{const}$



# Control Flow - GOTO

- At this level, all we have is `goto` and conditional `goto`
- Loops and conditional statements are synthesized from these
- **Note:** random jumps play havoc with pipeline efficiency; much work is done in modern compilers and processors to minimize this impact

# Unconditional Jumps

```
jmp dst
```

```
%rip ← address of dst
```

- dst is usually a label in the code (which can be on a line by itself)
- dst address can also be indirect using the address in a register or memory location ( \*reg or \*(reg) ) – use for method calls, switch

# Conditional Jumps

- Most arithmetic instructions set “**condition code**” bits to record information about the result (zero, non-zero, >0, etc.)
  - True of addq, subq, andq, orq; but not imulq, idivq, leaq
- Other instructions that set condition codes

```
cmpq src,dst # compare dst to src (e.g., dst-src)
testq src,dst      # calculate dst & src (logical and)
```

  - These do not alter src or dst

## Conditional Jumps Following Arithmetic Operations

```
jz    label    # jump if result == 0
jnz   label    # jump if result != 0
jg    label    # jump if result > 0
jng   label    # jump if result <= 0
jge   label    # jump if result >= 0
jnge  label    # jump if result < 0
jl    label    # jump if result < 0
jnl   label    # jump if result >= 0
jle   label    # jump if result <= 0
jnle  label    # jump if result > 0
```

- Obviously, the assembler is providing multiple opcode mnemonics for several actual instructions

# Compare and Jump Conditionally

- **Want:** compare two operands and jump if a relationship holds between them
- Would like to do this

`jmpcond op1, op2, label`

but can't, because 3-operand instructions can't be encoded in x86-64

(also true of most other machines for that matter)

## cmp and jcc

- Instead, we use a 2-instruction sequence

```
cmpq op1, op2
```

```
jcc label
```

where  $j_{cc}$  is a conditional jump that is taken if the result of the comparison matches the condition  $cc$

## Conditional Jumps Following Arithmetic Operations

<code>je</code>	<code>label</code>	<code># jump if op1 == op2</code>
<code>jne</code>	<code>label</code>	<code># jump if op1 != op2</code>
<code>jg</code>	<code>label</code>	<code># jump if op1 &gt; op2</code>
<code>jng</code>	<code>label</code>	<code># jump if op1 &lt;= op2</code>
<code>jge</code>	<code>label</code>	<code># jump if op1 &gt;= op2</code>
<code>jnge</code>	<code>label</code>	<code># jump if op1 &lt; op2</code>
<code>jl</code>	<code>label</code>	<code># jump if op1 &lt; op2</code>
<code>jnl</code>	<code>label</code>	<code># jump if op1 &gt;= op2</code>
<code>jle</code>	<code>label</code>	<code># jump if op1 &lt;= op2</code>
<code>jnle</code>	<code>label</code>	<code># jump if op1 &gt; op2</code>

- Again, the assembler is mapping more than one mnemonic to some machine instructions

# Function Call and Return

- The x86-64 instruction set itself only provides for transfer of control (jump) and return
- Stack is used to capture return address and recover it
- Everything else – parameter passing, stack frame organization, register usage – is a matter of convention and not defined by the hardware



# call and ret Instructions

## call label

- Push address of next instruction and jump

$\%rsp \leftarrow \%rsp - 8;$

$\text{memory}[\%rsp] \leftarrow \%rip$

$\%rip \leftarrow \text{address of label}$

- Call address can be in a register or memory as with jumps

## ret

- Pop address from top of stack and jump

$\%rip \leftarrow \text{memory}[\%rsp];$

$\%rsp \leftarrow \%rsp + 8$

- **WARNING!** The word on the top of the stack had better be an address and not some leftover data

## enter and leave

- Complex instructions for languages with nested procedures
  - `enter` can be slow on current processors – best avoided – i.e., don't use it in your project
  - `leave` is equivalent to

```
mov %rsp, %rbp
pop %rbp
```

and is generated by many compilers. Fits in 1 byte, saves space. Not clear if it's any faster.

# X86-64-Register Usage

- `%rax` – function result
- Arguments 1-6 passed in these registers in order
  - `%rdi, %rsi, %rdx, %rcx, %r8, %r9`
  - For Java/C++ “this” pointer is first argument, in `%rdi`
    - More about “this” later
- `%rsp` – stack pointer; value must be 8-byte aligned always and 16-byte aligned when calling a function
- `%rbp` – frame pointer (optional use)
  - We’ll use it

# x86-64 Register Save Conventions

- A called function must preserve these registers (or save/restore them if it wants to use them)
  - `%rbx`, `%rbp`, `%r12–%r15`
- `%rsp` isn't on the “callee save list”, but needs to be properly restored for return
- All other registers can change across a function call

## The Nice Thing About Standards...

- The above is the System V/AMD64 ABI convention (used by Linux, OS X)
- Microsoft's x64 calling conventions are slightly different (sigh...)
  - First four parameters in registers `%rcx`, `%rdx`, `%r8`, `%r9`; rest on the stack
  - Stack frame must include empty space for called function to use to store values passed in parameter registers if desired

## x86-64 Function Call

- Caller places up to 6 arguments in registers, rest on stack, then executes call instruction (which pushes 8-byte return address)
- On entry, called function prologue sets up the stack frame:

```
pushq    %rbp           # save old frame ptr
movq     %rsp, %rbp      # new frame ptr is top of
                        # stack after ret addr and
                        # old rbp pushed
subq     $framesize, %rsp # allocate stack frame
```

## x86-64 Function Return

- Called function:
  1. Puts result in %rax (if any)
  2. Restores any callee-saved registers (if needed)
- Called function returns with:

```
    movq  %rbp,%rsp    # or use leave instead
    popq  %rbp          # of movq/popq
    ret
```
- If caller allocated space for arguments it deallocates as needed

## Caller Example

- $n = \text{sumOf}(17, 42)$

```
    movq    $42, %rsi    # load arguments
    movq    $17, %rdi
    call    sumOf        # jump & push ret
addr
    movq    %rax, offsetn(%rbp) # store result
```



# Example Function

- Source code

```
int sumOf(int x, int y) {  
    int a, int b;  
    a = x;  
    b = a + y;  
    return b;  
}
```

# Assembly Language Version

```
int sumOf(int x, int y) {  
    int a,  
    int b;  
    # sumOf:  
    pushq %rbp #prologue  
    movq %rsp, %rbp  
    subq $16, %rsp  
  
    a = x;  
    # a = x;  
    movq %rdi, -8(%rbp)  
  
    b = a + y;  
    # b = a + y;  
    movq -8(%rbp), %rax  
    addq %rsi, %rax  
    movq %rax, -16(%rbp)  
  
    return b;  
    # return b;  
    movq -16(%rbp), %rax  
    movq %rbp, %rsp  
    popq %rbp  
    ret  
    # }
```

# Stack for sumOf

```
int sumOf(int x, int y) {  
    int a, int b;  
    a = x;  
    b = a + y;  
    return b;  
}
```

# Code Shape

## Review: Variables

- For us, all data is either:
  - In a stack frame (method local variables)
  - In an object (instance variables)
- **Local variables** accessed via `%rbp`  

```
movq    -16(%rbp), %rax
```
- **Object instance variables** accessed via an offset from an object address in a register  
(details later today)

# Conventions for Examples

- Examples show code snippets in isolation
  - Much the way we'll generate code for different parts of the AST in a compiler visitor pass
- Register %rax used here as a generic example
  - Rename as needed for more complex code
- 64-bit data used everywhere
- Examples show a few peephole optimizations
  - Some might be easy to do in the compiler project

## What We're Skipping for Now

- Real code generator deals with many things like:
  - Which registers are busy at which point in the program
  - Which registers to spill into memory when a new register is needed, and no free ones are available
  - Different sizes of data
  - Exploiting the full instruction set

# Code Generation for Constants

- Source

17

- x86-64

`movq $17, %rax`

– Idea: realize constant value in a register

- Optimization: if constant is 0

`xorq %rax, %rax`

(but some processors do better with `movq $0, %rax` –  
and this has changed over time, too)



# Assignment Statement

- Source

`var = exp;`

- x86-64

`<code to evaluate exp into, say, %rax>`

`movq %rax, offsetvar(%rbp)`

# Unary Minus

- Source

-exp

- x86-64

<code evaluating exp into %rax>  
negq %rax

- Optimization

– Collapse  $-(-\text{exp})$  to  $\text{exp}$

- Unary plus is a no-op

## Binary +

- Source

$\text{exp}_1 + \text{exp}_2$

- x86-64

<code evaluating  $\text{exp}_1$  into  $\%rax$ >

<code evaluating  $\text{exp}_2$  into  $\%rdx$ >

`addq  $\%rdx, \%rax$`

## Binary +

- Some optimizations

- If  $\text{exp}_2$  is a simple variable or constant, don't need to load it into another register first. Instead:

`addq  $\text{exp}_2$ , %rax`

- Change  $\text{exp}_1 + (-\text{exp}_2)$  into  $\text{exp}_1 - \text{exp}_2$

- If  $\text{exp}_2$  is 1

`incq %rax`

- **Somewhat surprising:** whether this is better than `addq $1, %rax` depends on processor implementation and has changed over time

## Binary -, \*

- Same as +
  - Use `subq` for `-` (but not commutative!)
  - Use `imulq` for `*`
- Some optimizations
  - Use left shift to multiply by powers of 2
  - If your multiplier is slow or you've got free scalar units and multiplier is busy, you can do  $10^*x = (8^*x) + (2^*x)$ 
    - But might be slower depending on microarchitecture
  - Use `x+x` instead of `2*x`, etc. (often faster)
  - Can use:  

```
leaq (%rax,%rax,4),%rax
```

 to compute `5*x`, then  

```
addq %rax,%rax
```

 to get `10*x`, etc. etc.
  - Use `decq` for `x-1`

# Signed Integer Division

- Ghastly on x86-64

- Only works on 128-bit int divided by 64-bit int
  - (similar instructions for 64-bit divided by 32-bit in 32-bit x86)
- Requires use of specific registers
- Very slow (~50 clocks)

- Source

$\text{exp}_1 / \text{exp}_2$

- x86-64

<code evaluating  $\text{exp}_1$  into %rax **ONLY**>

<code evaluating  $\text{exp}_2$  into %ebx>

cqto                      # extend to %rdx:%rax, clobbers  
%rdx

idivq %ebx    # quotient in %rax, remainder in %rdx

# Control Flow

# Control Flow

- **Basic idea:** decompose higher level operation into conditional and unconditional gotos
- In the following,  $j_{\text{false}}$  is used to mean jump when a condition is false
  - **No such instruction on x86-64**
  - Will have to realize it with appropriate instruction to set condition codes followed by conditional jump
  - Normally don't need to actually generate the value "true" or "false" in a register
    - **But this is a useful shortcut hack for the project**



# While

- Source

while (cond) stmt

- x86-64

```
test:  <code evaluating cond>
```

```
    jfalse done
```

```
    <code for stmt>
```

```
    jmp test
```

```
done:
```

- **Note:** In generated asm code we need to have unique labels for each loop, conditional statement, etc.

## Optimization for While

- Put the test at the end:

```
    jmp  test
loop:<code for stmt>
test:<code evaluating cond>
    j_true loop
```

- Why bother?
  - Pulls one `jmp` instruction out of the loop
  - May avoid a pipeline stall on `jmp` on each iteration
    - Although modern processors will often predict control flow and avoid the stall – x86-64 does this particularly well
- Easy to do from AST or other IR; not so easy if generating code on the fly (e.g., recursive descent 1-pass compiler)

# Do-While

- Source

```
do stmt while(cond)
```

- x86-64

```
loop:  <code for stmt>  
      <code evaluating cond>  
      jtrue loop
```

# If

- Source

if (cond) stmt

- x86-64

<code evaluating cond>

j<sub>false</sub> skip

<code for stmt>

skip:

# If-Else

- Source

if (cond) stmt<sub>1</sub> else stmt<sub>2</sub>

- x86-64

<code evaluating cond>

j<sub>false</sub> else

<code for stmt<sub>1</sub>>

jmp done

else: <code for stmt<sub>2</sub>>

done:

# Jump Chaining

- **Observation:** naïve implementation can produce jumps to jumps (if-else if-...-else; or nested loops or conditionals, ...)
- **Optimization:** if a jump has as its target an unconditional jump, change the target of the first jump to the target of the second
  - Repeat until no further changes
  - Often done in peephole optimization pass after initial code generation

# Boolean Expressions

- What do we do with this?

$x > y$

- Expression that evaluates to true or false
  - Could generate the value (0/1 or whatever the local convention is)
  - But normally we don't want/need the value – we're only trying to decide whether to jump

## Code for $\text{exp1} > \text{exp2}$

- **Basic idea:** generated code depends on context:
  - What is the jump target?
  - Jump if the condition is true or if false?
- **Example:** evaluate  $\text{exp1} > \text{exp2}$ , jump on false, target if jump taken is L123

```
<evaluate exp1 to %rax>
<evaluate exp2 to %rdx>
cmpq  %rdx,%rax
jng   L123
```



# Boolean Operators

# Boolean Operators: !

- Source  
! exp
- Context: evaluate exp and jump to L123 if false (or true)
- To compile !, just reverse the sense of the test: evaluate exp and jump to L123 if true (or false)

## Boolean Operators: && and ||

- In C/C++/Java/C#/many others, these are **short-circuit operators**
  - Right operand is evaluated only if needed
- Basically, generate the if statements that jump appropriately and only evaluate operands when needed

## Example: Code for &&

- Source

if ( $\text{exp}_1$  &&  $\text{exp}_2$ ) stmt

- x86-64

<code for  $\text{exp}_1$ >

j<sub>false</sub> skip

<code for  $\text{exp}_2$ >

j<sub>false</sub> skip

<code for stmt>

skip:

## Example: Code for ||

- Source

if ( $\text{exp}_1 \parallel \text{exp}_2$ ) stmt

- x86-64

```
<code for  $\text{exp}_1$ >
j_true doit
<code for  $\text{exp}_2$ >
j_false skip
doit: <code for stmt>
skip:
```

## Realizing Boolean Values

- If a boolean value needs to be stored in a variable or method call parameter, generate code needed to actually produce it
- Typical representations: 0 for false, +1 or -1 for true
  - C specifies 0 and 1 if stored; we'll use that
  - Best choice can depend on machine instructions; normally some convention is established during the primeval history of the architecture

# Boolean Values: Example

- Source

var = bexp;

- x86-64

```
<code for bexp>
j_false      genFalse
movq $1,%rax
jmp  storeIt
genFalse:
movq $0,%rax          # or xorq
storeIt:
movq %rax,offset_var(%rbp) # generated
by asg stmt
```

# Better, If Enough Registers

- Source

var = bexp;

- x86-64

```
xorq    %rax,%rax
```

```
<code for bexp>
```

```
j_false store
```

```
incq    %rax
```

```
store:
```

```
movq    %rax,offset_var(%rbp)
```

```
#generated by asg
```

- Better: use movecc instruction to avoid conditional jump
- Can also use conditional move instruction for sequences like  $x = y < z ? y : z$



## Better yet: setcc

- Source

var = x < y;

- x86-64

```
movq    offset_x(%rbp), %rax    # load x
cmpq    offset_y(%rbp), %rax    # compare to y
setl    %al                    # set low byte %rax
to 0/1
movzbq  %al, %rax              # zero-extend to 64
bits
movq    %rax, offset_var(%rbp)  # gen. by asg
stmt
```

## Other Control Flow: switch

- **Naïve:** generate a chain of nested if-else if statements
- **Better:** switch statement is intended to allow  $O(1)$  selection, provided the set of switch values is reasonably compact
- **Idea:** create a 1-D array of jumps or labels and use the switch expression to select the right one
  - Need to generate equivalent of an if to ensure expr. value is within bounds (& avoid wild jump/segfault)

# Switch

- Source

```
switch (exp) {  
    case 0: stmts0;  
    case 1: stmts1;  
    case 2: stmts2;  
}
```

“break” is an unconditional  
jump to the end of switch

- x86-64:

```
<put exp in %rax>  
"if (%rax < 0 ||  
%rax > 2)  
    jmp  
defaultLabel"  
movq  
    swtab(, %rax, 4), %rax  
    jmp *%rax  
    .data  
swtab:  
    .quad L0  
    .quad L1  
    .quad L2  
    .text  
L0: <stmts0>  
L1: <stmts1>  
L2: <stmts2>
```

# Arrays

# Arrays

- Several variations
- C/C++/Java
  - 0-origin: an array with  $n$  elements contains variables  $a[0] \dots a[n-1]$
  - 1 dimension (Java); 1 or more dimensions using row major order (C/C++)
- Key step is evaluate subscript expression, then calculate the location of the corresponding array element

# 0-Origin 1-D Integer Arrays

- Source

$\text{exp}_1[\text{exp}_2]$

- x86-64

`<evaluate exp1 (array address) in %rax>`

`<evaluate exp2 in %rdx>`

`address is (%rax,%rdx,8) # if 8 byte  
elements`

## 2-D Arrays

- Subscripts start with 0
- C/C++, etc. specify row-major order
  - E.g., an array with 3 rows and 2 columns is stored in sequence:  $a(0,0)$ ,  $a(0,1)$ ,  $a(1,0)$ ,  $a(1,1)$ ,  $a(2,0)$ ,  $a(2,1)$
- Fortran specifies column-major order
  - Exercises: What is the layout? How do you calculate location of  $a[i][j]$ ? What happens when you pass array references between Fortran and C/C++ code?
- Java does not have “real” 2-D arrays. A Java 2-D array is a pointer to a list of pointers to the rows
  - And rows may have different lengths (ragged arrays)

## $a[i][j]$ in C/C++/etc.

- If  $a$  is a “real” 0-origin, 2-D array, to find  $a[i][j]$ , we need to know:
  - Values of  $i$  and  $j$
  - How many columns (but not rows!) the array has
- Location of  $a[i][j]$  is:
  - Location of  $a$  +  $(i * (\text{\#of columns}) + j) * \text{sizeof}(\text{elt})$
- Can factor to pull out allocation-time constant part and evaluate that once – no recalculating at runtime; only calculate part depending on  $i, j$



# What does this program print?

```

class One {
    int tag;
    int it;

    void setTag() {
        tag = 1; }

    int getTag() {
        return tag; }

    void setIt(int
it) {
        this.it = it; }

    int getIt() {
        return it; }
}

class Two extends One {
    int it;
    void setTag() {
        tag = 2;
        it = 3;
    }

    int getThat() {
        return it;
    }

    void resetIt() {
        super.setit(42);
    }
}

public static void main(String[] args)
{
    Two two = new Two();
    One one = two;

    one.setTag();

    System.out.println(one.getTag());

    one.setIt(17);
    two.setTag();

    System.out.println(two.getIt());

    System.out.println(two.getThat());
    two.resetIt();

    System.out.println(two.getIt());

    System.out.println(two.getThat());

}

```

Your Answer Here

# Object Representation

- The naïve explanation is that an object contains
  - **Fields** declared in its class and in all superclasses
    - Redecclaration of a field hides (shadows) superclass instance – but the superclass field is still there
  - **Methods** declared in its class and all superclasses
    - Redecclaration of a method overrides (replaces) – but overridden methods can still be accessed by super...
- When a method is called, the method “inside” that particular object is called
  - Regardless of the static (compile-time) type of the variable
    - (But we really don’t want to copy all those methods, do we?)

# Actual Representation

- Each object contains:
  - An entry (“slot”) for each field (instance variable)
    - Including shadowed and private fields in superclasses
  - A pointer to a runtime data structure for its class
    - Key component: method dispatch table (next slide)
- Basically a C struct
- Fields hidden by declarations in subclasses are *still* allocated in the object and are accessible from superclass methods

# Method Dispatch Tables

- One of these per class, not per object
- Often called “vtable”, “vtbl”, or “vtab”
  - (virtual function table – term from C++)
- One pointer per method – points to beginning of method code
- Dispatch table offsets fixed at compile time in  $O(1)$  implementations

# Method Tables and Inheritance

- A really simple implementation
  - Method table for each class has pointers to all methods declared in it (a dictionary)
  - Method table also contains a pointer to parent class method table
  - Method dispatch:
    - Look in current table and use if method declared locally
    - Look in parent class table if not local
    - Repeat
    - “Message not understood” if you can’t find it after search
  - Actually used/needed in typical implementations of some dynamic languages (e.g. Ruby, Smalltalk, etc.)

## O(1) Method Dispatch

- Idea: first part of method table for extended class has pointers for the same methods in the same order as the parent class
  - BUT pointers actually refer to overriding methods if these exist
  - ∴ Method dispatch can be done with indirect jump using fixed offsets known at compile time –  $O(1)$ 
    - In C: `*(object->vtbl[offset])(parameters)`
- Pointers to methods added in subclass are after ptrs to inherited/overridden ones in vtable

## Method Dispatch Footnotes

- Don't need vtable pointer to parent class vtable for method calls, but still useful for other purposes
  - Casts and instanceof
- Multiple inheritance requires more complex mechanisms
  - Also true for multiple interfaces



# Confusing Example Revisited

```
class One {
    int tag;
    int it;
    void setTag() { tag = 1; }
    int getTag() { return tag; }
    void setIt(int it) {this.it = it;}
    int getIt() { return it; }
}
class Two extends One {
    int it;
    void setTag() {
        tag = 2; it = 3;
    }
    int getThat() { return it; }
    void resetIt() { super.setIt(42); }
}
```

```
public static void main(String[] args) {
    Two two = new Two();
    One one = two;

    one.setTag();
    System.out.println(one.getTag());

    one.setIt(17);
    two.setTag();
    System.out.println(two.getIt());
    System.out.println(two.getThat());
    two.resetIt();
    System.out.println(two.getIt());
    System.out.println(two.getThat());
}
```

## Now What?

- Need to explore
  - Object layout in memory
  - Compiling field references
    - Implicit and explicit use of “this”
  - Representation of vtables
  - Object creation – new
  - Code for dynamic dispatch
  - Runtime type information – instanceof and casts

# Object Layout

- Typically, allocate fields sequentially
- Follow processor/OS alignment conventions for struct/object when appropriate/available
  - Include padding bytes for alignment as needed
- Use first word of object for pointer to method table/class information
- Objects are allocated on the heap
  - No actual bits in the generated code

# Object Field Access

- Source

int n = obj.slot;

- x86-64

- Assuming that obj is a local variable in the current method's stack frame

```
movq    offset_obj(%rbp),%rax    # load obj ptr
movq    offset_slot(%rax),%rax   # load slot
movq    %rax,offset_n(%rbp)      # store n
```

- Same idea used to reference fields of “this”
  - Use implicit “this” parameter passed to method instead of a local variable to get object address

## Local Fields

- A method can refer to fields in the receiving object either explicitly as “this.f” or implicitly as “f”
  - Both compile to the same code – an implicit “this.” is assumed if not present explicitly

# Source Level View

What you write:

```
int getIt() {  
    return it;  
}  
void setIt(int it) {  
    this.it = it;  
}  
...  
obj.setIt(42);  
k = obj.getIt();
```

What you really get:

```
int getIt(ObjType this) {  
    return this.it;  
}  
void setIt(ObjType this, int it) {  
    this.it = it;  
}  
...  
setIt(obj, 42);  
k = getIt(obj);
```



## x86-64 “`this`” Convention (C++)

- “`this`” is an implicit first parameter to every non-static method
- Address of object placed in `%rdi` for every non-static method call
- Remaining parameters (if any) in `%rsi`, etc.
- We'll use this convention in our project

## MiniJava Method Tables (vtbls)

- Generate these as initialized data in the assembly language source program
- Need to pick a naming convention for assembly language labels; suggest:
  - For methods, classname\$methodname
    - Would need something more sophisticated for overloading
  - For the vtables themselves, classname\$\$
- First method table entry points to superclass table (we might not use this in our project, but is helpful if you add instanceof or type cast checks)



# Method Tables For Convoluted Example (gcc/as syntax)

```

class One {
    void setTag()      { ... }
    int getTag()       { ... }
    void setIt(int it) {...}
    int getIt()        { ... }
}

class Two extends One {
    void setTag() { ... }
    int getThat() { ... }
    void resetIt() { ... }
}

                                .data
                                One$$: .quad 0                #
                                no superclass
                                .quad One$setTag
                                .quad One$getTag
                                .quad One$setIt
                                .quad One$getIt
                                Two$$: .quad One$$            #
                                superclass
                                .quad Two$setTag
                                .quad One$getTag
                                .quad One$setIt
                                .quad One$getIt
                                .quad Two$getThat
                                .quad Two$resetIt
    
```

## Method Table Layout

Key point: first method entries in `Two`'s method table are pointers to methods declared in `One` in *exactly the same order*

- Actual pointers reference code appropriate for objects of each class (inherited or overridden)

∴ Compiler knows correct offset for a particular method pointer *regardless of whether that method is overridden* and regardless of the actual (dynamic) type of the object

# Object Creation – `new`

## Steps needed

- Call storage manager (malloc or similar) to get the raw bits
  - (and store 0's if required by the language, e.g., Java)
- Store pointer to method table in the first 8 bytes of the object
- Call a constructor with “`this`” pointer to the `new` object in `%rdi` and other parameters as needed
  - (Not in MiniJava since we don't have constructors)
- Result of `new` is a pointer to the `new` object

# Object Creation

- Source

```
One one = new One (...);
```

- x86-64

```
movq    $nBytesNeeded,%rdi    # obj size + 8 (include space for
vtbl ptr)                      #
call     mallocEquiv           # addr of allocated bits
returned in %rax
leaq     One$$,%rdx            # get method table address
movq     %rdx,0(%rax)          # store vtbl ptr at beginning
of object
movq     %rax,%rdi            # set up "this" for constructor
movq     %rax,offset_temp(%rbp) # save "this" for later
<load constructor arguments>   # arguments (if needed)
call     One$One              # call ctr if we have one (no
vtbl lookup)
movq     offset_temp(%rbp),%rax # recover ptr to object
movq     %rax,offset_one(%rbp)  # store object reference in
variable
```

# Constructor

- Why don't we need a vtable lookup to find the right constructor to call?
- Because at compile time we know the actual class (it says so right after “new”), so we can generate a call instruction to a known label
  - Same with `super.method(...)` or superclass constructor calls – at compile time we know all of the superclasses (need this to compile subclass and construct method tables), so we know statically what class “`super.method`” belongs to

# Method Calls

- Steps needed
  - Parameter passing: just like an ordinary C function, except load a pointer to the object in %rdi as the first (“this”) argument
  - Get a pointer to the object’s method table from the first 8 bytes of the object
  - Jump indirectly through the method table

# Method Call

- Source

obj.m(...);

- x86-64

<load arguments in registers as usual> # as needed

movq offset<sub>obj</sub>(%rbp),%rdi # first argument is obj ptr (“this”)

movq 0(%rdi),%rax # load vtable address into %rax

call \*offset<sub>m</sub>(%rax) # call function whose address is at  
# known offset in the vtable \*

- \* Or can use: addq \$offset<sub>m</sub>,%rax  
call \*(%rax)  
or : movq \$offset<sub>m</sub>(%rax),%rax  
call \*%rax


# Runtime Type Checking

- We can use the method table for the class as a “runtime representation” of the class
  - Each class has one vtable at a unique address
- The test for “o instanceof C” is
  - Is o’s method table pointer == &C\$\$ ?
    - If so, result is “true”
  - Recursively, get pointer to superclass method table from the method table and check that
  - Stop when you reach Object (or a null pointer, depending on whether there is a ultimate superclass of everything)
    - If no match by the top of the chain, result is “false”
- Same test as part of check for legal downcast (e.g., how to test for ClassCastException in (type)obj cast)



## Coming (& Past) Attractions

- Code analysis and optimization
- Industrial-strength back end (register allocation, instruction selection & scheduling)
- Other topics as time allows
  - GC? Dynamic languages? JVM? What else?
- And simple code generation for the project



[Meme credit: imgflip.com]