

CS 6410: Compilers

Fall 2023

Tamara Bonaci
t.bonaci@northeastern.edu

Thank you to UW faculty Hal Perkins. Today lecture notes are a modified version of his lecture notes.

Credits For Course Material

- Big thank you to UW CSE faculty member, Hal Perkins
- Some direct ancestors of this course:
 - UW CSE 401 (Chambers, Snyder, Notkin, Perkins, Ringenburt, Henry, ...)
 - UW CSE PMP 582/501 (Perkins)
 - Cornell CS 412-3 (Teitelbaum, Perkins)
 - Rice CS 412 (Cooper, Kennedy, Torczon)
 - Many books (Appel; Cooper/Torczon; Aho, [[Lam,] Sethi,] Ullman [Dragon Book], Fischer, [Cytron ,] LeBlanc; Muchnick, ...)

Agenda

- Optimization and transformation
 - Survey of some code “optimizations” (improvements)
 - Some organizing concepts
 - Basic blocks
 - Control-flow and dataflow graph
 - Analysis vs. transformation
 - A closer look at some common optimizing transformations

Reading:

Cooper and Torczon, chapters 4.1-4.4, 5.5, 6.2-6.5 and 7.1-7.4, 8.1-8.6

Dragon book, chapters 6.3-6.5, 7.1-7.7, 8.1-8.4, 9.1

Review: Optimizations

Review: Kinds of Optimizations

- **Peephole** - look at adjacent instructions
- **Local** - look at individual *basic blocks*
 - straight-line sequence of statements
- **Intraprocedural** - look at the whole procedure
 - Commonly called “global”
- **Interprocedural** - look across procedures
 - “whole program” analysis
 - gcc’s “link time optimization” is a version of this
- **Larger scope** - usually better optimization but more cost and complexity
 - Analysis is often less precise because of more possibilities

Review: Peephole Optimization

- Look at adjacent instructions (a “peephole” on the code stream)
 - try to replace adjacent instructions with something faster

movq %r9,16(%rsp) movq 16(%rsp),%r12	movq %r9,16(%rsp) movq %r9,%r12
---	--

- Jump chaining can also be considered a form of peephole optimization (removing jump to jump)

Review: Algebraic Simplification

- “constant folding”, “strength reduction”
 - $z = 3 + 4;$ $\rightarrow z = 7$
 - $z = x + 0;$ $\rightarrow z = x$
 - $z = x * 1;$ $\rightarrow z = x$
 - $z = x * 2;$ $\rightarrow z = x \ll 1$ or $z = x + x$
 - $z = x * 8;$ $\rightarrow z = x \ll 3$
 - $z = x / 8;$ $\rightarrow z = x \gg 3$ (only if $x \geq 0$ known)
 - $z = (x + y) - y;$ $\rightarrow z = x$ (maybe; not doubles, might change
int overflow)
- Can be done at many levels from peephole on up
- Why do these examples happen?
 - Often created during conversion to lower-level IR, by other optimizations, code gen, etc.

Review: Local Optimizations

- Analysis and optimizations within a basic block
- *Basic block*: straight-line sequence of statements
 - no control flow into or out of middle of sequence
- Better than peephole
- Not too hard to implement with reasonable IR
- Machine-independent, if done on IR

Review: Local Constant Propagation

- If variable assigned a constant, replace downstream uses of the variable with constant (until variable reassigned)
- Can enable more constant folding
 - Code; unoptimized intermediate code:

```
count = 10;  
...    // count not changed  
x = count * 5;  
y = x ^ 3;  
x = 7;
```

```
count = 10;  
t1 = count;  
t2 = 5;  
t3 = t1 * t2;  
x = t3;  
t4 = x;  
t5 = 3;  
t6 = exp(t4, t5);  
y = t6;  
x = 7
```

Review: Local Constant Propagation

- If variable assigned a constant, replace downstream uses of the variable with constant (until variable reassigned)
- Can enable more constant folding
 - Code; constant propagation:

```
count = 10;
... // count not changed
x = count * 5;
y = x ^ 3;
x = 7;
```

```
count = 10;
t1 = 10;           // cp count
t2 = 5;
t3 = 10 * t2;      // cp t1
x = t3;
t4 = x;
t5 = 3;
t6 = exp(t4,3);    // cp t5
y = t6;
x = 7
```

Review: Local Constant Propagation

- If variable assigned a constant, replace downstream uses of the variable with constant (until variable reassigned)
- Can enable more constant folding
 - Code; constant folding:

```
count = 10;
...    // count not changed
x = count * 5;
y = x ^ 3;
x = 7;
```

```
count = 10;
t1 = 10;
t2 = 5;
t3 = 50;           // 10*t2
x = t3;
t4 = x;
t5 = 3;
t6 = exp(t4, 3);
y = t6;
x = 7;
```

Review: Local Constant Propagation

- If variable assigned a constant, replace downstream uses of the variable with constant (until variable reassigned)
- Can enable more constant folding
 - Code; repropagated intermediate code

```
count = 10;
...    // count not changed
x = count * 5;
y = x ^ 3;
x = 7;
```

```
count = 10;
t1 = 10;
t2 = 5;
t3 = 50;
x = 50;          // cp t3
t4 = 50;         // cp x
t5 = 3;
t6 = exp(50,3);  // cp t4
y = t6;
x = 7;
```

Review: Local Constant Propagation

- If variable assigned a constant, replace downstream uses of the variable with constant (until variable reassigned)
- Can enable more constant folding
 - Code; refold intermediate code

```
count = 10;  
... // count not changed  
x = count * 5;  
y = x ^ 3;  
x = 7;
```

```
count = 10;  
t1 = 10;  
t2 = 5;  
t3 = 50;  
x = 50;  
t4 = 50;  
t5 = 3;  
t6 = 125000; // cf 50^3  
y = t6;  
x = 7;
```

Review: Local Constant Propagation

- If variable assigned a constant, replace downstream uses of the variable with constant (until variable reassigned)
- Can enable more constant folding
 - Code; repropagated intermediate code

```
count = 10;  
... // count not changed  
x = count * 5;  
y = x ^ 3;  
x = 7;
```

```
count = 10;  
t1 = 10;  
t2 = 5;  
t3 = 50;  
x = 50;  
t4 = 50;  
t5 = 3;  
t6 = 125000;  
y = 125000; // cp t6  
x = 7;
```

Local Dead Assignment Elimination

- If l.h.s. of assignment never referenced again before being overwritten, then can delete assignment
 - Why would this happen?
Clean-up after previous optimizations, often

```
count = 10;  
... // count not changed  
x = count * 5;  
y = x ^ 3;  
x = 7;
```

```
count = 10;  
t1 = 10;  
t2 = 5;  
t3 = 50;  
x = 50;  
t4 = 50;  
t5 = 3;  
t6 = 125000;  
y = 125000;  
x = 7;
```

Local Dead Assignment Elimination

- If l.h.s. of assignment never referenced again before being overwritten, then can delete assignment
 - Why would this happen?
Clean-up after previous optimizations, often

```
count = 10;  
... // count not changed  
x = count * 5;  
y = x ^ 3;  
x = 7;
```

```
count = 10;  
t1 = 10;  
t2 = 5;  
t3 = 50;  
x = 50;  
t4 = 50;  
t5 = 3;  
t6 = 125000;  
y = 125000;  
x = 7;
```


Local Common Subexpression Elimination

- Look for repetitions of the same computation. Eliminate them if result won't have changed and no side effects
 - Avoid repeated calculation and eliminates redundant loads
- Idea: walk through basic block keeping track of available expressions

```
... a[i] + b[i] ...
```

```
t1 = *(fp + ioffset);  
t2 = t1 * 4;  
t3 = fp + t2;  
t4 = *(t3 + aoffset);  
t5 = *(fp + ioffset);  
t6 = t5 * 4;  
t7 = fp + t6;  
t8 = *(t7 + boffset);  
t9 = t4 + t8;
```

Local Common Subexpression Elimination

- Look for repetitions of the same computation. Eliminate them if result won't have changed and no side effects
 - Avoid repeated calculation and eliminates redundant loads
- Idea: walk through basic block keeping track of available expressions

```
... a[i] + b[i] ...
```

```
t1 = *(fp + ioffset);  
t2 = t1 * 4;  
t3 = fp + t2;  
t4 = *(t3 + aoffset);  
t5 = t1;    // CSE  
t6 = t5 * 4;  
t7 = fp + t6;  
t8 = *(t7 + boffset);  
t9 = t4 + t8;
```

Local Common Subexpression Elimination

- Look for repetitions of the same computation. Eliminate them if result won't have changed and no side effects
 - Avoid repeated calculation and eliminates redundant loads
- Idea: walk through basic block keeping track of available expressions

```
... a[i] + b[i] ...
```

```
t1 = *(fp + ioffset);  
t2 = t1 * 4;  
t3 = fp + t2;  
t4 = *(t3 + aoffset);  
t5 = t1;  
t6 = t1 * 4; // CP  
t7 = fp + t6;  
t8 = *(t7 + boffset);  
t9 = t4 + t8;
```

Local Common Subexpression Elimination

- Look for repetitions of the same computation. Eliminate them if result won't have changed and no side effects
 - Avoid repeated calculation and eliminates redundant loads
- Idea: walk through basic block keeping track of available expressions

```
... a[i] + b[i] ...
```

```
t1 = *(fp + ioffset);  
t2 = t1 * 4;  
t3 = fp + t2;  
t4 = *(t3 + aoffset);  
t5 = t1;  
t6 = t2;           // CSE  
t7 = fp + t2;      // CP  
t8 = *(t7 + boffset);  
t9 = t4 + t8;
```

Local Common Subexpression Elimination

- Look for repetitions of the same computation. Eliminate them if result won't have changed and no side effects
 - Avoid repeated calculation and eliminates redundant loads
- Idea: walk through basic block keeping track of available expressions

```
... a[i] + b[i] ...
```

```
t1 = *(fp + ioffset);  
t2 = t1 * 4;  
t3 = fp + t2;  
t4 = *(t3 + aoffset);  
t5 = t1;  
t6 = t2;  
t7 = t3; // CSE  
t8 = *(t3 + boffset); //CP  
t9 = t4 + t8;
```

Local Common Subexpression Elimination

- Look for repetitions of the same computation. Eliminate them if result won't have changed and no side effects
 - Avoid repeated calculation and eliminates redundant loads
- Idea: walk through basic block keeping track of available expressions

```
... a[i] + b[i] ...
```

```
t1 = *(fp + ioffset);  
t2 = t1 * 4;  
t3 = fp + t2;  
t4 = *(t3 + aoffset);  
t5 = t1; // DAE  
t6 = t2; // DAE  
t7 = t3; // DAE  
t8 = *(t3 + boffset);  
t9 = t4 + t8;
```

Intraprocedural Optimizations

- Enlarge scope of analysis to whole procedure
 - more opportunities for optimization
 - have to deal with branches, merges, and loops
- Can do constant propagation, common subexpression elimination, etc. at “global” level
- Can do new things, e.g. loop optimizations
- Optimizing compilers usually work at this level (-O2)

Code Motion

- Goal: move loop-invariant calculations out of loops
- Can do at source level or at intermediate code level

```
for (i = 0; i < 10; i = i+1) {  
    a[i] = a[i] + b[j];  
    z = z + 10000;  
}
```

```
t1 = b[j];  
t2 = 10000;  
for (i = 0; i < 10; i = i+1) {  
    a[i] = a[i] + t1;  
    z = z + t2;  
}
```


Code Motion at Intermediate Level

```
for (i = 0; i < 10; i = i+1) {  
    a[i] = b[j];  
}
```

```
*(fp + ioffset) = 0;  
label top;  
    t0 = *(fp + ioffset);  
    iffalse (t0 < 10) goto done;  
    t1 = *(fp + joffset);  
    t2 = t1 * 4;  
    t3 = fp + t2;  
    t4 = *(t3 + boffset);  
    t5 = *(fp + ioffset);  
    t6 = t5 * 4;  
    t7 = fp + t6;  
    *(t7 + aoffset) = t4;  
    t9 = *(fp + ioffset);  
    t10 = t9 + 1;  
    *(fp + ioffset) = t10;  
    goto top;  
label done;
```

Code Motion at Intermediate Level

```
for (i = 0; i < 10; i = i+1) {  
    a[i] = b[j];  
}
```

```
t11 = fp + ioffset; t13 = fp + aoffset;  
t12 = fp + joffset; t14 = fp + boffset  
*(fp + ioffset) = 0;  
label top;  
    t0 = *t11;  
    iffalse (t0 < 10) goto done;  
    t1 = *t12;  
    t2 = t1 * 4;  
t3 = t14;  
    t4 = *(t14 + t2);  
    t5 = *t11;  
    t6 = t5 * 4;  
t7 = t13;  
    *(t13 + t6) = t4;  
    t9 = *t11;  
    t10 = t9 + 1;  
    *t11 = t10;  
    goto top;  
label done;
```

Loop Induction Variable Elimination

- A special and common case of loop-based strength reduction
- For-loop index is *induction variable*
 - incremented each time around loop
 - offsets & pointers calculated from it
- If used only to index arrays, can rewrite with pointers
 - compute initial offsets/pointers before loop
 - increment offsets/pointers each time around loop
 - no expensive scaling in loop
 - can then do loop-invariant code motion

```
for (i = 0; i < 10; i = i+1) {  
    a[i] = a[i] + x;  
}  
=> transformed to  
for (p = &a[0]; p < &a[10]; p = p+4) {  
    *p = *p + x;  
}
```

Interprocedural Optimization

- Expand scope of analysis to procedures calling each other
- Can do local & intraprocedural optimizations at larger scope
- Can do new optimizations, e.g. inlining

Inlining: Replace Call With Body

- Replace procedure call with body of called procedure

- Source:

```
final double pi = 3.1415927;  
double circle_area(double radius) {  
    return pi * (radius * radius);  
}
```

...

```
double r = 5.0;
```

...

```
double a = circle_area(r);
```

- After inlining:

...

```
double r = 5.0;
```

...

```
double a = pi * r * r;
```

- (Then what? Constant propagation/folding)

An example

```
x = a[i] + b[2];  
c[i] = x - 5;
```

```
t1 = *(fp + ioffset); // i  
t2 = t1 * 4;  
t3 = fp + t2;  
t4 = *(t3 + aoffset); // a[i]  
t5 = 2;  
t6 = t5 * 4;  
t7 = fp + t6;  
t8 = *(t7 + boffset); // b[2]  
t9 = t4 + t8;  
*(fp + xoffset) = t9; // x = ...  
t10 = *(fp + xoffset); // x  
t11 = 5;  
t12 = t10 - t11;  
t13 = *(fp + ioffset); // i  
t14 = t13 * 4;  
t15 = fp + t14;  
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];
c[i] = x - 5;
```

Strength reduction: shift
often cheaper than multiply

```
t1 = *(fp + ioffset); // i
t2 = t1 << 2; // was t1 * 4
t3 = fp + t2;
t4 = *(t3 + aoffset); // a[i]
t5 = 2;
t6 = t5 << 2; // was t5 * 4
t7 = fp + t6;
t8 = *(t7 + boffset); // b[2]
t9 = t4 + t8;
*(fp + xoffset) = t9; // x = ...
t10 = *(fp + xoffset); // x
t11 = 5;
t12 = t10 - t11;
t13 = *(fp + ioffset); // i
t14 = t13 << 2; // was t13 * 4
t15 = fp + t14;
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];
c[i] = x - 5;
```

Constant propagation:
replace variables with
known constant values

```
t1 = *(fp + ioffset); // i
t2 = t1 << 2;
t3 = fp + t2;
t4 = *(t3 + aoffset); // a[i]
t5 = 2;
t6 = 2 << 2; // was t5 << 2
t7 = fp + t6;
t8 = *(t7 + boffset); // b[2]
t9 = t4 + t8;
*(fp + xoffset) = t9; // x = ...
t10 = *(fp + xoffset); // x
t11 = 5;
t12 = t10 - 5; // was t10 - t11
t13 = *(fp + ioffset); // i
t14 = t13 << 2;
t15 = fp + t14;
*(t15 + coffset) = t12; // c[i] := ...
```


An example

```
x = a[i] + b[2];
c[i] = x - 5;
```

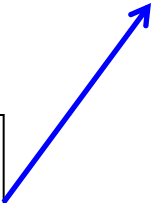
Dead store (or dead assignment) elimination:
remove assignments to provably unused variables

```
t1 = *(fp + ioffset); // i
t2 = t1 << 2;
t3 = fp + t2;
t4 = *(t3 + aoffset); // a[i]
t5 = 2;
t6 = 2 << 2;
t7 = fp + t6;
t8 = *(t7 + boffset); // b[2]
t9 = t4 + t8;
*(fp + xoffset) = t9; // x = ...
t10 = *(fp + xoffset); // x
t11 = 5;
t12 = t10 - 5;
t13 = *(fp + ioffset); // i
t14 = t13 << 2;
t15 = fp + t14;
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];  
c[i] = x - 5;
```

Constant folding: statically
compute operations
with known constant values



```
t1 = *(fp + ioffset); // i  
t2 = t1 << 2;  
t3 = fp + t2;  
t4 = *(t3 + aoffset); // a[i]  
t6 = 8; // was 2 << 2  
t7 = fp + t6;  
t8 = *(t7 + boffset); // b[2]  
t9 = t4 + t8;  
*(fp + xoffset) = t9; // x = ...  
t10 = *(fp + xoffset); // x  
t12 = t10 - 5;  
t13 = *(fp + ioffset); // i  
t14 = t13 << 2;  
t15 = fp + t14;  
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];  
c[i] = x - 5;
```

Constant propagation then
dead store elimination



```
t1 = *(fp + ioffset); // i  
t2 = t1 << 2;  
t3 = fp + t2;  
t4 = *(t3 + aoffset); // a[i]  
t6 = 8;  
t7 = fp + 8; // was fp + t6  
t8 = *(t7 + boffset); // b[2]  
t9 = t4 + t8;  
*(fp + xoffset) = t9; // x = ...  
t10 = *(fp + xoffset); // x  
t12 = t10 - 5;  
t13 = *(fp + ioffset); // i  
t14 = t13 << 2;  
t15 = fp + t14;  
(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];  
c[i] = x - 5;
```

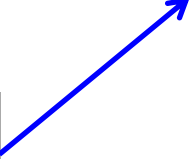
Arithmetic identities: + is commutative & associative. boffset is typically a known, compile-time constant (say -32), so this enables...

```
t1 = *(fp + ioffset); // i  
t2 = t1 << 2;  
t3 = fp + t2;  
t4 = *(t3 + aoffset); // a[i]  
t7 = boffset + 8; // was fp + 8  
t8 = *(t7 + fp); // b[2] (was t7 + boffset)  
t9 = t4 + t8;  
*(fp + xoffset) = t9; // x = ...  
t10 = *(fp + xoffset); // x  
t12 = t10 - 5;  
t13 = *(fp + ioffset); // i  
t14 = t13 << 2;  
t15 = fp + t14;  
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];  
c[i] = x - 5;
```

... more **constant folding**,
which in turn enables ...



```
t1 = *(fp + ioffset); // i  
t2 = t1 << 2;  
t3 = fp + t2;  
t4 = *(t3 + aoffset); // a[i]  
t7 = -24;           // was boffset (-32) + 8  
t8 = *(t7 + fp);     // b[2]  
t9 = t4 + t8;  
*(fp + xoffset) = t9; // x = ...  
t10 = *(fp + xoffset); // x  
t12 = t10 - 5;  
t13 = *(fp + ioffset); // i  
t14 = t13 << 2;  
t15 = fp + t14;  
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];  
c[i] = x - 5;
```

More constant propagation
and dead store elimination

```
t1 = *(fp + ioffset); // i  
t2 = t1 << 2;  
t3 = fp + t2;  
t4 = *(t3 + aoffset); // a[i]  
t7 = 24;  
t8 = *(fp - 24); // b[2] (was t7+fp)  
t9 = t4 + t8;  
*(fp + xoffset) = t9; // x = ...  
t10 = *(fp + xoffset); // x  
t12 = t10 - 5;  
t13 = *(fp + ioffset); // i  
t14 = t13 << 2;  
t15 = fp + t14;  
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];
c[i] = x - 5;
```

Common subexpression
elimination – no need to
compute $*(fp + ioffset)$ again
if we know it won't change

```
t1 = *(fp + ioffset); // i
t2 = t1 << 2;
t3 = fp + t2;
t4 = *(t3 + aoffset); // a[i]
t8 = *(fp - 24);      // b[2]
t9 = t4 + t8;
*(fp + xoffset) = t9; // x = ...
t10 = *(fp + xoffset); // x
t12 = t10 - 5;
t13 = t1;      // i (was *(fp + ioffset))
t14 = t13 << 2;
t15 = fp + t14;
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];  
c[i] = x - 5;
```

Copy propagation: replace assignment targets with their values (e.g., replace t13 with t1)

```
t1 = *(fp + ioffset); // i  
t2 = t1 << 2;  
t3 = fp + t2;  
t4 = *(t3 + aoffset); // a[i]  
t8 = *(fp - 24);      // b[2]  
t9 = t4 + t8;  
*(fp + xoffset) = t9; // x = ...  
t10 = t9;             // x (was *(fp + xoffset))  
t12 = t10 - 5;  
t13 = t1;             // i  
t14 = t1 << 2;        // was t13 << 2  
t15 = fp + t14;  
*(t15 + coffset) = t12; // c[i] := ...
```


An example

```
x = a[i] + b[2];
c[i] = x - 5;
```

Common subexpression
elimination



```
t1 = *(fp + ioffset); // i
t2 = t1 << 2;
t3 = fp + t2;
t4 = *(t3 + aoffset); // a[i]
t8 = *(fp - 24);      // b[2]
t9 = t4 + t8;
*(fp + xoffset) = t9; // x = ...
t10 = t9;             // x
t12 = t10 - 5;
t13 = t1;             // i
t14 = t2;             // was t1 << 2
t15 = fp + t14;
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];  
c[i] = x - 5;
```

More [copy propagation](#)



```
t1 = *(fp + ioffset); // i  
t2 = t1 << 2;  
t3 = fp + t2;  
t4 = *(t3 + aoffset); // a[i]  
t8 = *(fp - 24);      // b[2]  
t9 = t4 + t8;  
*(fp + xoffset) = t9; // x = ...  
t10 = t9;             // x  
t12 = t9 - 5; // was t10 - 5  
t13 = t1;             // i  
t14 = t2;  
t15 = fp + t14;  
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];
c[i] = x - 5;
```

More [copy propagation](#)



```
t1 = *(fp + ioffset); // i
t2 = t1 << 2;
t3 = fp + t2;
t4 = *(t3 + aoffset); // a[i]
t8 = *(fp - 24);      // b[2]
t9 = t4 + t8;
*(fp + xoffset) = t9; // x = ...
t10 = t9;             // x
t12 = t9 - 5;
t13 = t1;             // i
t14 = t2;
t15 = fp + t2; // was fp + t14
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];
c[i] = x - 5;
```

Dead assignment
elimination

```
t1 = *(fp + ioffset); // i
t2 = t1 << 2;
t3 = fp + t2;
t4 = *(t3 + aoffset); // a[i]
t8 = *(fp - 24); // b[2]
t9 = t4 + t8;
*(fp + xoffset) = t9; // x = ...
t10 = t9; // x
t12 = t9 - 5;
t13 = t1; // i
t14 = t2;
t15 = fp + t2;
*(t15 + coffset) = t12; // c[i] := ...
```

An example

```
x = a[i] + b[2];  
c[i] = x - 5;
```

```
t1 = *(fp + ioffset); // i  
t2 = t1 << 2;  
t3 = fp + t2;  
t4 = *(t3 + aoffset); // a[i]  
t8 = *(fp - 24);      // b[2]  
t9 = t4 + t8;  
*(fp + xoffset) = t9; // x = ...  
t12 = t9 - 5;  
t15 = fp + t2;  
*(t15 + coffset) = t12; // c[i] := ...
```

- **Final:** 3 loads (i, a[i], b[2]), 2 stores (x, c[i]), 5 register-only moves, 9 +/-, 1 shift
- **Original:** 5 loads, 2 stores, 10 register-only moves, 12 +/-, 3 *
- **Optimizer note:** we usually leave assignment of actual registers to later stage of the compiler and assume as many “pseudo registers” as we need here

Some Frequent Compiler Optimization Techniques

- **Strength reduction** – replace an “expensive” operation with an equivalent, but less expensive operation (e.g., multiplication → summation/shift)
- **Constant propagation** – substitute values of known constants at compile time
- **Constant folding** – recognize and evaluate a constant at compile time rather than run time
- **Dead assignment elimination** – recognize assignments that never referenced, and remove them from the code
- **Common subexpression elimination** - find repetitions of same computations, and eliminate them if result won't changed
- **Code motion** - move loop-invariant calculations out of loops
- **Inlining** – replace some function calls with the body of the function (e.g., some getters)

Data Structures for Optimizations

- Need to represent control and data flow
- **Control flow graph (CFG)** captures flow of control:
 - nodes are IL statements, or whole basic blocks
 - edges represent (all possible) control flow
 - node with multiple successors = branch/switch
 - node with multiple predecessors = merge
 - loop in graph = loop
- **Data flow graph (DFG)** captures flow of data (e.g. def/use chains):
 - nodes are def(inition)s and uses
 - edge from def to use
 - a def can reach multiple uses
 - a use can have multiple reaching defs (different control flow paths, possible aliasing, etc.)
- **SSA**: another widely used way of linking defs and uses

Summary

- **Optimizations** organized as collections of passes, each rewriting IL in place into (hopefully) better version
- Each pass does analysis to determine what is possible, followed by transformation(s) that (hopefully) improve the program
 - Sometimes “analysis-only” passes are helpful
 - Often redo analysis/transformations again to take advantage of possibilities revealed by previous changes
- Presence of optimizations makes other parts of compiler (e.g. intermediate and target code generation) easier to write

Analysis and Transformation

Analysis and Transformation

- Each **optimization** is made up of
 - Some number of **analyses**
 - Followed by a **transformation**
- Analyze CFG and/or DFG by propagating info forward or backward along CFG and/or DFG edges
 - Merges in graph require combining info
 - Loops in graph require *iterative approximation*
- Perform (improving) transformations based on info computed
- Analysis must be conservative/safe/sound so that transformations preserve program behavior


Role of Transformations

- Dataflow analysis discovers opportunities for code improvement
- Compiler rewrites the (IR) to make these improvements
 - Transformation may reveal additional opportunities for further optimization
 - May also block opportunities by obscuring information

Organizing Transformations in a Compiler

- Typically middle end consists of many phases
 - Analyze IR
 - Identify optimization
 - Rewrite IR to apply optimization
 - And repeat (50 phases in a commercial compiler is typical)
- Each individual optimization is supported by rigorous formal theory
- But no formal theory for what order or how often to apply them(!)
 - Some rules of thumb and best practices
 - May apply some transformations several times as different phases reveal opportunities for further improvement

Optimization 'Phases'



- Each optimization requires a 'pass' (linear scan) over the IR
- IR may sometimes shrink, sometimes expand
- Some optimizations may be repeated
- 'Best' ordering is heuristic
- Don't try to *beat* an optimizing compiler - you will lose!
- **Note:** not all programs are written by humans!
- Machine-generated code can pose a challenge for optimizers
 - eg: a single function with 10,000 statements, 1,000+ local variables, loops nested 15 deep, spaghetti of "GOTOs", etc

A Taxonomy

- Machine Independent Transformations
 - Mostly independent of target machine
(e.g., loop unrolling will likely make it faster regardless of target)
 - “Mostly”? – e.g., vectorize only if target has SIMD ops
 - Worthwhile investment – applies to all targets
- Machine Dependent Transformations
 - Mostly concerned with instruction selection & scheduling, register allocation
 - Need to tune for different targets
 - Most of this in the back end, but some in the optimizer

Machine Independent Transformations

- Dead code elimination
 - unreachable or not actually used later
- Code motion
 - “hoist” loop-invariant code out of a loop
- Specialization
- Strength reduction
 - $2 * x \Rightarrow x + x$; $@A + ((i * \text{numcols} + j) * \text{eltsize}) \Rightarrow p += 4$
- Enable *other* transformations
- Eliminate redundant computations
 - Value numbering, GCSE

Machine Dependent Transformations

- Take advantage of special hardware
 - e.g., expose instruction-level parallelism (ILP)
 - e.g., use special instructions (VAX polyf; x86 sqrt, strings)
 - e.g., use SIMD instructions and registers
- Manage or hide latencies
 - e.g., tiling/blocking and loop interchange
 - Improves cache behavior – hugely important
- Deal with finite resources - # functional units
- Compilers generate for a vanilla machine, e.g., SSE2
 - But provide switches to tune (arch:AVX, arch:IA32)
 - JIT compiler knows its target architecture!

Optimizer Contracts

- **Prime directive**

- No optimization will change observable program behavior!
- This can be subtle. e.g.:
 - What is "observable"? (via IO? to another thread?)
 - Dead-Code-Eliminate a *throw*?
 - Language Reference Manual may be ambiguous/undefined/negotiable for edge cases

- **Avoid harmful optimizations**

- If an optimization does not improve code significantly, don't do it: it harms throughput
- If an optimization degrades code quality, don't do it

Is this *hoist* legal?

```
for (int i = start; i < finish; ++i) a[i] += 7;
```

```
i = start
loop:
  if (i >= finish) goto done
  if (i < 0 || i >= a.length) throw OutOfBounds
  a[i] += 7
  goto loop
done:
```

```
if (start < 0 || finish >= a.length) throw OutOfBounds
i = start
loop:
  if (i >= finish) goto done
  a[i] += 7
  goto loop
done:
```

Another example: "volatile" pretty much kills all attempts to optimize

Dead Code Elimination


- If a compiler can prove that a computation has no external effect, it can be removed
 - Unreachable operations – always safe to remove
 - Useless operations – reachable, may be executed, but results not actually required
- Dead code often results from other transformations
 - Often want to do DCE several times

Dead Code Elimination


- Classic algorithm is similar to garbage collection
 - Pass I – Mark all useful operations
 - Instructions whose result does, or can, affect visible behavior:
 - Input or Output
 - Updates to object fields that might be used later
 - Instructions that may throw an exception (e.g.: array bounds check)
 - Calls to functions that might perform IO or affect visible behavior
 - (Remember, for many languages, compiler does not process entire program at one time – but a JIT compiler might be able to)
 - Mark all useful instructions
 - Repeat until no more changes
 - Pass II – delete all unmarked operations

Code Motion

- **Idea:** move an operation to a location where it is executed less frequently
 - **Classic situation:** *hoist* loop-invariant code: execute once, rather than on every iteration
- Lazy code motion & *partial* redundancy



`a` must be re-calculated - wasteful if control took right-hand arm



Replicate, so `a` need not be re-calculated

Specialization I

- **Idea:** replace general operation in IR with more specific
 - Constant folding:
 - $\text{feet_per_minute} = \text{mph} * \text{feet_per_mile} / \text{minutes_per_hour}$
 - $\text{feet_per_minute} = \text{mph} * 5280 / 60$
 - $\text{feet_per_minute} = \text{mph} * 88$
 - Replacing multiplications and division by constants with shifts (when safe)
 - Peephole optimizations
 - `movl $0,%eax => xorl %eax,%eax`

Specialization:2 - Eliminate Tail Recursion

- Factorial - recursive
`int fac(n) = if (n <= 2) return 1; else return n * fac(n - 1);`
- 'accumulating' Factorial - tail-recursive
`facaux(n, r) = if (n <= 2) return 1; else return facaux(n - 1, n*r)`
`call facaux(n, 1)`
- Optimize-away the call overhead; replace with simple jump
`facaux(n, r) = if (n <= 2) return 1;`
`else n = n - 1; r = n*r; jump back to start of facaux`
 - So replace recursive call with a loop and just one stack frame
- Issue?
 - Avoid stack overflow - good! - "observable" change?

Strength Reduction

- Classic example: Array references in a loop

```
for (k = 0; k < n; k++) a[k] = 0;
```

- Naive codegen for $a[k] = 0$ in loop body

```
movl $4,%eax           // elemsize = 4 bytes
imull offsetk(%rbp),%eax // k * elemsize
addl offseta(%rbp),%eax // &a[0] + k * elemsize
mov  $0,(%eax)         // a[k] = 0
```

- Better!

```
movl offseta(%rbp),eax // &a[0], once-off
```

```
movl $0,(%eax)         // a[k] = 0
addl $4,%eax           // eax = &a[k+1]
```

Note: *pointers* allow a user to do this directly in C or C++
 Eg: for (p = a; p < a + n;) *p++ = 0;

Implementing Strength Reduction

- **Idea:** look for operations in a loop involving:
 - A value that does not change in the loop, the *region constant*, and
 - A value that varies systematically from iteration to iteration, the *induction variable*
- Create a new induction variable that directly computes the sequence of values produced by the original one; use an addition in each iteration to update the value

Other Common Transformations

- Inline substitution (procedure bodies)
- Cloning / Replicating
- Loop Unrolling
- Loop Unswitching

Inline Substitution - "inlining"

Class with trivial *getter*

```
class C {  
    int x;  
    int getx() { return x; }  
}
```

Method **f** calls **getx**

```
class X {  
    void f() {  
        C c = new C();  
        int total = c.getx() + 42;  
    }  
}
```

Compiler *inlines* body of **getx** into **f**

```
class X {  
    void f() {  
        C c = new C();  
        int total = c.x + 42;  
    }  
}
```

- Eliminates **call** overhead
- Opens opportunities for more optimizations
- Can be applied to large method bodies too
- Aggressive optimizer will inline 2 or more deep
- Increases total code size (memory & cache issues)
- With care, is a huge win for OO code

Code Replication

Original

```
if (x < y) {  
    p = x + y;  
} else {  
    p = z + 1;  
}  
q = p * 3;  
w = y + x;
```

Replicated code

```
if (x < y) {  
    p = x + y;  
    q = p * 3;  
    w = y + x;  
} else {  
    p = z + 1;  
    q = p * 3;  
    w = y + x;  
}
```

- + : extra opportunities to optimize in larger basic blocks (eg: LVN)
- - : increase total code size - may impact effectiveness of I-cache

Loop Unrolling

- Idea: replicate the loop body
 - More opportunity to optimize loop body
 - Increases chances for good schedules and instruction level parallelism
 - Reduces loop overhead (reduce test/jumps by 75%)
- Catches
 - must ensure unrolled code produces the same answer: "loop-carried dependency analysis"
 - code bloat
 - don't overwhelm registers

Loop Unroll Example

Original

```
for (i = 1, i <= n, i++) {  
    a[i] = a[i] + b[i];  
}
```

- Unroll 4x
- Need tidy-up loop for remainder

Unrolled

```
i = 1;  
while (i + 3 <= n) {  
    a[i] = a[i] + b[i];  
    a[i+1] = a[i+1] + b[i+1];  
    a[i+2] = a[i+2] + b[i+2];  
    a[i+3] = a[i+3] + b[i+3];  
    i += 4;  
}
```

```
while (i <= n) {  
    a[i] = a[i] + b[i];  
    i++;  
}
```

Loop Unswitching

- **Idea:** if the condition in an if-then-else is loop invariant, rewrite the loop by pulling the if-then-else out of the loop and generating a tailored copy of the loop for each half of the new conditional
 - After this transformation, both loops have simpler control flow – more chances for rest of compiler to do better

Loop Unswitch Example

Original

```
for (i = 1, i <= n, i++) {  
    if (x > y) {  
        a[i] = b[i]*x;  
    } else {  
        a[i] = b[i]*y;  
    }  
}
```

Unswitched

```
if (x > y) {  
    for (i = 1; i <= n; i++) {  
        a[i] = b[i]*x;  
    }  
} else {  
    for (i = 1; i <= n; i++) {  
        a[i] = b[i]*y;  
    }  
}
```

- IF condition does not change value in this code snippet
- No need to check $x > y$ on every iteration
- Do the IF check once!

Summary

- Just a sampler
 - 100s of transformations in the literature
 - Will examine several in more detail, particularly involving loops
- Big part of engineering a compiler is:
 - decide which transformations to use
 - decide in what order
 - decide if & when to repeat each transformation
- Compilers offer options:
 - optimize for speed
 - optimize for codesize
 - optimize for specific target micro-architecture
 - optimize for power consumption(!)
- Competitive bench-marking will investigate many permutations

Value Numbering

Optimizations

- Big part of engineering a compiler is:
 - Deciding which transformations to use
 - Deciding in what order
 - Deciding if & when to repeat each transformation
- Compilers offer options to:
 - Optimize for speed
 - Optimize for codesize
 - Optimize for specific target micro-architecture
 - Optimize for power consumption(!)
- Competitive bench-marking will investigate many permutations

“But...”

- None of these improvements are truly “optimal”
 - Hard problems (in theory-of-computation sense)
 - Proofs of optimality assume artificial restrictions
- Best we can do is to improve things
 - Most (much?) (some?) of the time
 - Realistically: try to do better for common idioms both in the code and on the machine

Issues (1)

- Safety – transformation must not change program meaning
 - Must generate correct results
 - Can't generate spurious errors
 - Optimizations must be conservative
 - Large part of analysis goes towards proving safety
 - Can pay off to speculate (be optimistic) but then need to recover if reality is different

Issues (2)

- Profitability

- If a transformation is possible, is it profitable?
- Example: loop unrolling
 - Can increase amount of work done on each iteration, i.e., reduce loop overhead
 - Can eliminate duplicate operations done on separate iterations

Issues (3)

- Downside risks
 - Even if a transformation is generally worthwhile, need to think about potential problems
 - For example:
 - Transformation might need more temporaries, putting additional pressure on registers
 - Increased code size could cause cache misses, or, in bad cases, increase page working set

Example: Redundancy Elimination

- An expression $x+y$ is *redundant* at a program point if and only if, along every path from the procedure's entry, it has been evaluated and its constituent subexpressions (x and y) have not been redefined
- If the compiler can prove the expression is redundant:
 - Can store the result of the earlier evaluation
 - Can replace the redundant computation with a reference to the earlier (stored) result

Value Numbering

- Technique for eliminating redundant expressions:
 - Assign an identifying number $VN(n)$ to each expression
 - $VN(x + y) = VN(j)$ if $x+y$ and j have the same value
 - Use hashing over value numbers for efficiency
- Old idea (Balke 1968, Ershov 1954)
 - Invented for low-level, linear IRs
 - Equivalent methods exist for tree IRs, e.g., build a DAG

Local Value Numbering

- Algorithm
 - For each operation $o = \langle op, o1, o2 \rangle$ in a block
 1. Get value numbers for operands from hash lookup
 2. Hash $\langle op, VN(o1), VN(o2) \rangle$ to get a value number for o
(If op is commutative, sort $VN(o1), VN(o2)$ first)
 3. If o already has a value number, replace o with a reference to the value
 4. If $o1$ and $o2$ are constant, evaluate o at compile time and replace with an immediate load
- If hashing behaves well, this runs in linear time

Example

Code

$a = x + y$

$b = x + y$

$a = 17$

$c = x + y$

Rewritten

Bug in Simple Example

- If we use the original names, we get in trouble when a name is reused
- Solutions
 - Be clever about which copy of the value to use (e.g., use `c=b` in last statement)
 - Create an extra temporary
 - Rename around it (best!)

Renaming

- Idea: give each value a unique name
 - a_i^j means i^{th} definition of a with $\text{VN} = j$
- Somewhat complex notation, but meaning is clear
- This is the idea behind SSA (Static Single Assignment)
 - Popular modern IR – exposes many opportunities for optimizations

Example Revisited

Code

$a = x + y$

$b = x + y$

$a = 17$

$c = x + y$

Rewritten

Simple Extensions to Value Numbering

- Constant folding
 - Add a bit that records when a value is constant
 - Evaluate constant values at compile time
 - Replace op with load immediate
- Algebraic identities: $x+0$, $x*1$, $x-x$, ...
 - Many special cases
 - Switch on op to narrow down checks needed
 - Replace result with input VN

Larger Scopes

- This algorithm works on straight-line blocks of code (basic blocks)
 - Best possible results for single basic blocks
 - Loses all information when control flows to another block
- To go further, we need to represent multiple blocks of code and the control flow between them

Optimization Categories (1)

- *Local methods*
 - Usually confined to basic blocks
 - Simplest to analyze and understand
 - Most precise information

Optimization Categories (2)

- *Superlocal methods*
 - Operate over *Extended Basic Blocks* (EBBs)
 - An EBB is a set of blocks b_1, b_2, \dots, b_n where b_1 has multiple predecessors and each of the remaining blocks b_i ($2 \leq i \leq n$) have only b_{i-1} as its unique predecessor
 - The EBB is entered only at b_1 , but may have multiple exits
 - A single block b_i can be the head of multiple EBBs (these EBBs form a tree rooted at b_i)
 - Use information discovered in earlier blocks to improve code in successors

Optimization Categories (3)

- *Regional methods*
 - Operate over scopes larger than an EBB but smaller than an entire procedure/function/method
 - Typical example: loop body
 - Difference from superlocal methods is that there may be merge points in the graph (i.e., a block with two or more predecessors)
 - Facts true at merge point are facts known to be true on all possible paths to that point

Optimization Categories (4)


- *Global methods*
 - Operate over entire procedures
 - Sometimes called *intraprocedural* methods
 - Motivation is that local optimizations sometimes have bad consequences in larger context
 - Procedure/method/function is a natural unit for analysis, separate compilation, etc.
 - Almost always need global *data-flow* analysis information for these

Optimization Categories (5)

- *Whole-program methods*
 - Operate over more than one procedure
 - Sometimes called *interprocedural* methods
 - Challenges: name scoping and parameter binding issues at procedure boundaries
 - Classic examples: inline method substitution, interprocedural constant propagation
 - Common in aggressive JIT compilers and optimizing compilers for object-oriented languages

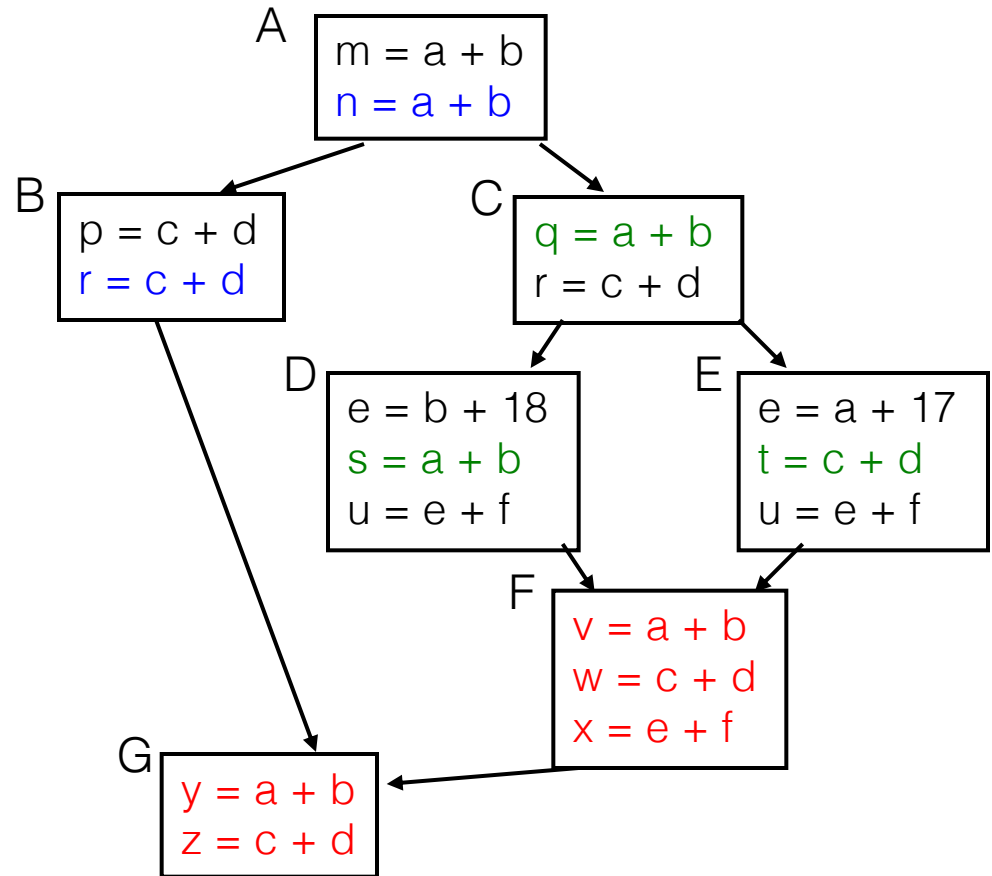
Value Numbering Revisited

- Local Value Numbering
 - 1 block at a time
 - Strong local results
 - No cross-block effects
- Missed opportunities



Superlocal Value Numbering

- Idea: apply local method to EBBs
 - {A,B}, {A,C,D}, {A,C,E}
- Final info from A is initial info for B, C; final info from C is initial for D, E
- Gets reuse from ancestors
- Avoid reanalyzing A, C
- Doesn't help with F, G



SSA Name Space

- Two Principles
 - Each name is defined by exactly one operation
 - Each operand refers to exactly one definition
- Need to deal with merge points
 - Add Φ functions at merge points to reconcile names
 - Use subscripts on variable names for uniqueness

SSA Name Space (from before)

Code

$$a_0^3 = x_0^1 + y_0^2$$

$$b_0^3 = x_0^1 + y_0^2$$

$$a_1^4 = 17$$

$$c_0^3 = x_0^1 + y_0^2$$

Rewritten

$$a_0^3 = x_0^1 + y_0^2$$

$$b_0^3 = a_0^3$$


$$a_1^4 = 17$$

$$c_0^3 = a_0^3$$

- Unique name for each definition
- Name \Leftrightarrow VN
- a_0^3 is available to assign to c_0^3


Superlocal Value Numbering with All Bells & Whistles

- Finds more redundancies
- Little extra cost
- Still does nothing for F and G



Larger Scopes

- Still have not helped F and G
- Problem: multiple predecessors
- Must decide what facts hold in F and in G
 - For G, combine B & F?
 - Merging states is expensive
 - Fall back on what we know



Dominators

- Definition
 - x *dominates* y if and only if every path from the entry of the control-flow graph to y includes x
- By definition, x dominates x
- Associate a Dom set with each node
 - $|\text{Dom}(x)| \geq 1$
- Many uses in analysis and transformation
 - Finding loops, building SSA form, code motion

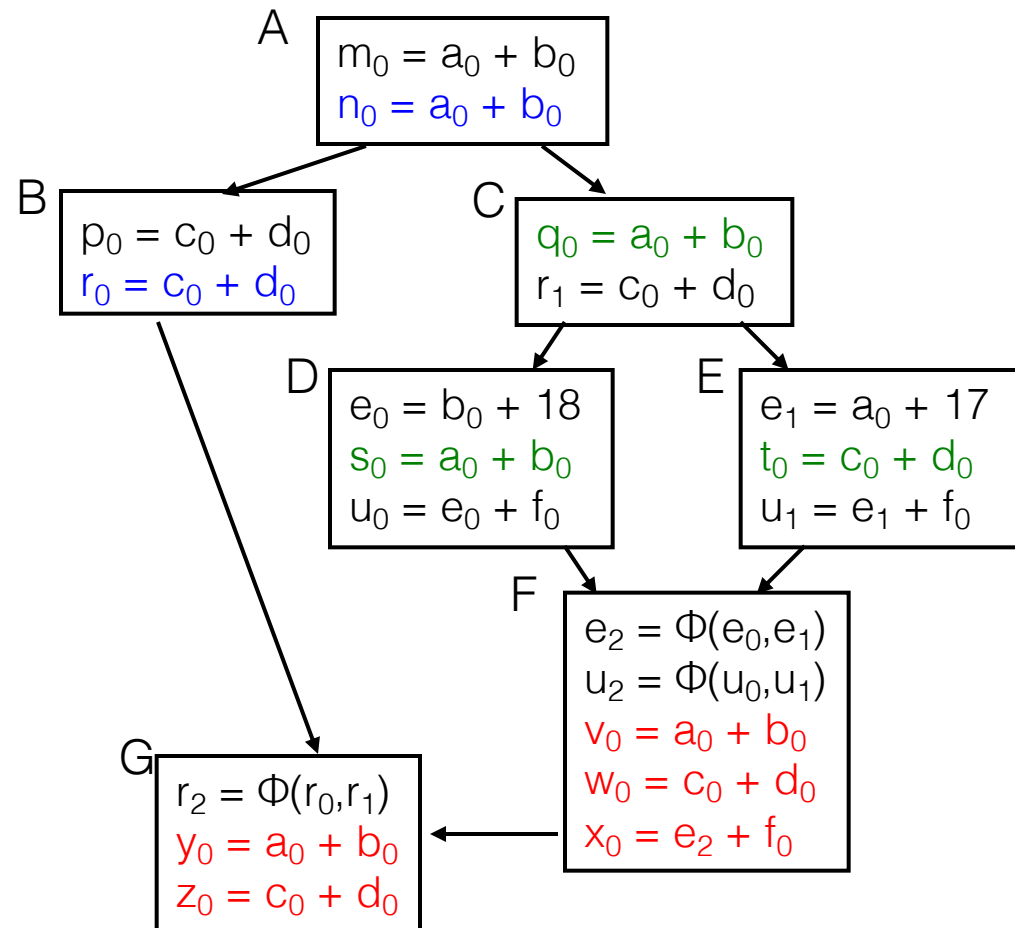
Immediate Dominators

- For any node x , there is a y in $\text{Dom}(x)$ closest to x
- This is the *immediate dominator* of x
 - Notation: $\text{IDom}(x)$

Dominator Sets

Block Dom

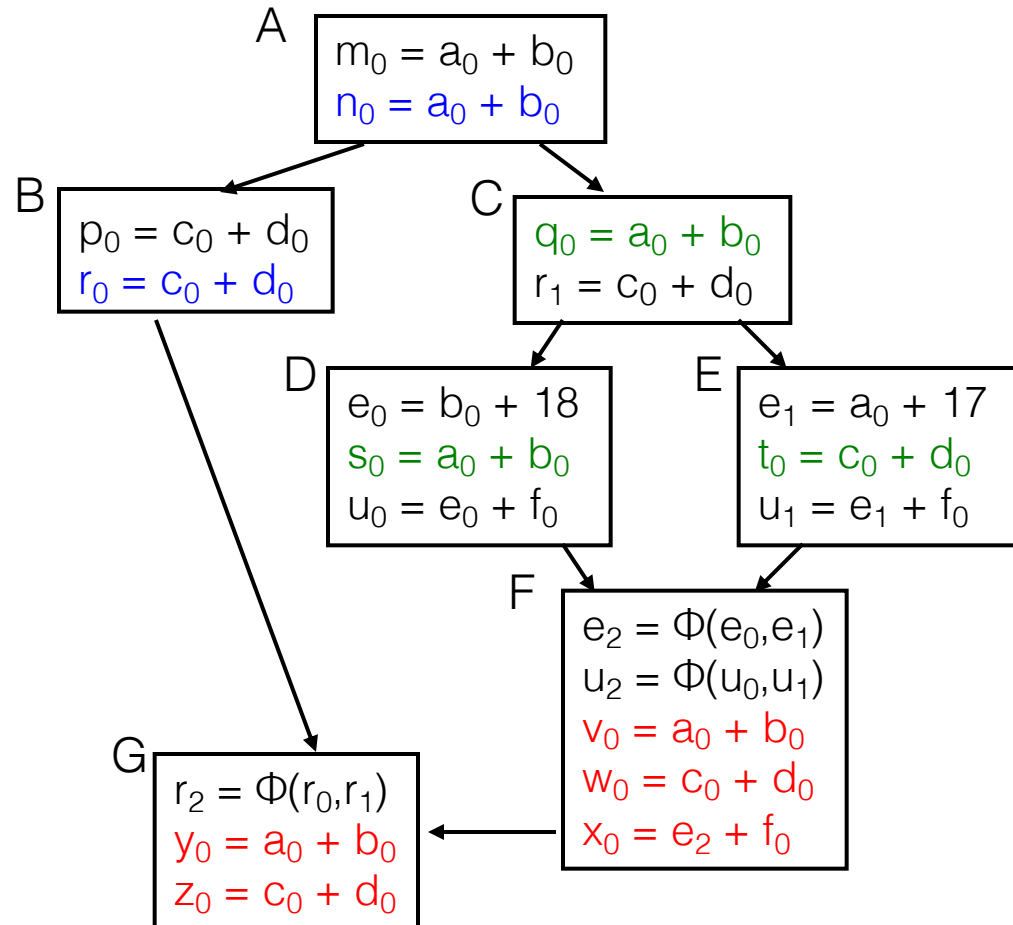
IDom



Note that the IDOM relation defines a tree!

Dominator Value Numbering

- Still looking for a way to handle F and G
- Idea: Use info from $IDom(x)$ to start analysis of x
 - Use C for F and A for G
- Dominator VN
Technique (DVNT)

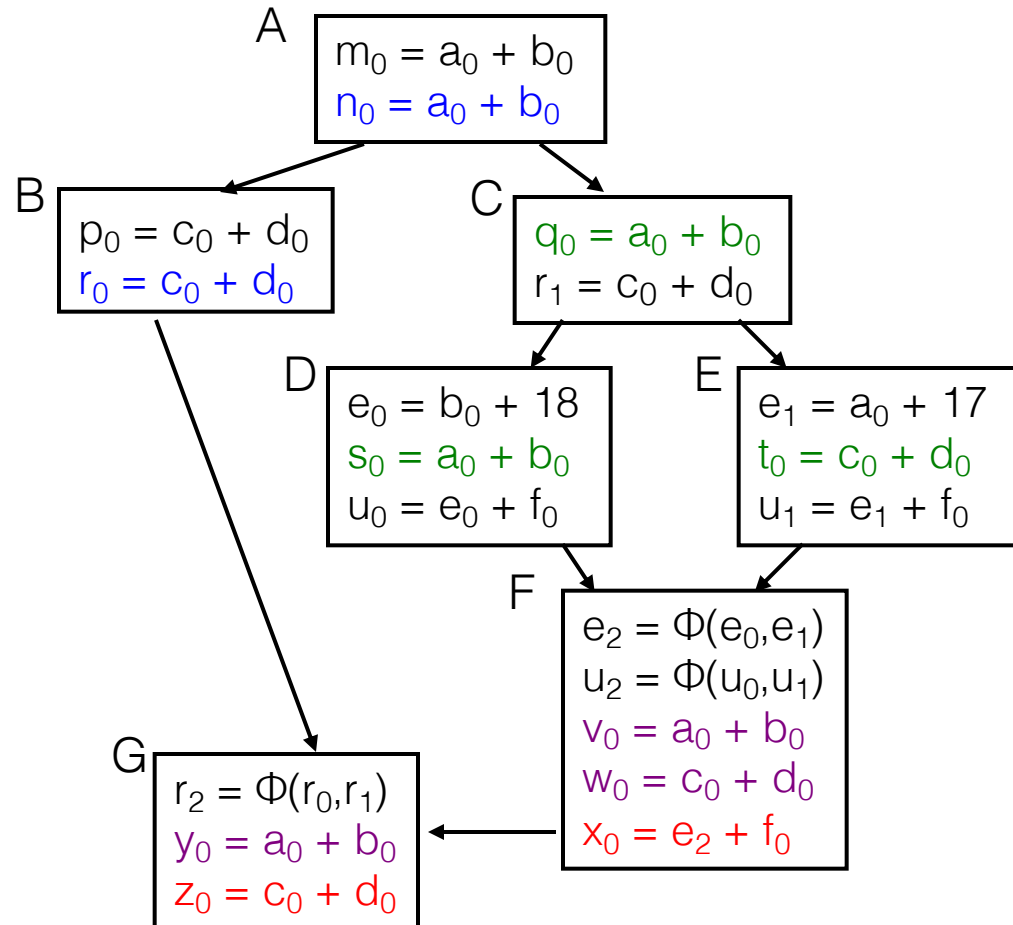


DVNT Algorithm

- Use superlocal algorithm on extended basic blocks
 - Use scoped hash tables & SSA name space as before
- Start each node with table from its IDOM
- No values flow along back edges (i.e., loops)
- Constant folding, algebraic identities as before

Dominator Value Numbering

- Advantages
 - Finds more redundancy
 - Little extra cost
- Shortcomings
 - Misses some opportunities (common calculations in ancestors that are not IDOMs)
 - Doesn't handle loops or other back edges



The Story So Far...


- Local algorithm
- Superlocal extension
 - Some local methods extend cleanly to superlocal scopes
- Dominator VN Technique (DVNT)
- All of these propagate along forward edges
- None are global

Available Expressions

- **Goal:** use dataflow analysis to find common sub-expressions whose range spans basic blocks
- **Idea:** calculate *available expressions* at beginning of each basic block
- Avoid re-evaluation of an available expression – use a copy operation

“Available” and Other Terms

- An expression e is *defined* at point p in the CFG if its value is computed at p
 - Sometimes called *definition site*
- An expression e is *killed* at point p if one of its operands is defined at p
 - Sometimes called *kill site*
- An expression e is *available* at point p if every path leading to p contains a prior definition of e and e is not killed between that definition and p



Available Expression Sets

- To compute available expressions, for each block b , define
 - $AVAIL(b)$ – the set of expressions available on entry to b
 - $NKILL(b)$ – the set of expressions not killed in b
 - i.e., all expressions in the program *except* for those killed in b
 - $DEF(b)$ – the set of expressions defined in b and not subsequently killed in b

Computing Available Expressions

- $AVAIL(b)$ is the set

$$AVAIL(b) = \bigcap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (AVAIL(x) \cap \text{NKILL}(x)))$$

- $\text{preds}(b)$ is the set of b 's predecessors in the CFG
- The set of expressions available on entry to b is the set of expressions that were available at the end of *every* predecessor basic block x
- The expressions available on exit from block b are those defined in b or available on entry to b and not killed in b
- This gives a system of simultaneous equations – a dataflow problem

Name Space Issues

- In previous value-numbering algorithms, we used a SSA-like renaming to keep track of versions
- In global dataflow problems, we use the original namespace
 - we require $a+b$ have the same value along *all* paths to its use
 - If a or b is updated along *any* path to its use, then $a+b$ has the “wrong” value
 - so original names are exactly what we want
- The KILL information captures when a value is no longer available

Computing Available Expressions

- Big Picture

- Build control-flow graph
- Calculate initial local data – $DEF(b)$ and $NKILL(b)$
 - This only needs to be done once for each block b and depends only on the statements in b
- Iteratively calculate $AVAIL(b)$ by repeatedly evaluating equations until nothing changes
 - Another fixed-point algorithm

Computing DEF and NKILL (1)

- For each block b with operations o_1, o_2, \dots, o_k
 $KILLED = \emptyset$ // killed *variables*, not expressions
 $DEF(b) = \emptyset$
 for $i = k$ to 1 // note: working back to front
 assume o_i is “ $x = y + z$ ”
 if ($y \notin KILLED$ and $z \notin KILLED$)
 add “ $y + z$ ” to $DEF(b)$
 add x to $KILLED$
 ...

Computing DEF and NKILL (2)

- After computing DEF and KILLED for a block b , compute set of all expressions in the program not killed in b

$NKILL(b) = \{ \text{all expressions} \}$

for each expression e

for each variable $v \in e$

if $v \in KILLED$ then

$NKILL(b) = NKILL(b) - e$


Data Flow Analysis

Available Expressions

- **Goal:** use dataflow analysis to find common sub-expressions whose range spans basic blocks
- **Idea:** calculate *available expressions* at beginning of each basic block
- Avoid re-evaluation of an available expression – use a copy operation

“Available” and Other Terms

- An expression e is *defined* at point p in the CFG if its value is computed at p
 - Sometimes called *definition site*
- An expression e is *killed* at point p if one of its operands is defined at p
 - Sometimes called *kill site*
- An expression e is *available* at point p if every path leading to p contains a prior definition of e and e is not killed between that definition and p



Available Expression Sets

- To compute available expressions, for each block b , define
 - $AVAIL(b)$ – the set of expressions available on entry to b
 - $NKILL(b)$ – the set of expressions not killed in b
 - i.e., all expressions in the program *except* for those killed in b
 - $DEF(b)$ – the set of expressions defined in b and not subsequently killed in b

Computing Available Expressions

- $AVAIL(b)$ is the set

$$AVAIL(b) = \bigcap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (AVAIL(x) \cap \text{NKILL}(x)))$$

- $\text{preds}(b)$ is the set of b 's predecessors in the CFG
- The set of expressions available on entry to b is the set of expressions that were available at the end of *every* predecessor basic block x
- The expressions available on exit from block b are those defined in b or available on entry to b and not killed in b
- This gives a system of simultaneous equations – a dataflow problem

Name Space Issues

- In previous value-numbering algorithms, we used a SSA-like renaming to keep track of versions
- In global dataflow problems, we use the original namespace
 - we require $a+b$ have the same value along *all* paths to its use
 - If a or b is updated along *any* path to its use, then $a+b$ has the “wrong” value
 - so original names are exactly what we want
- The KILL information captures when a value is no longer available

Computing Available Expressions

- Big Picture

- Build control-flow graph
- Calculate initial local data – $DEF(b)$ and $NKILL(b)$
 - This only needs to be done once for each block b and depends only on the statements in b
- Iteratively calculate $AVAIL(b)$ by repeatedly evaluating equations until nothing changes
 - Another fixed-point algorithm

Computing DEF and NKILL (1)

- For each block b with operations o_1, o_2, \dots, o_k
 $KILLED = \emptyset$ // killed *variables*, not expressions
 $DEF(b) = \emptyset$
 for $i = k$ to 1 // note: working back to front
 assume o_i is “ $x = y + z$ ”
 if ($y \notin KILLED$ and $z \notin KILLED$)
 add “ $y + z$ ” to $DEF(b)$
 add x to $KILLED$
 ...

Computing DEF and NKILL (2)

- After computing DEF and KILLED for a block b , compute set of all expressions in the program not killed in b

$NKILL(b) = \{ \text{all expressions} \}$

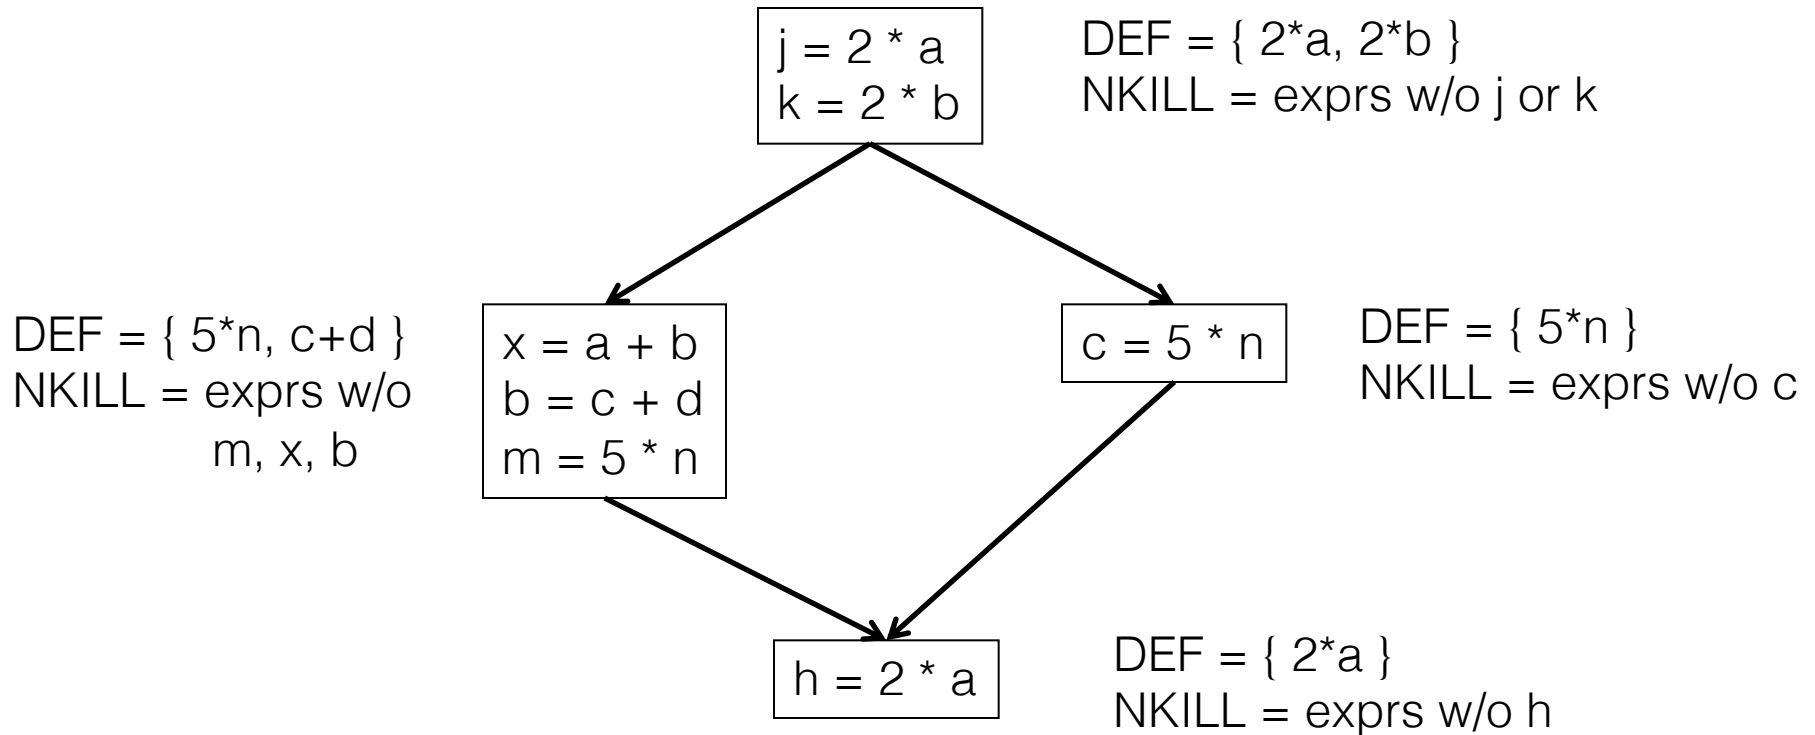
for each expression e

for each variable $v \in e$

if $v \in KILLED$ then

$NKILL(b) = NKILL(b) - e$

Example: Compute DEF and NKILL



Computing Available Expressions

Once $DEF(b)$ and $NKILL(b)$ are computed for all blocks b

Worklist = { all blocks b_i }

while (Worklist $\neq \emptyset$)

 remove a block b from Worklist

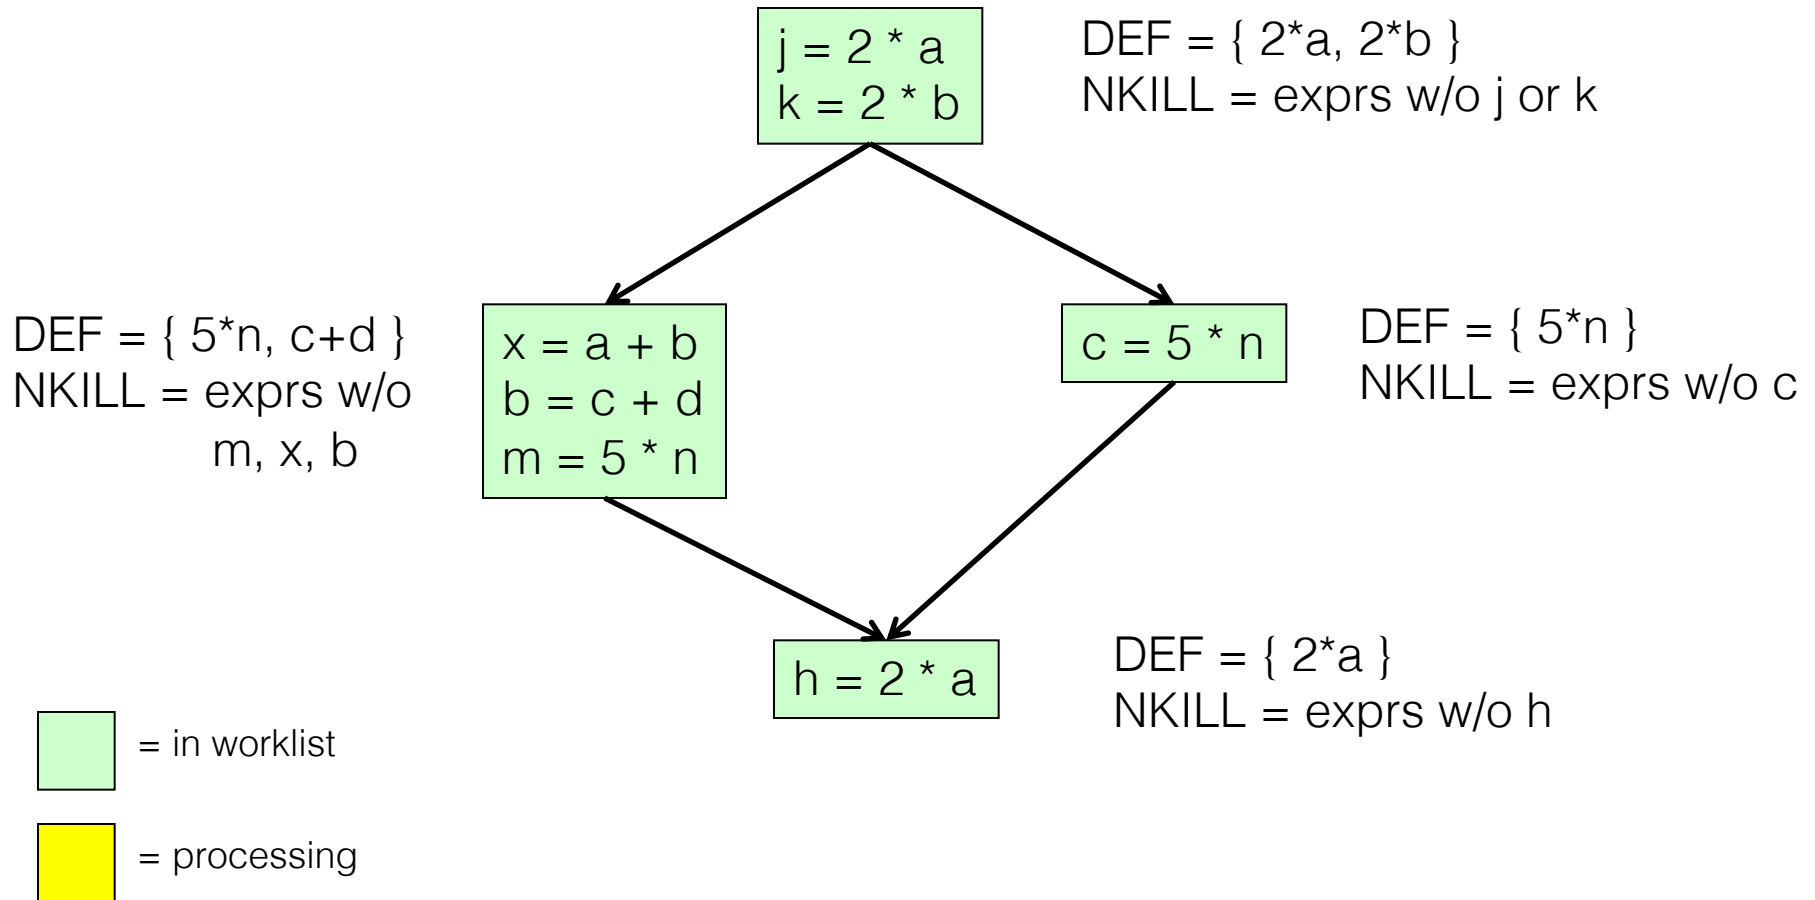
 recompute $AVAIL(b)$

 if $AVAIL(b)$ changed

 Worklist = Worklist \cup successors(b)

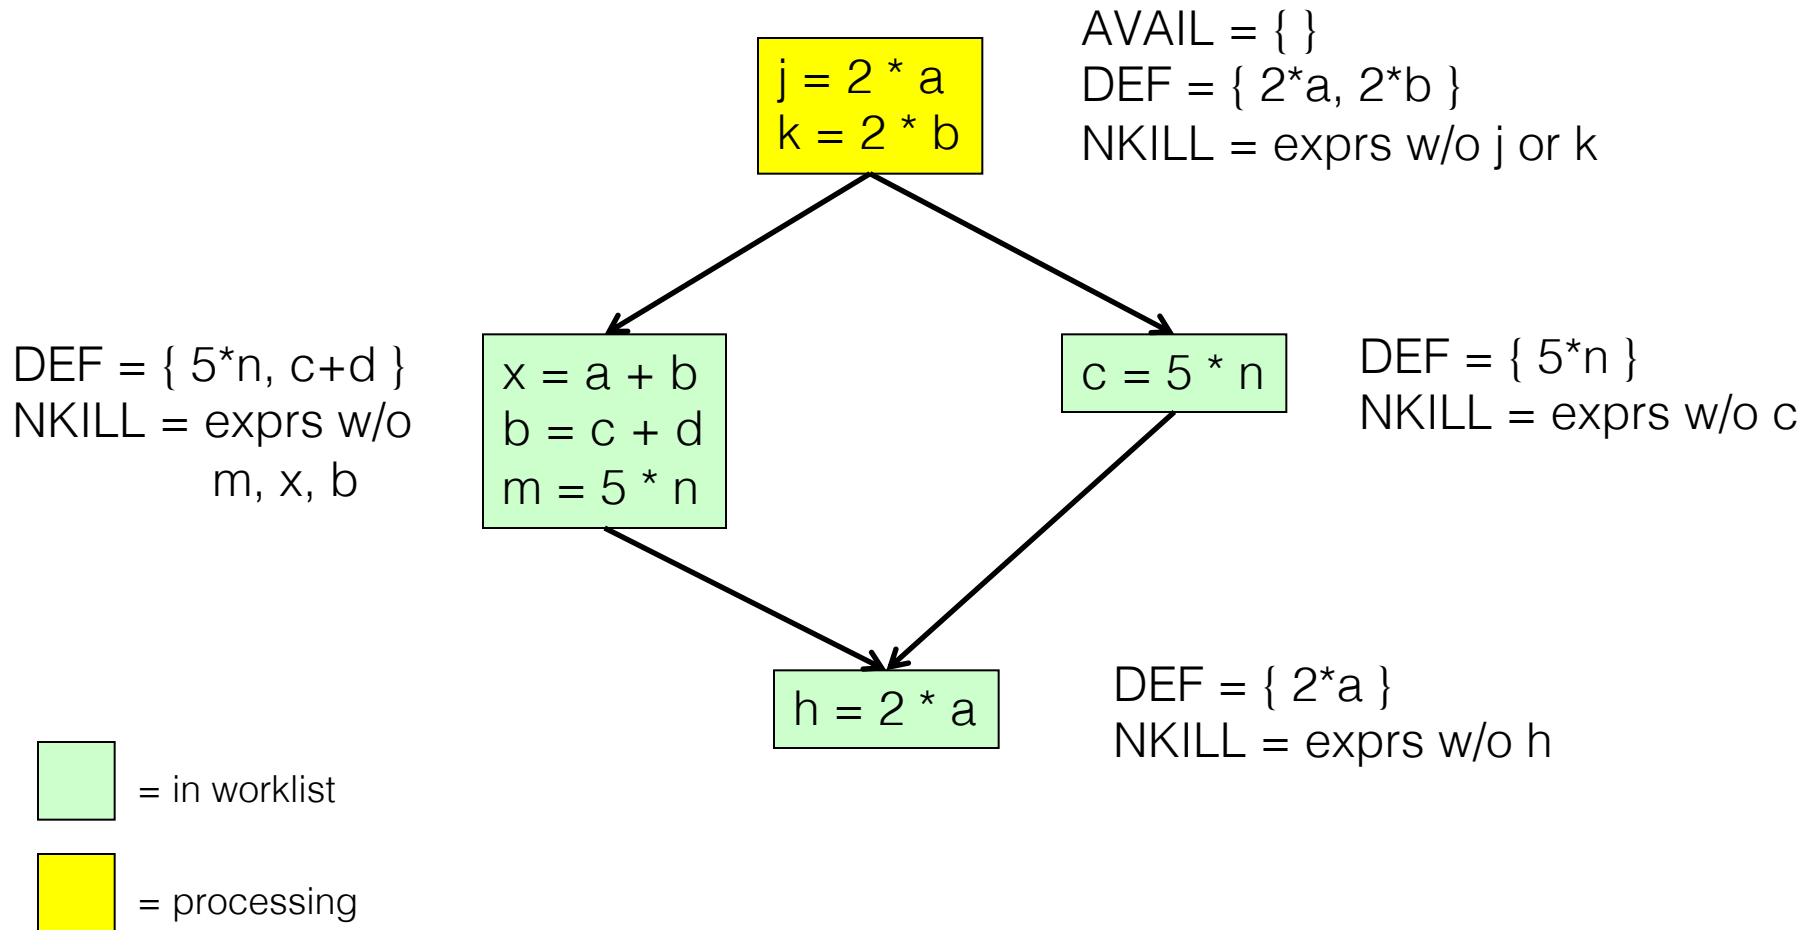
Example: Find Available Expressions

$$AVAIL(b) = \cap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (\text{AVAIL}(x) \cap \text{NKILL}(x)))$$



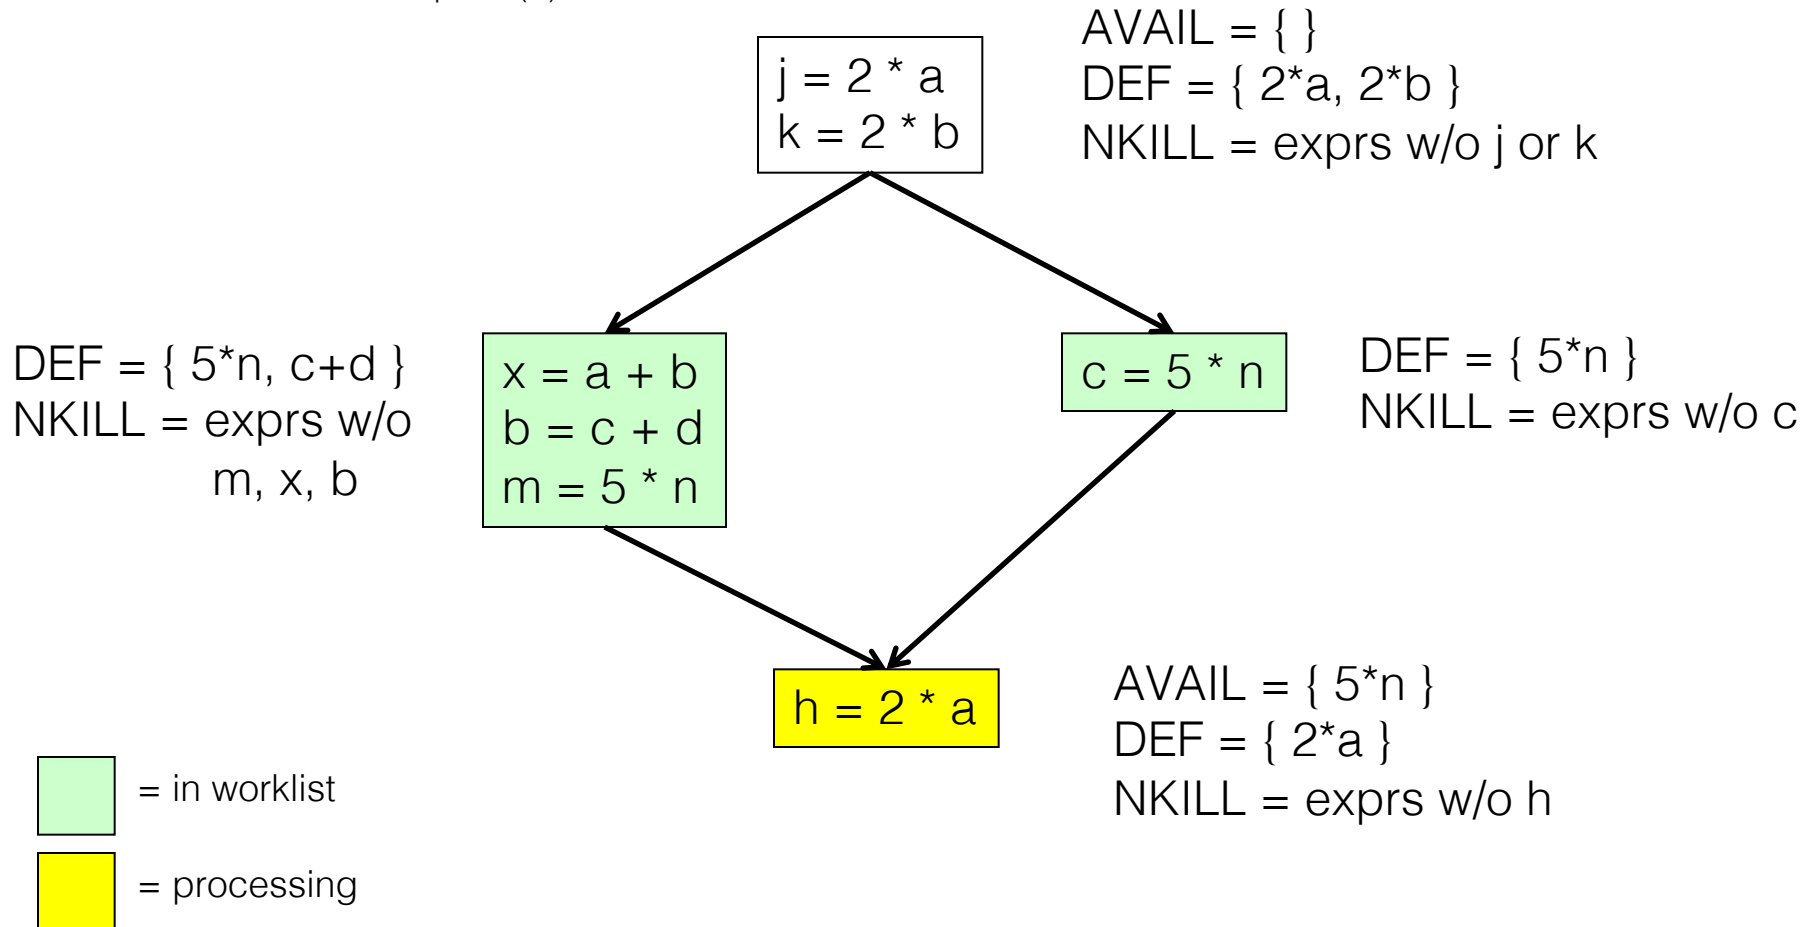
Example: Find Available Expressions

$$AVAIL(b) = \cap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (\text{AVAIL}(x) \cap \text{NKILL}(x)))$$




Example: Find Available Expressions

$$AVAIL(b) = \cap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (\text{AVAIL}(x) \cap \text{NKILL}(x)))$$




Example: Find Available Expressions

$$AVAIL(b) = \cap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (\text{AVAIL}(x) \cap \text{NKILL}(x)))$$



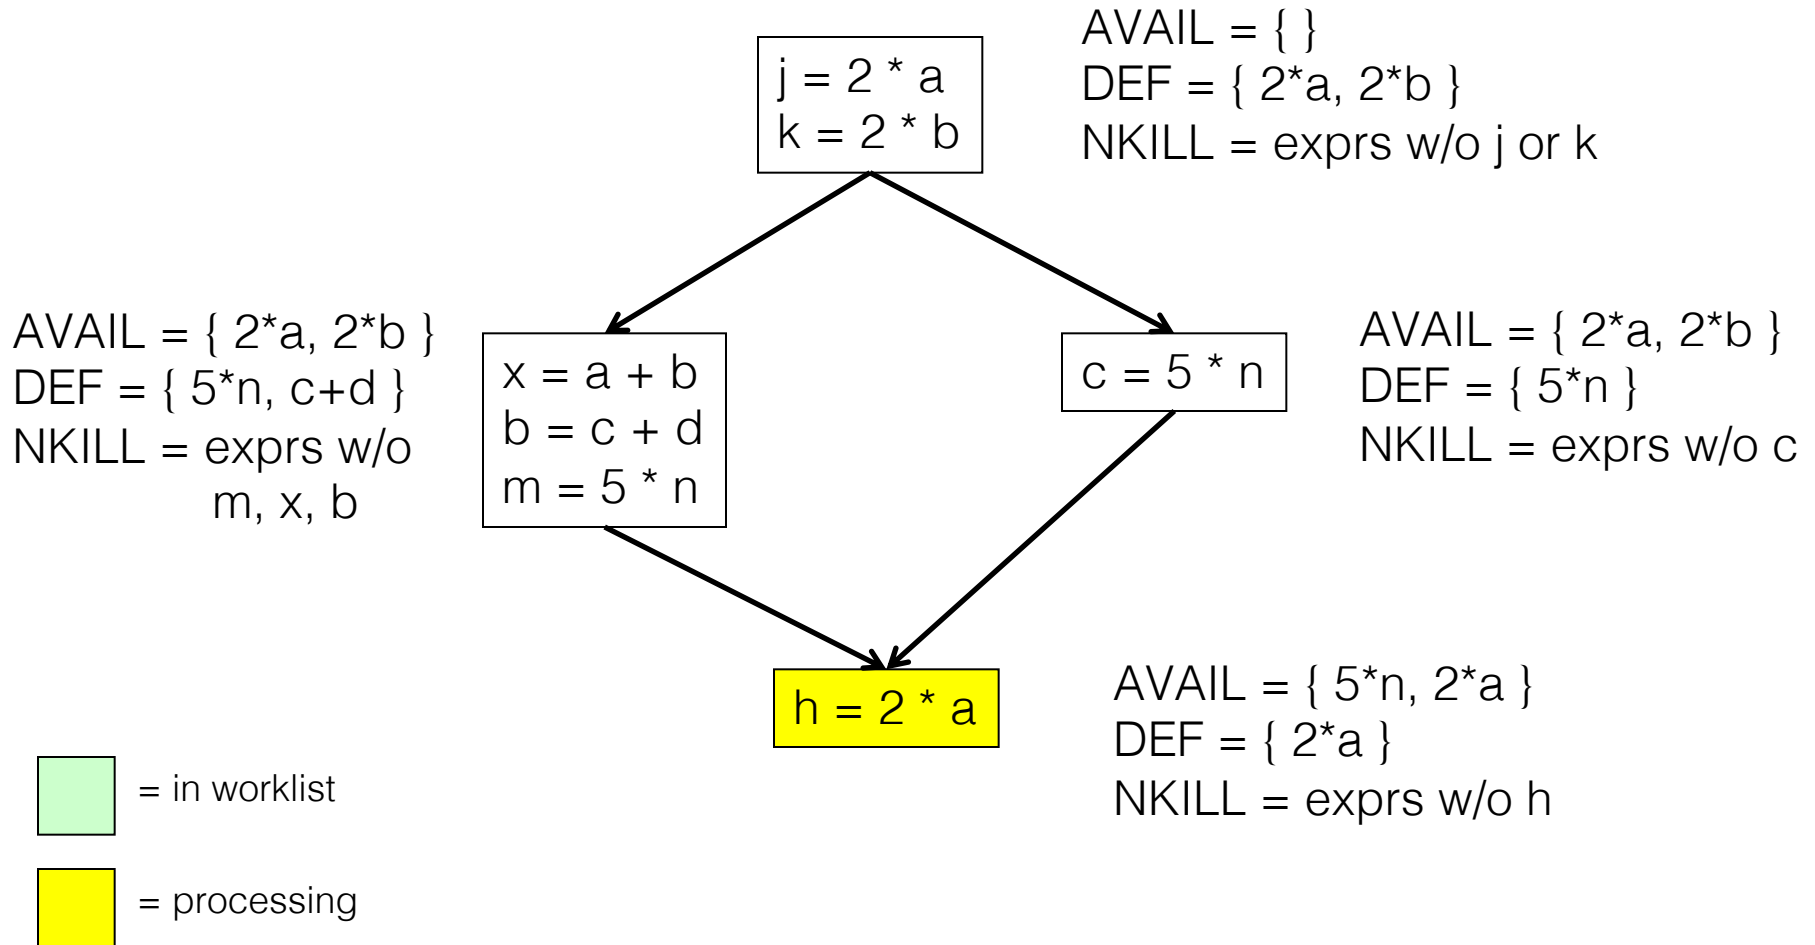
Example: Find Available Expressions

$$AVAIL(b) = \cap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (\text{AVAIL}(x) \cap \text{NKILL}(x)))$$



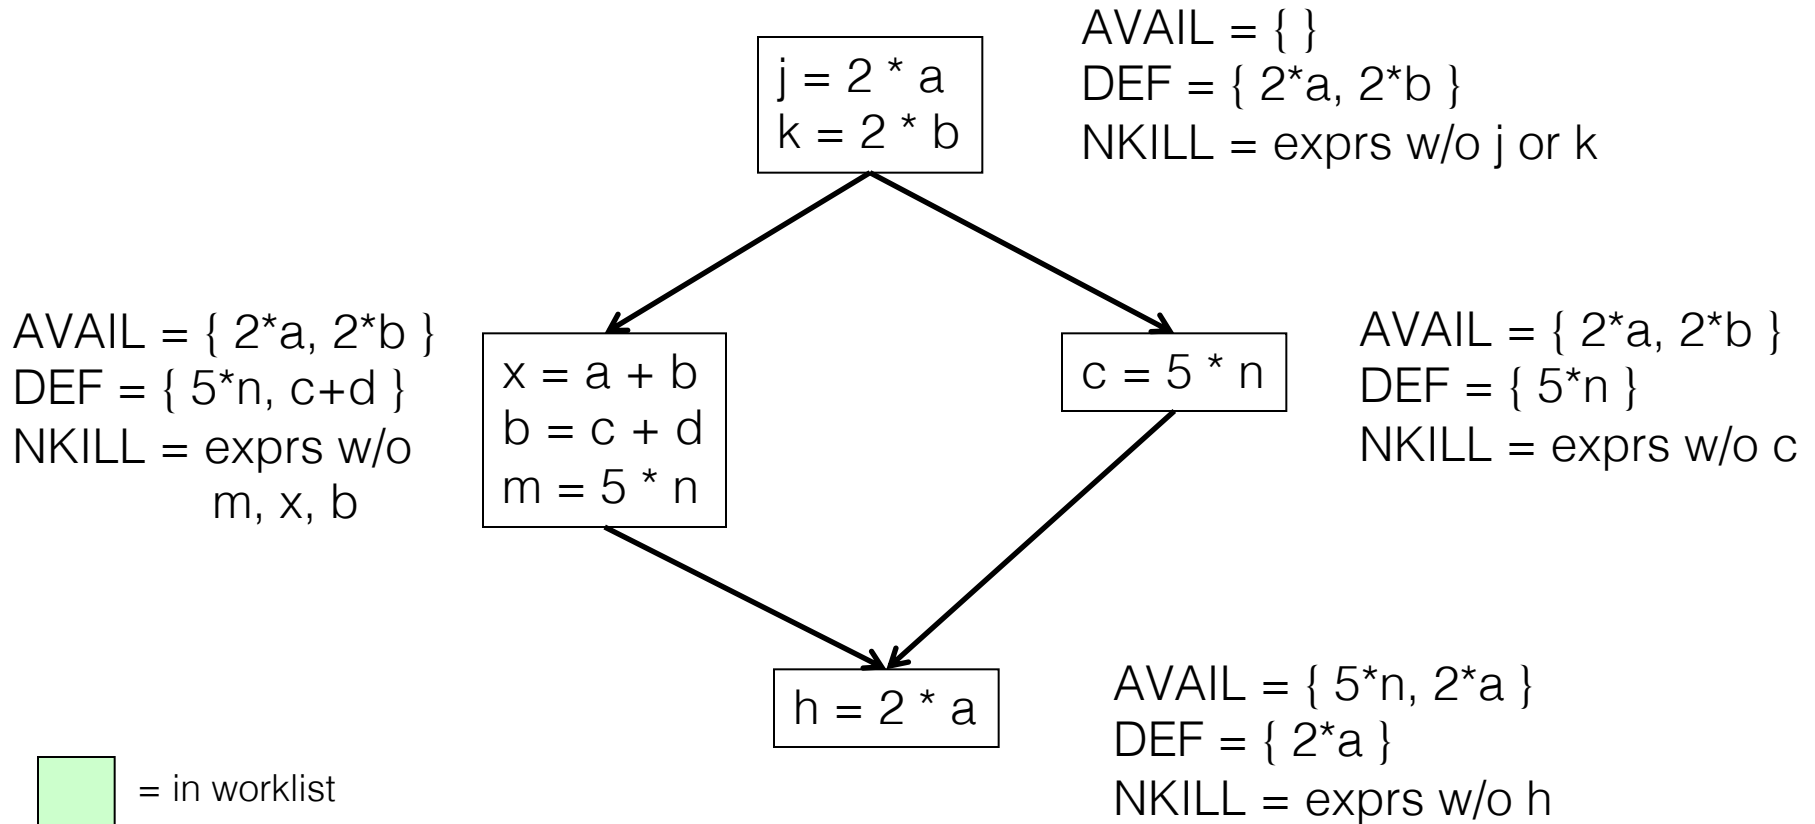
Example: Find Available Expressions

$$AVAIL(b) = \cap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (\text{AVAIL}(x) \cap \text{NKILL}(x)))$$



Example: Find Available Expressions

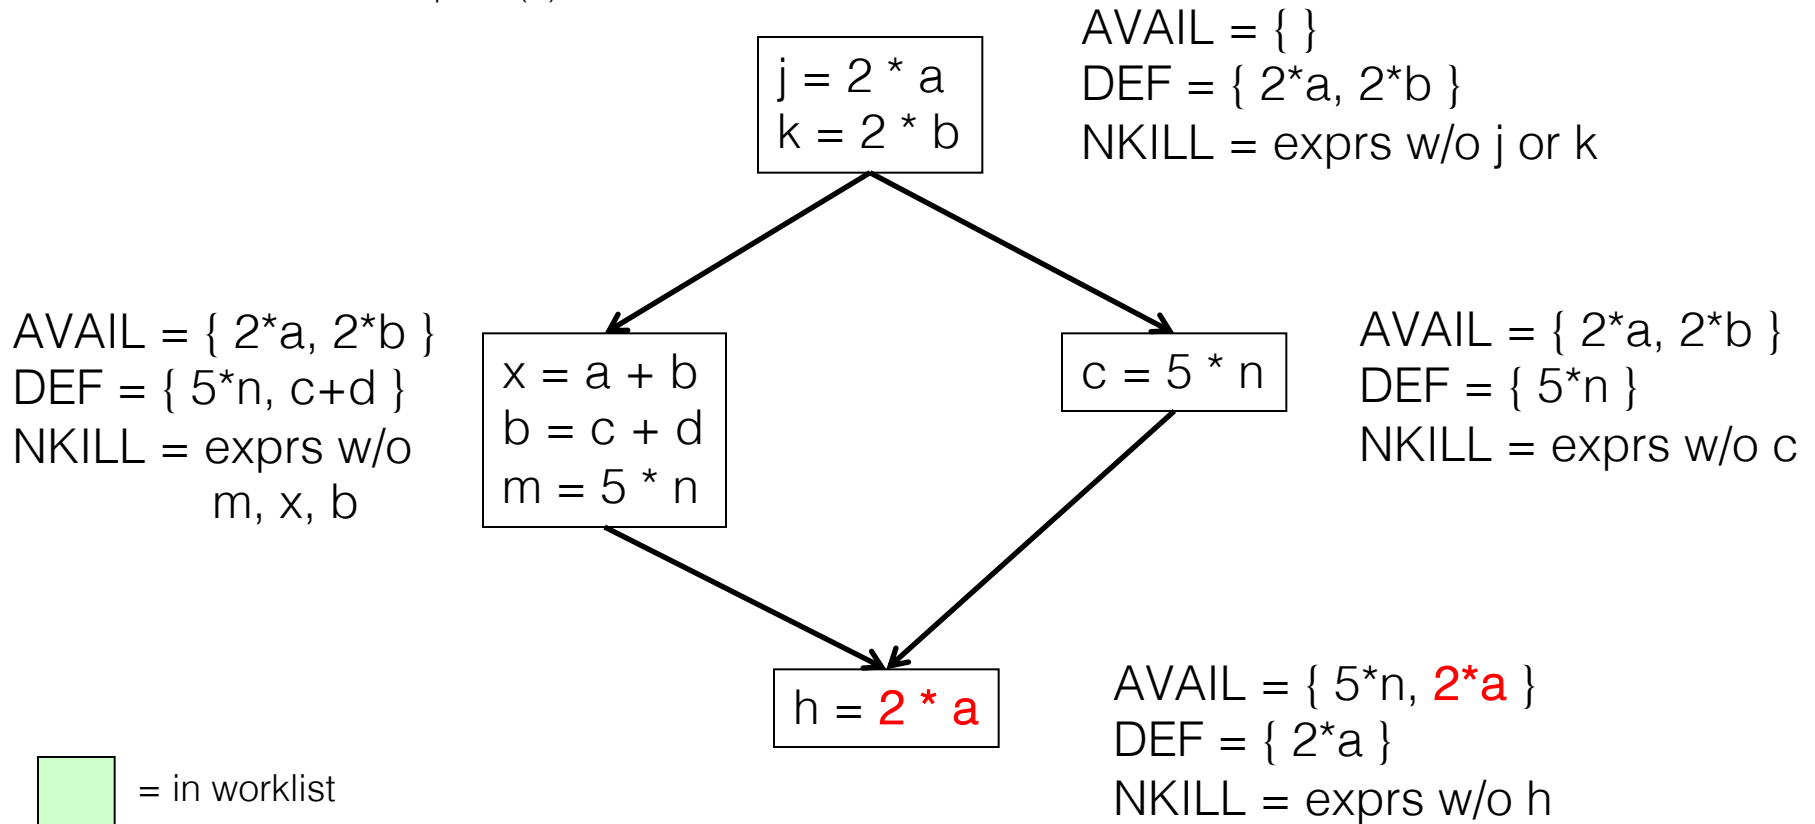
$$AVAIL(b) = \cap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (\text{AVAIL}(x) \cap \text{NKILL}(x)))$$



And the common subexpression is???


Example: Find Available Expressions

$$AVAIL(b) = \cap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (\text{AVAIL}(x) \cap \text{NKILL}(x)))$$



Comparing Algorithms

- LVN – Local Value Numbering
- SVN – Superlocal Value Numbering
- DVN – DominatoT-based Value Numbering
- GRE – Global Redundancy Elimination



Comparing Algorithms (2)

- $LVN \Rightarrow SVN \Rightarrow DVN$ form a strict hierarchy – later algorithms find a superset of previous information
- Global RE finds a somewhat different set
 - Discovers $e+f$ in F (computed in both D and E)
 - Misses identical values if they have different names (e.g.,
 $a+b$ and $c+d$ when $a=c$ and $b=d$)
 - Value Numbering catches this

Scope of Analysis

- Larger context (EBBs, regions, global, interprocedural) sometimes helps
 - More opportunities for optimizations
- But not always
 - Introduces uncertainties about flow of control
 - Usually only allows weaker analysis
 - Sometimes has unwanted side effects
 - Can create additional pressure on registers, for example


Code Replication

- Sometimes replicating code increases opportunities – modify the code to create larger regions with simple control flow
- Two examples
 - Cloning
 - Inline substitution

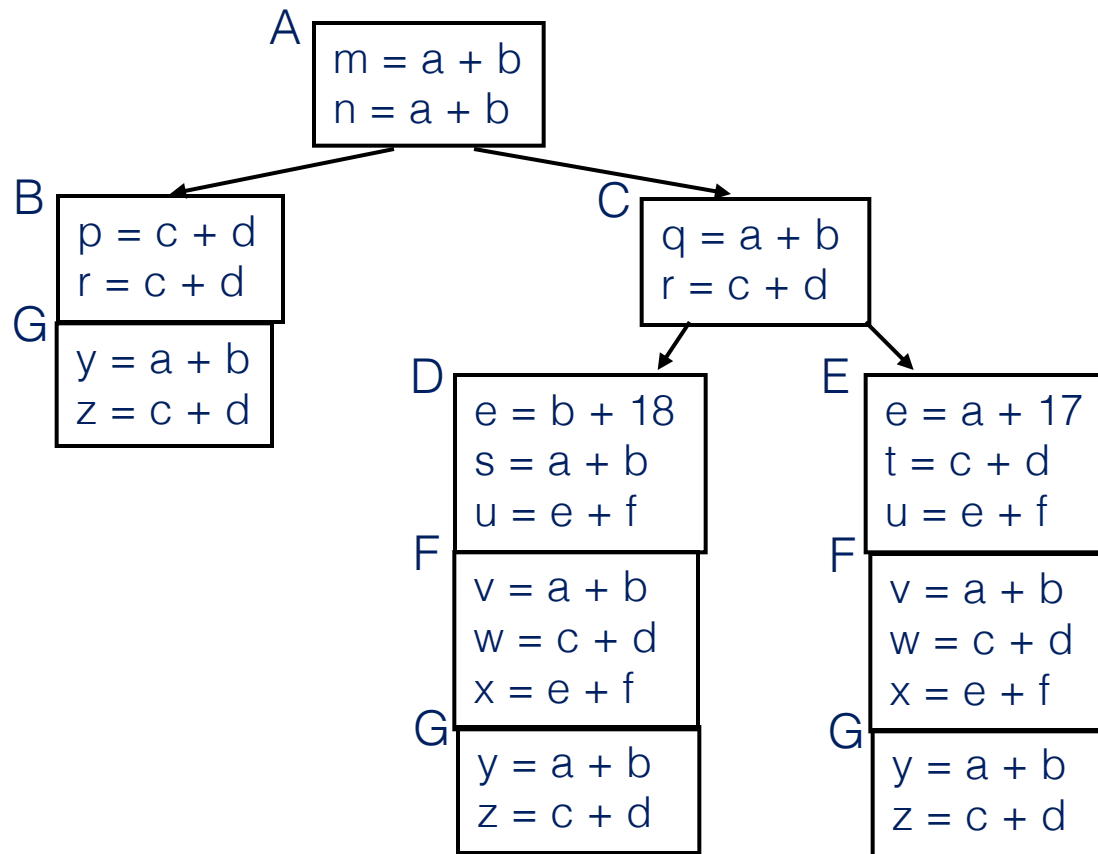
Cloning

- **Idea:** duplicate blocks with multiple predecessors
- **Tradeoff**
 - More local optimization possibilities – larger blocks, fewer branches
 - But: larger code size, may slow down if it interacts badly with cache

Original VN Example



Example with Cloning



Inline Substitution

- **Problem:** an optimizer has to treat a procedure call as if it (could have) modified all globally reachable data
 - Plus there is the basic expense of calling the procedure
- **Inline Substitution:** replace each call site with a copy of the called function body

Inline Substitution Issues

- Pro
 - More effective optimization – better local context and don't need to invalidate local assumptions
 - Eliminate overhead of normal function call
- Con
 - Potential code bloat
 - Need to manage recompilation when either caller or callee changes

Dataflow Analysis

- Available expressions are an example of a *dataflow analysis* problem
- Many similar problems can be expressed in a similar framework
- Only the first part of the story – once we've discovered facts, we then need to use them to improve code

Characterizing Dataflow Analysis

- All of these algorithms involve sets of facts about each basic block b
 - $IN(b)$ – facts true on entry to b
 - $OUT(b)$ – facts true on exit from b
 - $GEN(b)$ – facts created and not killed in b
 - $KILL(b)$ – facts killed in b
- These are related by the equation
$$OUT(b) = GEN(b) \cup (IN(b) - KILL(b))$$
 - Solve this iteratively for all blocks
 - Sometimes information propagates forward; sometimes backward

Dataflow Analysis (1)

- A collection of techniques for compile-time reasoning about run-time values
- Almost always involves building a graph
 - Trivial for basic blocks
 - Control-flow graph or derivative for global problems
 - Call graph or derivative for whole-program problems

Dataflow Analysis (2)

- Usually formulated as a set of *simultaneous equations* (dataflow problem)
 - Sets attached to nodes and edges
 - Need a lattice (or semilattice) to describe values
 - In particular, has an appropriate operator to combine values and an appropriate “bottom” or minimal value

Dataflow Analysis (3)

- Desired solution is usually a *meet over all paths* (MOP) solution
 - “What is true on every path from entry”
 - “What can happen on any path from entry”
 - Usually relates to safety of optimization

Dataflow Analysis (4)

- Limitations
 - Precision – “up to symbolic execution”
 - Assumes all paths taken
 - Sometimes cannot afford to compute full solution
 - Arrays – classic analysis treats each array as a single fact
 - Pointers – difficult, expensive to analyze
 - Imprecision rapidly adds up
 - But gotta do it to effectively optimize things like C/C++
- For scalar values we can quickly solve simple problems

Example: Live Variable Analysis

- A variable v is *live* at point p iff there is *any* path from p to a use of v along which v is not redefined
- Some uses:
 - Register allocation – only live variables need a register
 - Eliminating useless stores – if variable not live at store, then stored variable will never be used
 - Detecting uses of uninitialized variables – if live at declaration (before initialization) then it might be used uninitialized
 - Improve SSA construction – only need Φ -function for variables that are live in a block (later)

Liveness Analysis Sets

- For each block b , define
 - $\text{use}[b]$ = variable used in b before any def
 - $\text{def}[b]$ = variable defined in b & not killed
 - $\text{in}[b]$ = variables live on entry to b
 - $\text{out}[b]$ = variables live on exit from b

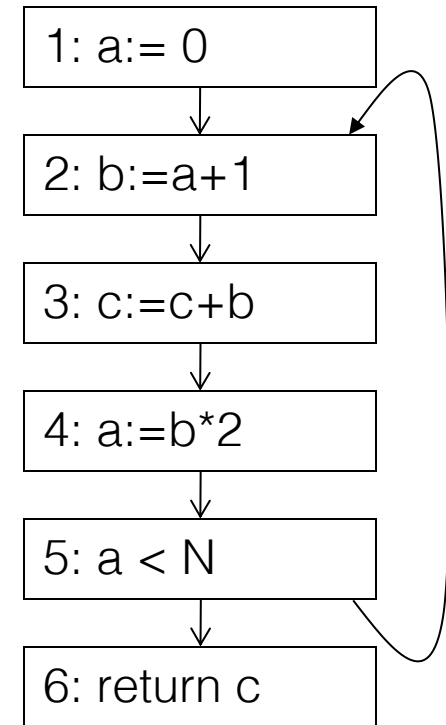
Equations for Live Variables

- Given the preceding definitions, we have
$$\text{in}[b] = \text{use}[b] \cup (\text{out}[b] - \text{def}[b])$$
$$\text{out}[b] = \bigcup_{s \in \text{succ}[b]} \text{in}[s]$$
- Algorithm
 - Set $\text{in}[b] = \text{out}[b] = \emptyset$
 - Update in, out until no change

Example (1 stmt per block)

- Code

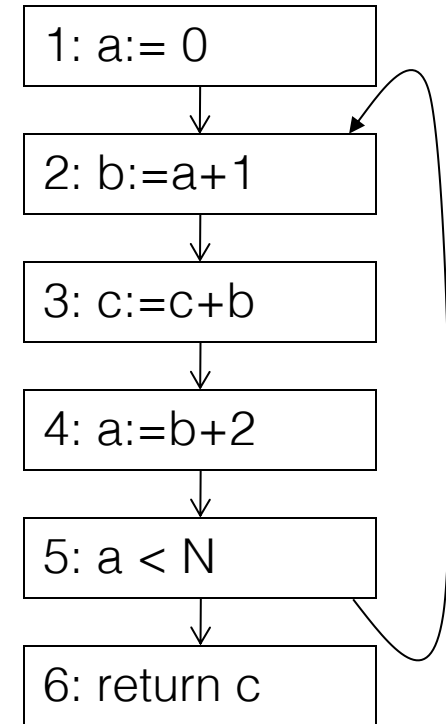
```
a := 0  
L: b := a+1  
  c := c+b  
  a := b*2  
  if a < N goto L  
  return c
```



$$\begin{aligned} \text{in}[b] &= \text{use}[b] \cup (\text{out}[b] - \text{def}[b]) \\ \text{out}[b] &= \bigcup_{s \in \text{succ}[b]} \text{in}[s] \end{aligned}$$

Calculation

	I			II			III	
block	use	def	out	in	out	in	out	in
6								
5								
4								
3								
2								
1								

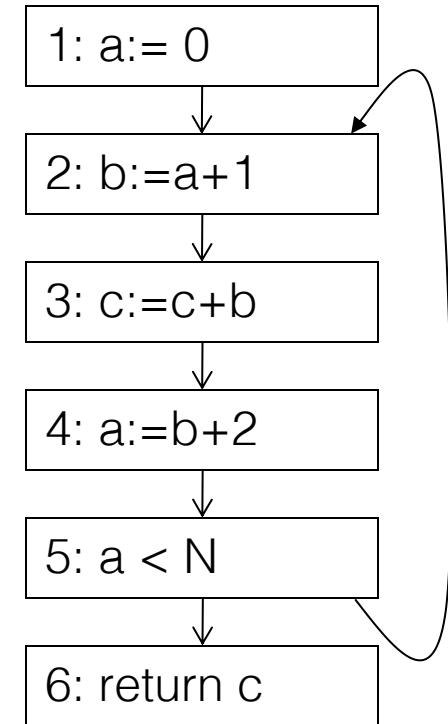


$$\text{in}[b] = \text{use}[b] \cup (\text{out}[b] - \text{def}[b])$$

$$\text{out}[b] = \bigcup_{s \in \text{succ}[b]} \text{in}[s]$$

Calculation

	I				II		III	
block	use	def	out	in	out	in	out	in
6	c	--	--	c	--	c		
5	a	--	c	a,c	a,c	a,c		
4	b	a	a,c	b,c	a,c	b,c		
3	b,c	c	b,c	b,c	b,c	b,c		
2	a	b	b,c	a,c	b,c	a,c		
1	--	a	a,c	c	a,c	c		



$$\text{in}[b] = \text{use}[b] \cup (\text{out}[b] - \text{def}[b])$$

$$\text{out}[b] = \bigcup_{s \in \text{succ}[b]} \text{in}[s]$$

Equations for Live Variables v2

- Many problems have more than one formulation. For example, Live Variables...
- Sets
 - USED(b) – variables used in b before being defined in b
 - NOTDEF(b) – variables not defined in b
 - LIVE(b) – variables live on *exit* from b
- Equation
$$\text{LIVE}(b) = \bigcup_{s \in \text{succ}(b)} \text{USED}(s) \cup (\text{LIVE}(s) \cap \text{NOTDEF}(s))$$

Efficiency of Dataflow Analysis

- The algorithms eventually terminate, but the expected time needed can be reduced by picking a good order to visit nodes in the CFG
 - Forward problems – reverse postorder
 - Backward problems – postorder

Example: Reaching Definitions

- A definition d of some variable v *reaches* operation i iff i reads the value of v and there is a path from d to i that does not define v
- Uses
 - Find all of the possible definition points for a variable in an expression

Equations for Reaching Definitions

- Sets
 - DEFOUT(b) – set of definitions in b that reach the end of b (i.e., not subsequently redefined in b)
 - SURVIVED(b) – set of all definitions not obscured by a definition in b
 - REACHES(b) – set of definitions that reach b
- Equation
$$\text{REACHES}(b) = \cup_{p \in \text{preds}(b)} \text{DEFOUT}(p) \cup (\text{REACHES}(p) \cap \text{SURVIVED}(p))$$

Example: Very Busy Expressions

- An expression e is considered *very busy* at some point p if e is evaluated and used along every path that leaves p , and evaluating e at p would produce the same result as evaluating it at the original locations
- Uses
 - Code hoisting – move e to p (reduces code size; no effect on execution time)

Equations for Very Busy Expressions

- Sets
 - USED(b) – expressions used in b before they are killed
 - KILLED(b) – expressions redefined in b before they are used
 - VERYBUSY(b) – expressions very busy on exit from b
- Equation
$$\text{VERYBUSY}(b) = \bigcap_{s \in \text{succ}(b)} \text{USED}(s) \cup (\text{VERYBUSY}(s) - \text{KILLED}(s))$$

Using Dataflow Information

- A few examples of possible transformations...

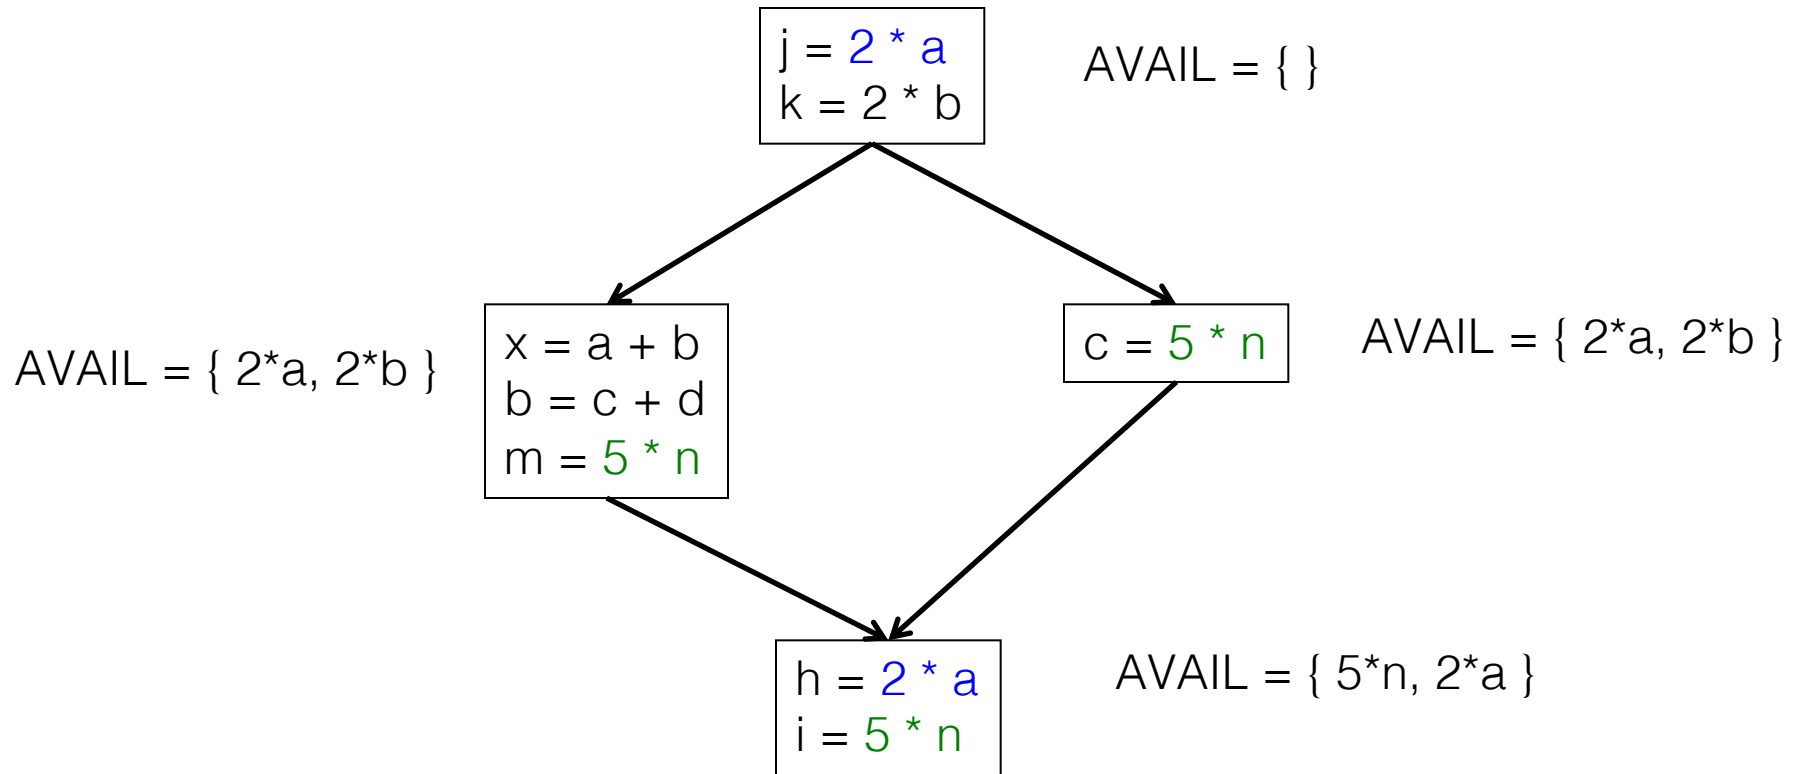
Classic Common-Subexpression Elimination (CSE)

- In a statement $s: t := x \text{ op } y$, if $x \text{ op } y$ is *available* at s then it need not be recomputed
- Analysis: compute *reaching expressions* i.e., statements $n: v := x \text{ op } y$ such that the path from n to s does not compute $x \text{ op } y$ or define x or y

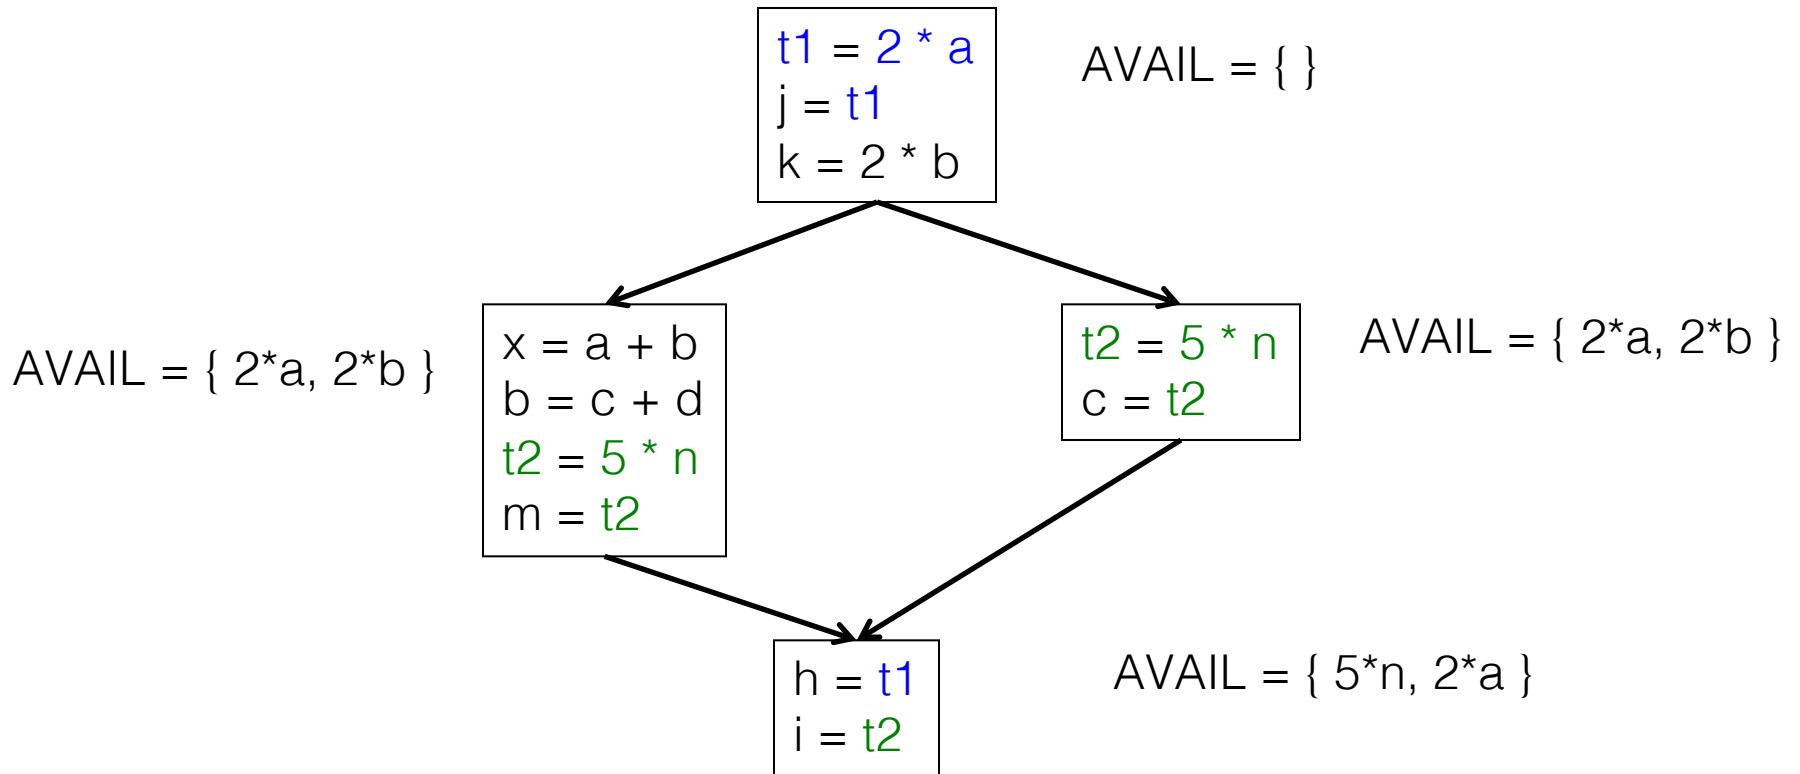
Classic CSE Transformation

- If $x \text{ op } y$ is defined at n and reaches s
 - Create new temporary w
 - Rewrite $n: v := x \text{ op } y$ as
$$\begin{array}{l} n: w := x \text{ op } y \\ n': v := w \end{array}$$
 - Modify statement s to be
$$s: t := w$$
 - (Rely on copy propagation to remove extra assignments that are not really needed)

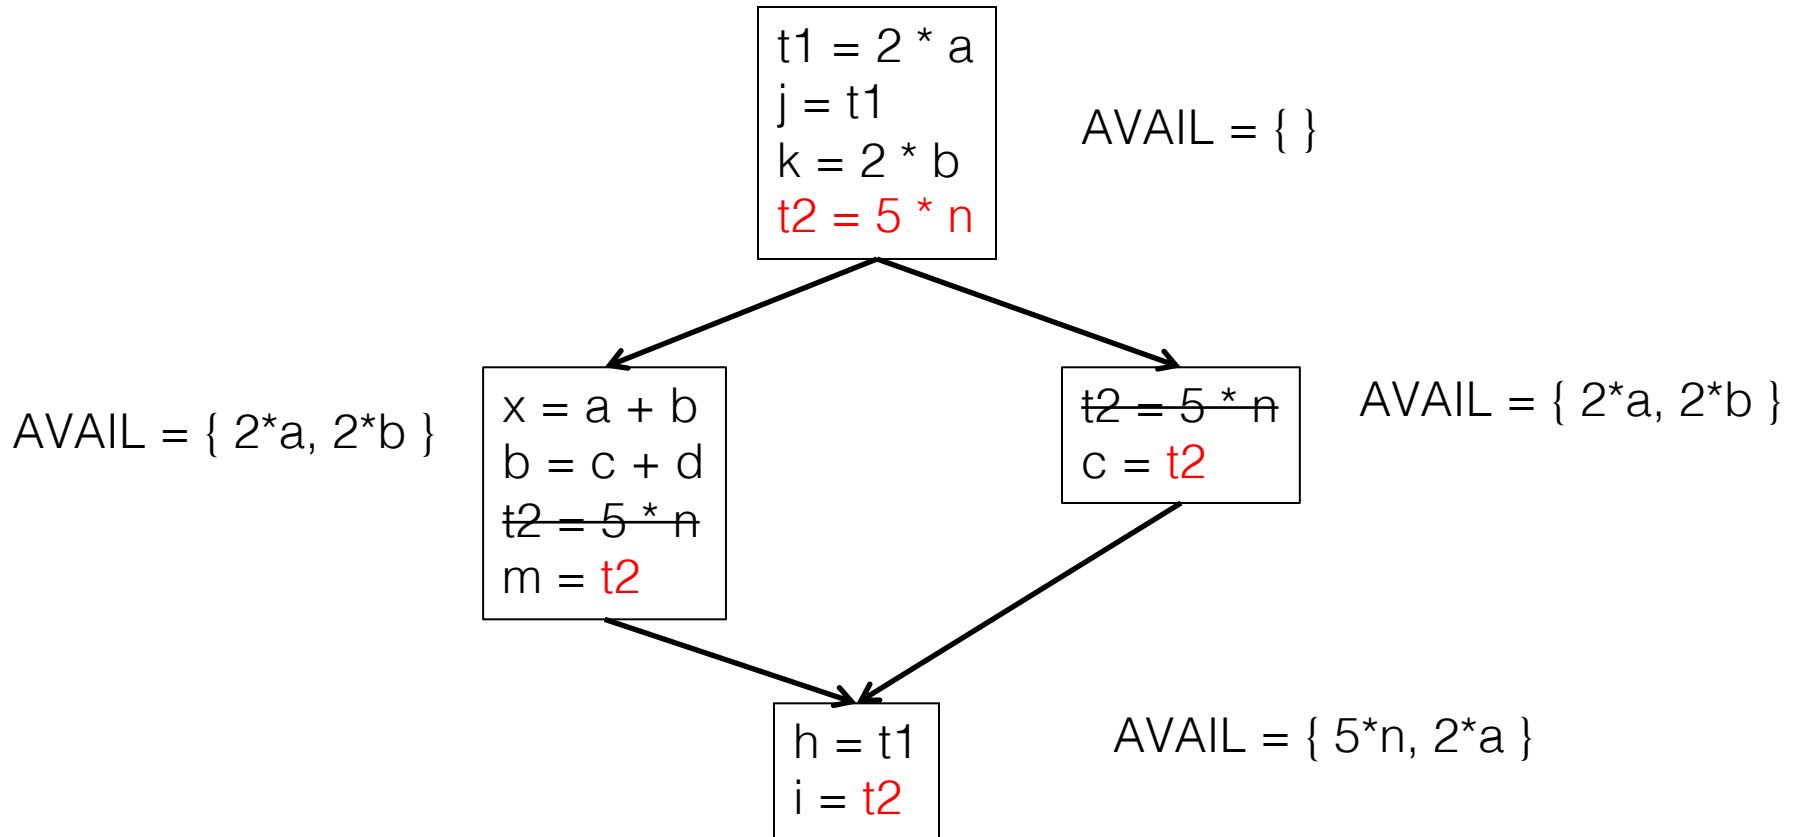
Revisiting Example (w/slight addition)



Revisiting Example (w/slight addition)



Then Apply Very Busy...



Constant Propagation

- Suppose we have
 - Statement d: $t := c$, where c is constant
 - Statement n that uses t
- If d reaches n and no other definitions of t reach n, then rewrite n to use c instead of t

Copy Propagation

- Similar to constant propagation
- Setup:
 - Statement d: $t := z$
 - Statement n uses t
- If d reaches n and no other definition of t reaches n, and there is no definition of z on any path from d to n, then rewrite n to use z instead of t
 - Recall that this can help remove dead assignments

Copy Propagation Tradeoffs

- Downside is that this can increase the lifetime of variable z and increase need for registers or memory traffic
- But it can expose other optimizations, e.g.,
 - $a := y + z$
 - $u := y$
 - $c := u + z$ // copy propagation makes this $y + z$
- After copy propagation we can recognize the common subexpression

Dead Code Elimination

- If we have an instruction
 $s: a := b \text{ op } c$
and a is not live-out after s , then s can be eliminated
 - Provided it has no implicit side effects that are visible (output, exceptions, etc.)
 - If b or c are function calls, they have to be assumed to have unknown side effects unless the compiler can prove otherwise

Aliases

- A variable or memory location may have multiple names or *aliases*
 - Call-by-reference parameters
 - Variables whose address is taken (&x)
 - Expressions that dereference pointers (p.x, *p)
 - Expressions involving subscripts (a[i])
 - Variables in nested scopes

Aliases vs Optimizations

- Example:

`p.x := 5; q.x := 7; a := p.x;`

- Does reaching definition analysis show that the definition of `p.x` reaches `a`?
- (Or: do `p` and `q` refer to the same variable/object?)
- (Or: *can* `p` and `q` refer to the same thing?)

Aliases vs Optimizations

- Example

```
void f(int *p, int *q) {  
    *p = 1; *q = 2;  
    return *p;  
}
```

- How do we account for the possibility that p and q might refer to the same thing?
- Safe approximation: since it's possible, assume it is true (but rules out a lot)
 - C programmers can use “restrict” to indicate no other pointer is an alias for this one

Types and Aliases (1)

- In Java, ML, MiniJava, and others, if two variables have incompatible types they cannot be names for the same location
 - Also helps that programmer cannot create arbitrary pointers to storage in these languages

Types and Aliases (2)

- Strategy: Divide memory locations into *alias classes* based on type information (every type, array, record field is a class)
- Implication: need to propagate type information from the semantics pass to optimizer
 - Not normally true of a minimally typed IR
- Items in different alias classes cannot refer to each other

Aliases and Flow Analysis

- Idea: Base alias classes on points where a value is created
 - Every new/malloc and each local or global variable whose address is taken is an alias class
 - Pointers can refer to values in multiple alias classes (so each memory reference is to a set of alias classes)
 - Use to calculate “may alias” information (e.g., p “may alias” q at program point s)

Using “may-alias” information

- Treat each alias class as a “variable” in dataflow analysis problems
- Example: framework for available expressions
 - Given statement $s: M[a] := b$,
 $gen[s] = \{ \}$
 $kill[s] = \{ M[x] \mid a \text{ may alias } x \text{ at } s \}$

May-Alias Analysis

- Without alias analysis, #2 kills $M[t]$ since x and t might be related
- If analysis determines that “ x may-alias t ” is false, $M[t]$ is still available at #3; can eliminate the common subexpression and use copy propagation
- Code
 - 1: $u := M[t]$
 - 2: $M[x] := r$
 - 3: $w := M[t]$
 - 4: $b := u + w$

Coming Attractions

- Dataflow analysis is the core of classical optimizations
 - Although not the only possible story
- Still to explore:
 - Discovering and optimizing loops
 - SSA – Static Single Assignment form

SSA Name Space

- Two Principles
 - Each name is defined by exactly one operation
 - Each operand refers to exactly one definition
- Need to deal with merge points
 - Add Φ functions at merge points to reconcile names
 - Use subscripts on variable names for uniqueness



[Meme credit: imgflip.com]