

CS 6410: Compilers

Fall 2023

Tamara Bonaci
t.bonaci@northeastern.edu

Thank you to UW faculty Hal Perkins. Today lecture notes are a modified version of his lecture notes.

Credits For Course Material

- Big thank you to UW CSE faculty member, Hal Perkins
- Some direct ancestors of this course:
 - UW CSE 401 (Chambers, Snyder, Notkin, Perkins, Ringenburt, Henry, ...)
 - UW CSE PMP 582/501 (Perkins)
 - Cornell CS 412-3 (Teitelbaum, Perkins)
 - Rice CS 412 (Cooper, Kennedy, Torczon)
 - Many books (Appel; Cooper/Torczon; Aho, [[Lam,] Sethi,] Ullman [Dragon Book], Fischer, [Cytron ,] LeBlanc; Muchnick, ...)

Agenda

- Dataflow analysis – review and finish
 - Framework for many common compiler analyses
 - Dataflow analysis for common subexpression elimination
 - Other analysis problems that work in the same framework
 - Some of these are optimizations we've seen, but more formally and with details
- Loops optimizations
 - Dominators – discovering loops
 - Loop invariant calculations
 - Loop transformations
 - A quick look at some memory hierarchy issues
 - Largely based on material in Appel ch. 18, 21; similar material in other books
- Overview of SSA IR
 - Constructing SSA graphs
 - Sample of SSA-based optimizations
 - Converting back from SSA form
 - Sources: Appel ch. 19, also an extended discussion in Cooper-Torczon sec. 9.3, Mike Ringenbourg's CSE 401 slides (13wi)

Value Numbering

Review: Value Numbering

- Technique for eliminating redundant expressions:
 - Assign an identifying number $VN(n)$ to each expression
 - $VN(x + y) = VN(j)$ if $x+y$ and j have the same value
 - Use hashing over value numbers for efficiency
- Old idea (Balke 1968, Ershov 1954)
 - Invented for low-level, linear IRs
 - Equivalent methods exist for tree IRs, e.g., build a DAG

Optimization Categories (1)

- *Local methods*
 - Usually confined to basic blocks
 - Simplest to analyze and understand
 - Most precise information

Optimization Categories (2)

- *Superlocal methods*
 - Operate over *Extended Basic Blocks* (EBBs)
 - An EBB is a set of blocks b_1, b_2, \dots, b_n where b_1 has multiple predecessors and each of the remaining blocks b_i ($2 \leq i \leq n$) have only b_{i-1} as its unique predecessor
 - The EBB is entered only at b_1 , but may have multiple exits
 - A single block b_i can be the head of multiple EBBs (these EBBs form a tree rooted at b_i)
 - Use information discovered in earlier blocks to improve code in successors

Optimization Categories (3)

- *Regional methods*
 - Operate over scopes larger than an EBB but smaller than an entire procedure/function/method
 - Typical example: loop body
 - Difference from superlocal methods is that there may be merge points in the graph (i.e., a block with two or more predecessors)
 - Facts true at merge point are facts known to be true on all possible paths to that point

Optimization Categories (4)

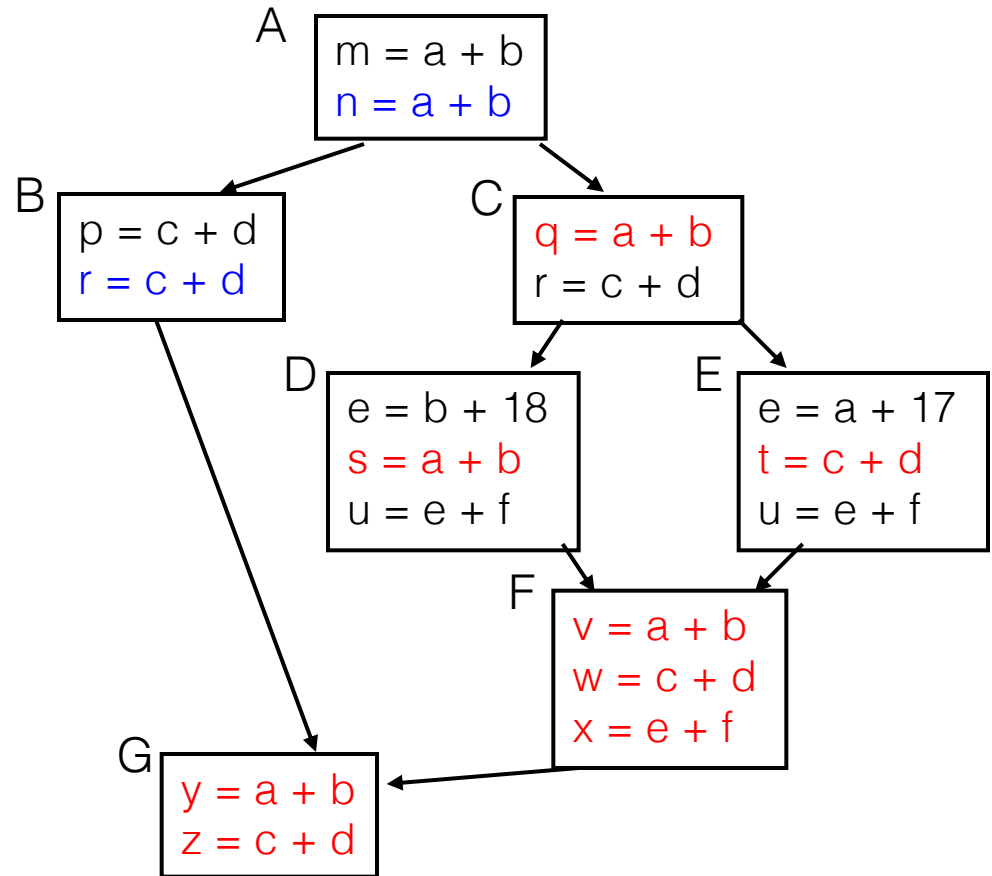
- *Global methods*
 - Operate over entire procedures
 - Sometimes called *intraprocedural* methods
 - Motivation is that local optimizations sometimes have bad consequences in larger context
 - Procedure/method/function is a natural unit for analysis, separate compilation, etc.
 - Almost always need global *data-flow* analysis information for these

Optimization Categories (5)

- *Whole-program methods*
 - Operate over more than one procedure
 - Sometimes called *interprocedural* methods
 - Challenges: name scoping and parameter binding issues at procedure boundaries
 - Classic examples: inline method substitution, interprocedural constant propagation
 - Common in aggressive JIT compilers and optimizing compilers for object-oriented languages


Value Numbering Revisited

- Local Value Numbering
 - 1 block at a time
 - Strong local results
 - No cross-block effects
- Missed opportunities



Superlocal Value Numbering

- Idea: apply local method to EBBs
 - {A,B}, {A,C,D}, {A,C,E}
- Final info from A is initial info for B, C; final info from C is initial for D, E
- Gets reuse from ancestors
- Avoid reanalyzing A, C
- Doesn't help with F, G



SSA Name Space

- Two Principles
 - Each name is defined by exactly one operation
 - Each operand refers to exactly one definition
- Need to deal with merge points
 - Add Φ functions at merge points to reconcile names
 - Use subscripts on variable names for uniqueness

SSA Name Space (from before)

Code

$$a_0^3 = x_0^1 + y_0^2$$

$$b_0^3 = x_0^1 + y_0^2$$

$$a_1^4 = 17$$

$$c_0^3 = x_0^1 + y_0^2$$

Rewritten

$$a_0^3 = x_0^1 + y_0^2$$

$$b_0^3 = a_0^3$$

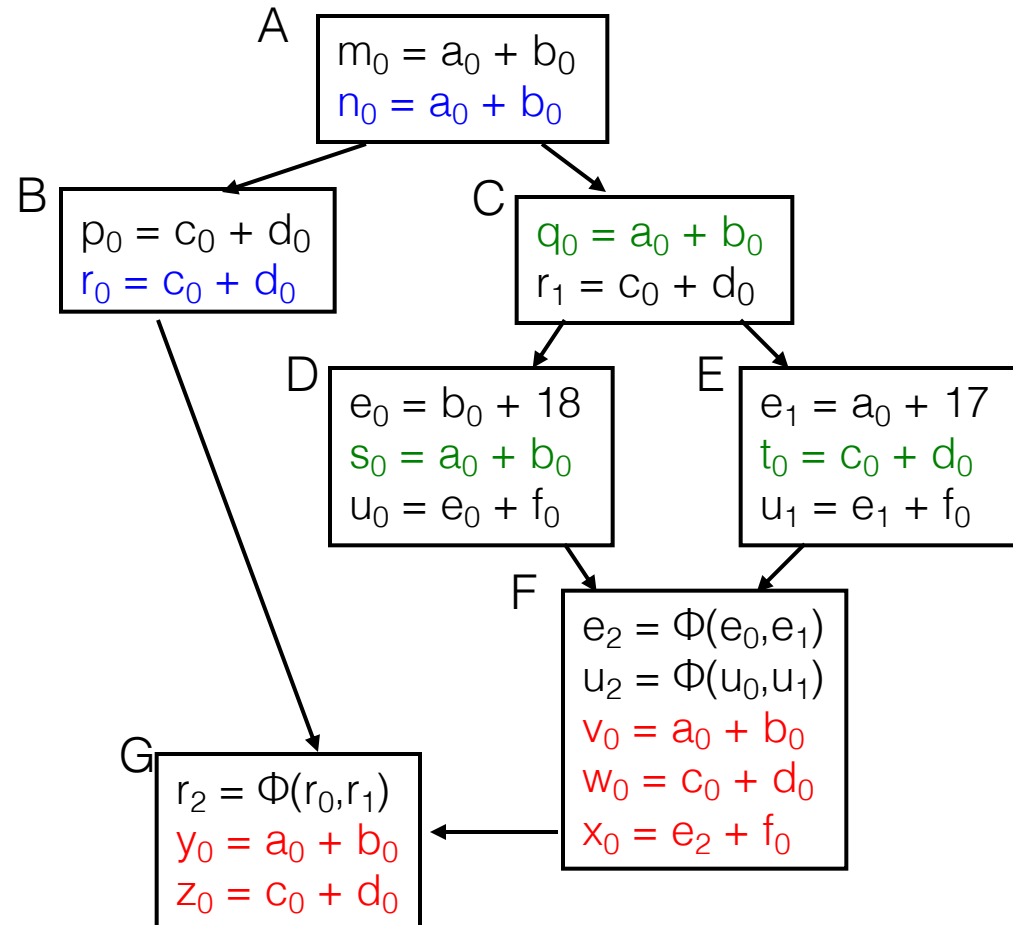
$$a_1^4 = 17$$

$$c_0^3 = a_0^3$$

- Unique name for each definition
- Name \Leftrightarrow VN
- a_0^3 is available to assign to c_0^3


Superlocal Value Numbering with All Bells & Whistles

- Finds more redundancies
- Little extra cost
- Still does nothing for F and G



Larger Scopes

- Still have not helped F and G
- Problem: multiple predecessors
- Must decide what facts hold in F and in G
 - For G, combine B & F?
 - Merging states is expensive
 - Fall back on what we know



Dominators

- Definition
 - x *dominates* y if and only if every path from the entry of the control-flow graph to y includes x
- By definition, x dominates x
- Associate a Dom set with each node
 - $|\text{Dom}(x)| \geq 1$
- Many uses in analysis and transformation
 - Finding loops, building SSA form, code motion

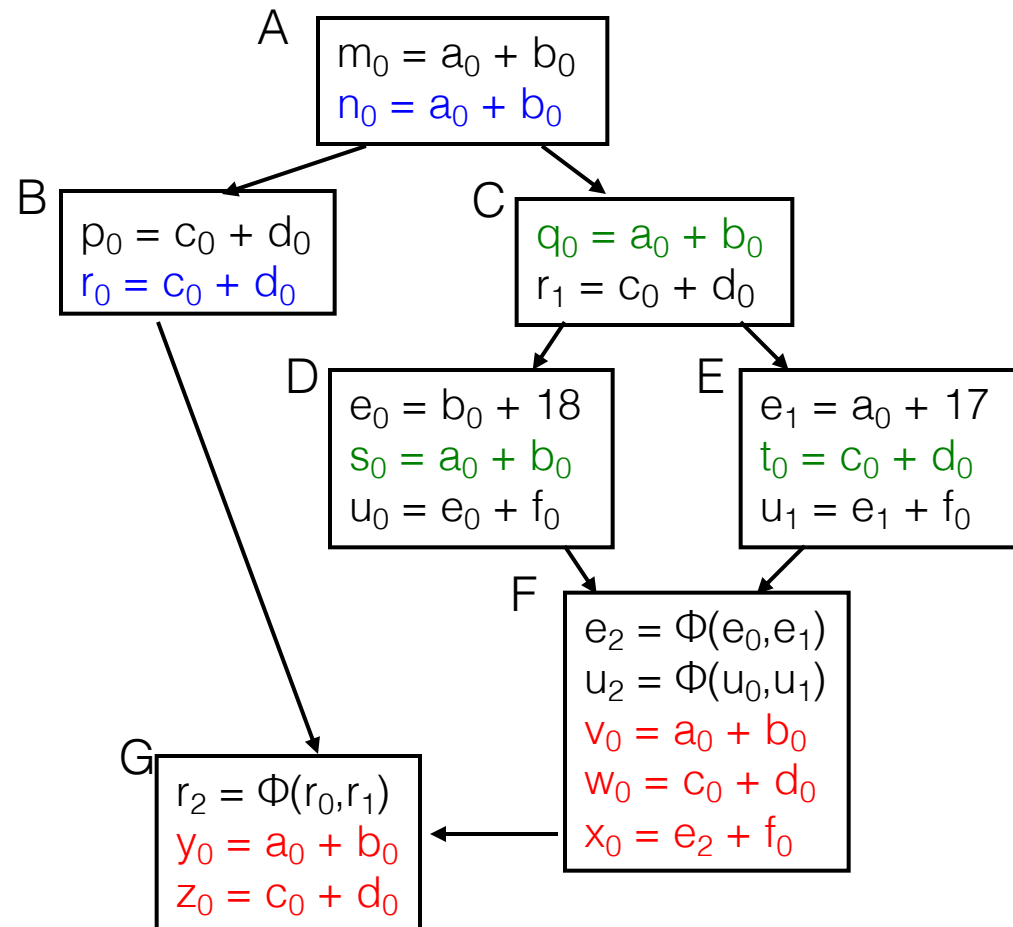
Immediate Dominators

- For any node x , there is a y in $\text{Dom}(x)$ closest to x
- This is the *immediate dominator* of x
 - Notation: $\text{IDom}(x)$

Dominator Sets

Block Dom

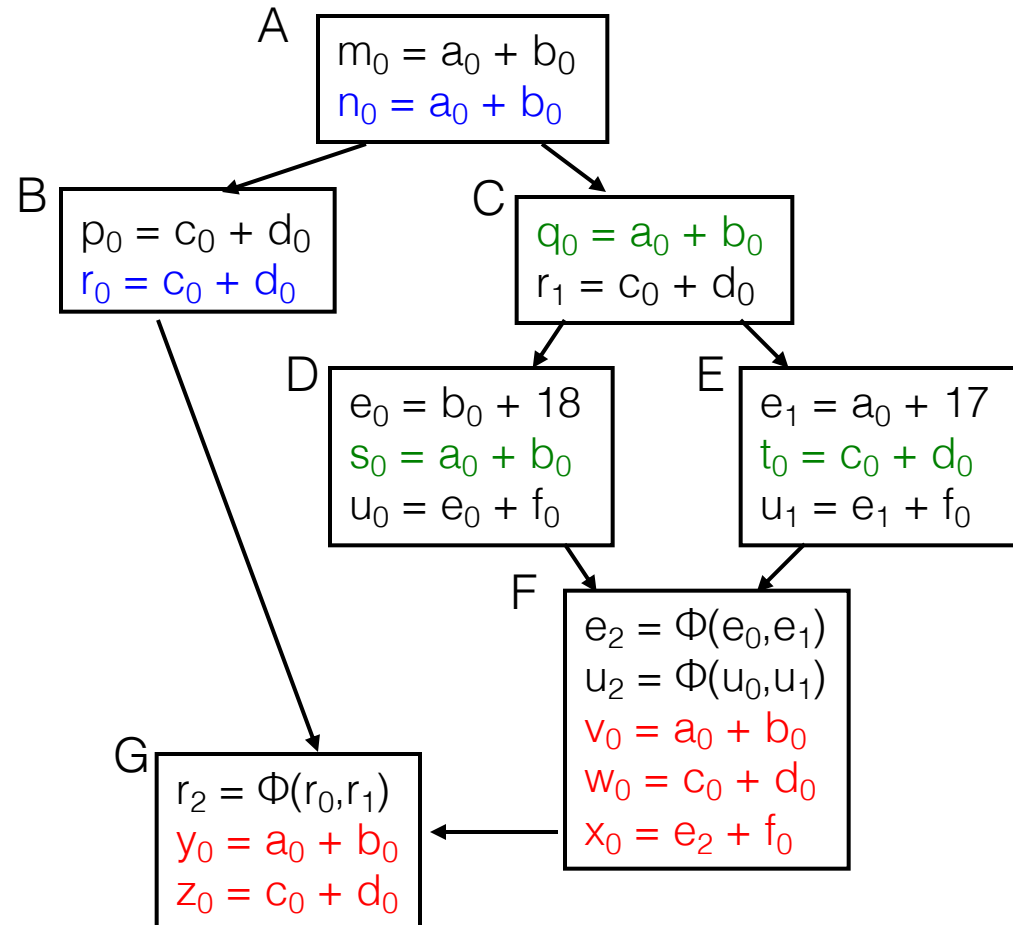
IDom



Note that the IDOM relation defines a tree!

Dominator Value Numbering

- Still looking for a way to handle F and G
- Idea: Use info from $IDom(x)$ to start analysis of x
 - Use C for F and A for G
- Dominator VN
Technique (DVNT)




DVNT Algorithm

- Use superlocal algorithm on extended basic blocks
 - Use scoped hash tables & SSA name space as before
- Start each node with table from its IDOM
- No values flow along back edges (i.e., loops)
- Constant folding, algebraic identities as before

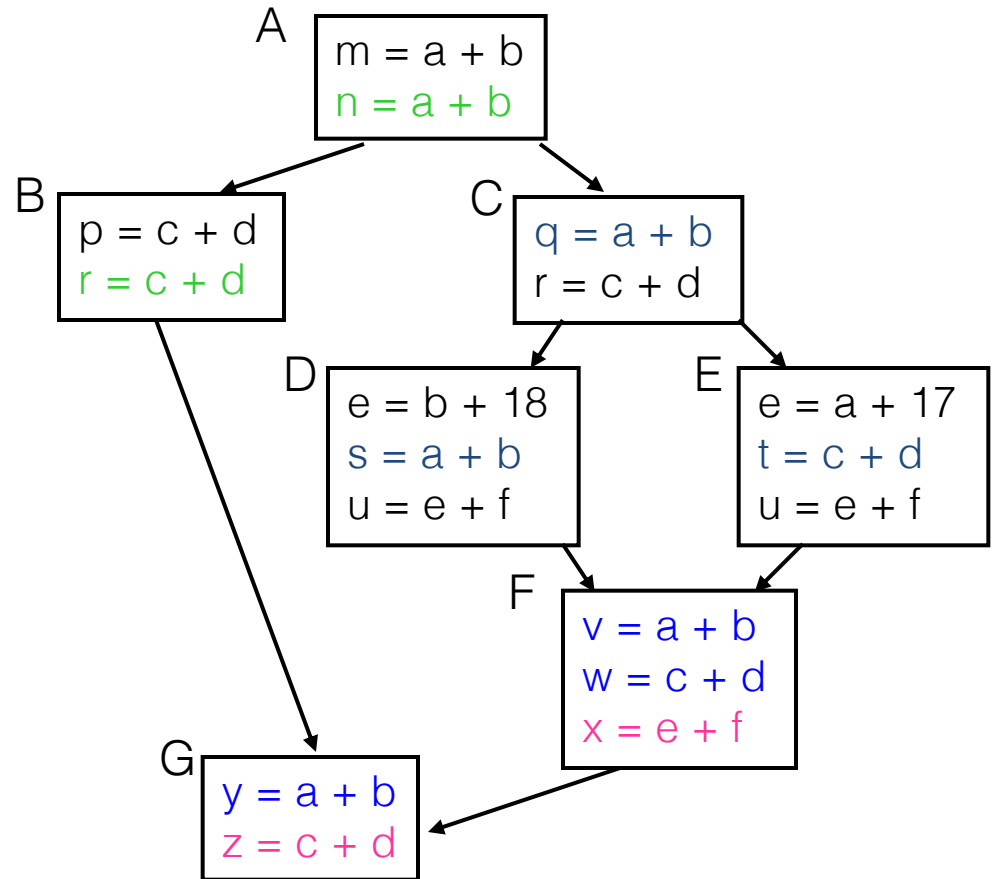
Dominator Value Numbering

- Advantages
 - Finds more redundancy
 - Little extra cost
- Shortcomings
 - Misses some opportunities (common calculations in ancestors that are not IDOMs)
 - Doesn't handle loops or other back edges



Comparing Algorithms

- LVN – Local Value Numbering
- SVN – Superlocal Value Numbering
- DVN – DominatoT-based Value Numbering
- GRE – Global Redundancy Elimination



Comparing Algorithms (2)

- $LVN \Rightarrow SVN \Rightarrow DVN$ form a strict hierarchy – later algorithms find a superset of previous information
- Global RE finds a somewhat different set
 - Discovers $e+f$ in F (computed in both D and E)
 - Misses identical values if they have different names (e.g.,
 $a+b$ and $c+d$ when $a=c$ and $b=d$)
 - Value Numbering catches this

Scope of Analysis

- Larger context (EBBs, regions, global, interprocedural) sometimes helps
 - More opportunities for optimizations
- But not always
 - Introduces uncertainties about flow of control
 - Usually only allows weaker analysis
 - Sometimes has unwanted side effects
 - Can create additional pressure on registers, for example

The Story So Far...

- Local algorithm
- Superlocal extension
 - Some local methods extend cleanly to superlocal scopes
- Dominator VN Technique (DVNT)
- All of these propagate along forward edges
- None are global

Dataflow Analysis

- A collection of techniques for **compile-time** reasoning about **run-time values**
- Almost always involves building a graph
 - Trivial for basic blocks
 - Control-flow graph or derivative for global problems
 - Call graph or derivative for whole-program problems

Dataflow Analysis

- Limitations
 - Precision – “up to symbolic execution”
 - Assumes all paths taken
 - Sometimes cannot afford to compute full solution
 - Arrays – classic analysis treats each array as a single fact
 - Pointers – difficult, expensive to analyze
 - Imprecision rapidly adds up
 - But gotta do it to effectively optimize things like C/C++
- For scalar values we can quickly solve simple problems

Dataflow Analysis

- Many different applications of dataflow analysis:
 - Available expressions
 - Live variables
 - Reaching definitions
 - Very busy expressions

Characterizing Dataflow Analysis

- All of these algorithms involve sets of facts about each basic block b
 - $IN(b)$ – facts true on entry to b
 - $OUT(b)$ – facts true on exit from b
 - $GEN(b)$ – facts created and not killed in b
 - $KILL(b)$ – facts killed in b
- These are related by the equation
$$OUT(b) = GEN(b) \cup (IN(b) - KILL(b))$$
 - Solve this iteratively for all blocks
 - Sometimes information propagates forward; sometimes backward


Available Expressions

Available Expressions

- **Goal:** use dataflow analysis to find common sub-expressions whose range spans basic blocks
- **Idea:** calculate *available expressions* at beginning of each basic block
- Avoid re-evaluation of an available expression – use a copy operation

“Available” and Other Terms

- An expression e is *defined* at point p in the CFG if its value is computed at p
 - Sometimes called *definition site*
- An expression e is *killed* at point p if one of its operands is defined at p
 - Sometimes called *kill site*
- An expression e is *available* at point p if every path leading to p contains a prior definition of e and e is not killed between that definition and p



Available Expression Sets

- To compute available expressions, for each block b , define
 - $AVAIL(b)$ – the set of expressions available on entry to b
 - $NKILL(b)$ – the set of expressions not killed in b
 - i.e., all expressions in the program *except* for those killed in b
 - $DEF(b)$ – the set of expressions defined in b and not subsequently killed in b

Computing Available Expressions

- Big Picture
 - Build control-flow graph
 - Calculate initial local data – $DEF(b)$ and $NKILL(b)$
 - This only needs to be done once for each block b and depends only on the statements in b
 - Iteratively calculate $AVAIL(b)$ by repeatedly evaluating equations until nothing changes
 - Another fixed-point algorithm

Computing Available Expressions

- $AVAIL(b)$ is the set

$$AVAIL(b) = \bigcap_{x \in \text{preds}(b)} (\text{DEF}(x) \cup (AVAIL(x) \cap \text{NKILL}(x)))$$

- $\text{preds}(b)$ is the set of b 's predecessors in the CFG
 - The set of expressions available on entry to b is the set of expressions that were available at the end of *every* predecessor basic block x
 - The expressions available on exit from block b are those defined in b or available on entry to b and not killed in b
- This gives a system of simultaneous equations – a dataflow problem

Computing DEF and NKILL (1)

- For each block b with operations o_1, o_2, \dots, o_k
 $KILLED = \emptyset$ // killed *variables*, not expressions
 $DEF(b) = \emptyset$
 for $i = k$ to 1 // note: working back to front
 assume o_i is “ $x = y + z$ ”
 if ($y \notin KILLED$ and $z \notin KILLED$)
 add “ $y + z$ ” to $DEF(b)$
 add x to $KILLED$
 ...

Computing DEF and NKILL (2)

- After computing DEF and KILLED for a block b , compute set of all expressions in the program not killed in b

$NKILL(b) = \{ \text{all expressions} \}$

for each expression e

for each variable $v \in e$

if $v \in KILLED$ then

$NKILL(b) = NKILL(b) - e$

Computing Available Expressions

Once $DEF(b)$ and $NKILL(b)$ are computed for all blocks b

Worklist = { all blocks b_i }

while (Worklist $\neq \emptyset$)

 remove a block b from Worklist

 recompute $AVAIL(b)$

 if $AVAIL(b)$ changed

 Worklist = Worklist \cup successors(b)

Live Variable Analysis

Live Variable Analysis

- A variable v is *live* at point p if and only if there is *any* path from p to a use of v along which v is not redefined
- Some uses:
 - Register allocation – only live variables need a register
 - Eliminating useless stores – if variable not live at store, then stored variable will never be used
 - Detecting uses of uninitialized variables – if live at declaration (before initialization) then it might be used uninitialized
 - Improve SSA construction – only need Φ -function for variables that are live in a block (later)

Liveness Analysis Sets

- For each block b , define
 - $\text{use}[b]$ = variable used in b before any def
 - $\text{def}[b]$ = variable defined in b & not killed
 - $\text{in}[b]$ = variables live on entry to b
 - $\text{out}[b]$ = variables live on exit from b

Equations for Live Variables

- Given the preceding definitions, we have:
$$\text{in}[b] = \text{use}[b] \cup (\text{out}[b] - \text{def}[b])$$
$$\text{out}[b] = \bigcup_{s \in \text{succ}[b]} \text{in}[s]$$
- Algorithm:
 - Set $\text{in}[b] = \text{out}[b] = \emptyset$
 - Update in, out until no change

Equations for Live Variables v2

- Many problems have more than one formulation. For example, Live Variables...
- Sets:
 - **USED**(b) – variables used in b before being defined in b
 - **NOTDEF**(b) – variables not defined in b
 - **LIVE**(b) – variables live on *exit* from b
- Equation:
$$\text{LIVE}(b) = \bigcup_{s \in \text{succ}(b)} \text{USED}(s) \cup (\text{LIVE}(s) \cap \text{NOTDEF}(s))$$

Reaching Definitions

Reaching Definitions

- A definition d of some variable v *reaches* operation i if and only if i reads the value of v and there is a path from d to i that does not define v
- Uses:
 - Find all of the possible definition points for a variable in an expression

Equations for Reaching Definitions

- Sets:
 - **DEFOUT**(b) – set of definitions in b that reach the end of b (i.e., not subsequently redefined in b)
 - **SURVIVED**(b) – set of all definitions not obscured by a definition in b
 - **REACHES**(b) – set of definitions that reach b
- Equation:

$$\text{REACHES}(b) = \bigcup_{p \in \text{preds}(b)} \text{DEFOUT}(p) \cup (\text{REACHES}(p) \cap \text{SURVIVED}(p))$$

Very Busy Expressions

Very Busy Expressions

- An expression e is considered *very busy* at some point p if e is evaluated and used along every path that leaves p , and evaluating e at p would produce the same result as evaluating it at the original locations
- Uses
 - Code hoisting – move e to p (reduces code size; no effect on execution time)

Equations for Very Busy Expressions

- Sets:
 - $USED(b)$ – expressions used in b before they are killed
 - $KILLED(b)$ – expressions redefined in b before they are used
 - $VERYBUSY(b)$ – expressions very busy on exit from b
- Equation:
$$VERYBUSY(b) = \bigcap_{s \in \text{succ}(b)} USED(s) \cup (VERYBUSY(s) - KILLED(s))$$

Using Dataflow Information

- A few examples of possible transformations that use dataflow information:
 - Common sub-expression elimination
 - Constant propagation
 - Copy propagation
 - Dead code elimination

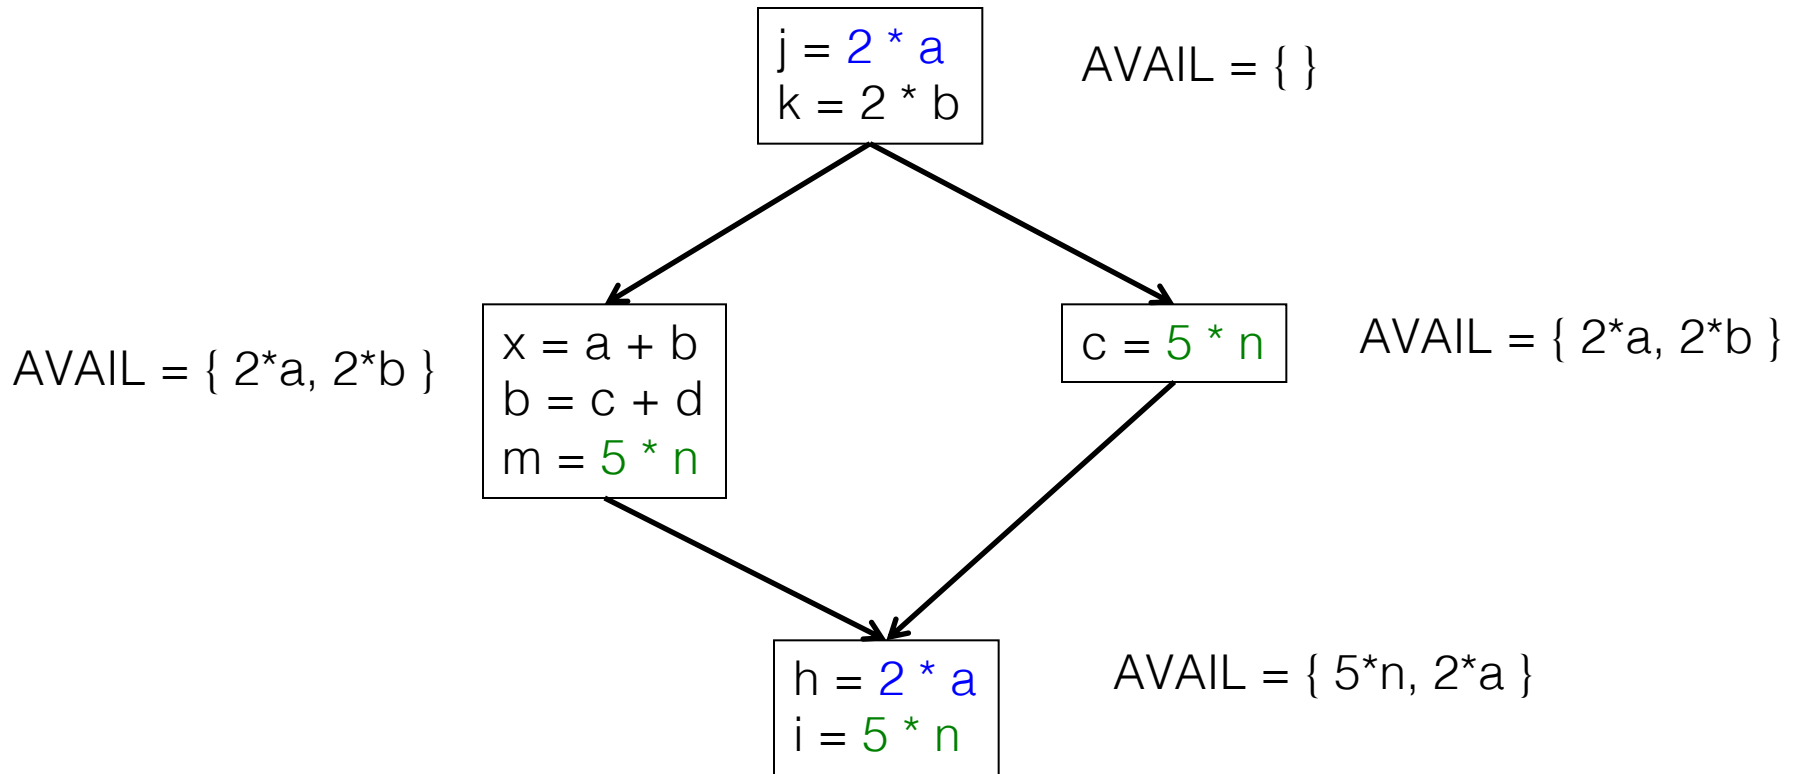
Classic Common-Subexpression Elimination (CSE)

- In a statement $s: t := x \text{ op } y$, if $x \text{ op } y$ is *available* at s , then it need not be recomputed
- **Analysis:** compute *reaching expressions* i.e., statements $n: v := x \text{ op } y$ such that the path from n to s does not compute $x \text{ op } y$ or define x or y

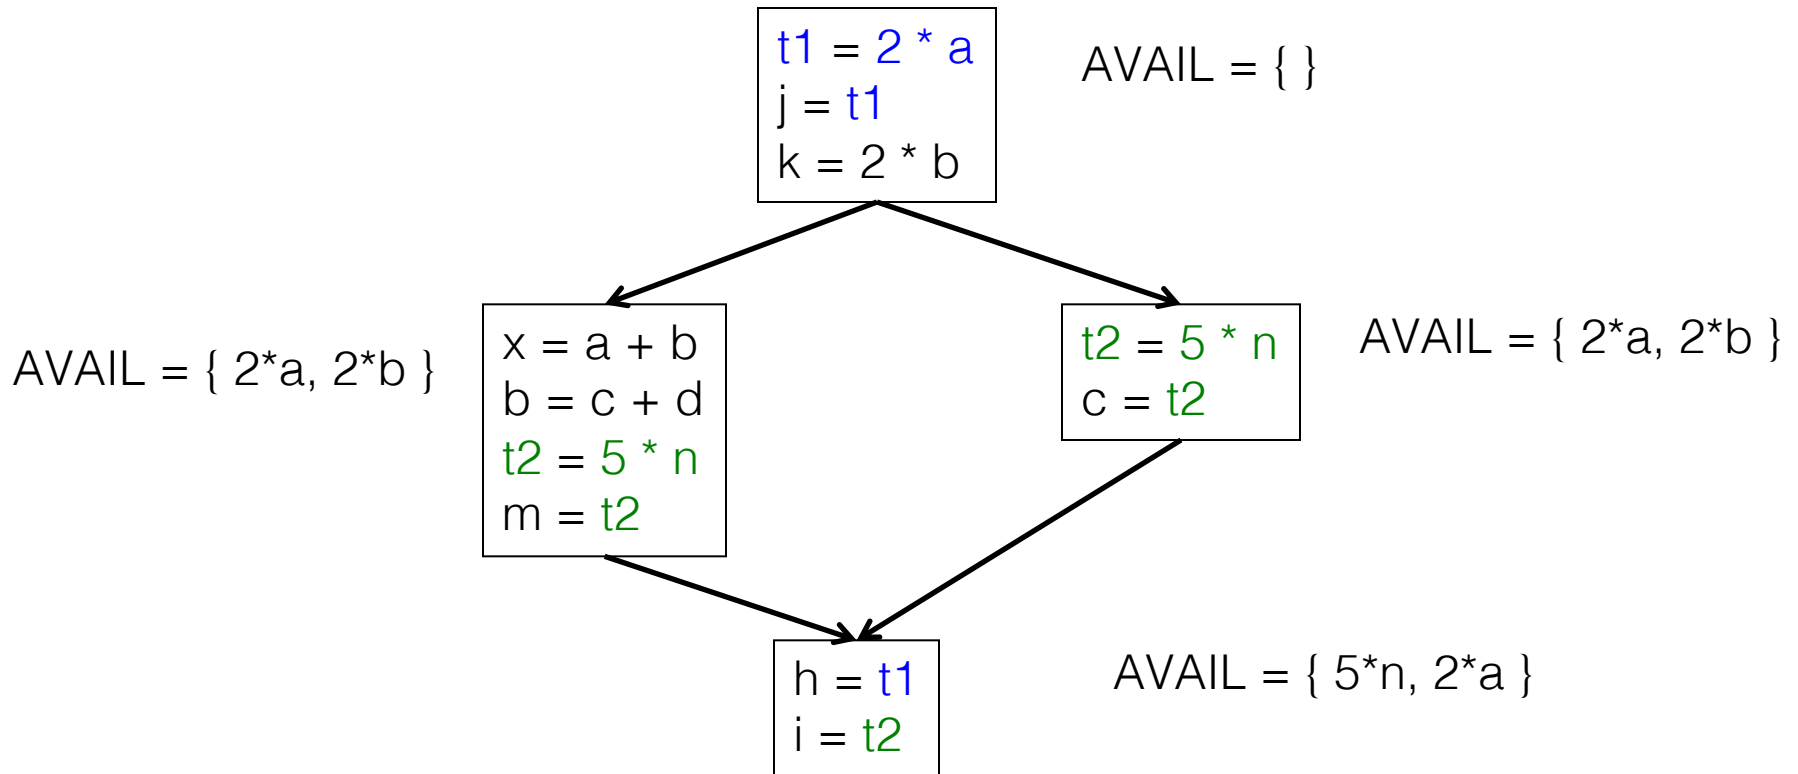
Classic CSE Transformation

- If $x \text{ op } y$ is defined at n and reaches s
 - Create new temporary w
 - Rewrite $n: v := x \text{ op } y$ as
$$\begin{array}{l} n: w := x \text{ op } y \\ n': v := w \end{array}$$
 - Modify statement s to be
$$s: t := w$$
 - (Rely on copy propagation to remove extra assignments that are not really needed)


Revisiting Example (w/slight addition)



Revisiting Example (w/slight addition)



Then Apply Very Busy...



Constant Propagation

- Suppose we have
 - Statement $d: t := c$, where c is constant
 - Statement n that uses t
- If d reaches n and no other definitions of t reach n , then rewrite n to use c instead of t

Copy Propagation

- Similar to constant propagation
- Setup:
 - Statement d: $t := z$
 - Statement n uses t
- If d reaches n and no other definition of t reaches n, and there is no definition of z on any path from d to n, then rewrite n to use z instead of t
 - Recall that this can help remove dead assignments

Copy Propagation Tradeoffs

- Downside is that this can increase the lifetime of variable z and increase need for registers or memory traffic
- But it can expose other optimizations, e.g.,

$a := y + z$

$u := y$

$c := u + z$ // copy propagation makes this $y + z$

- After copy propagation we can recognize the common subexpression

Dead Code Elimination

- If we have an instruction
 $s: a := b \text{ op } c$
and a is not live-out after s , then s can be eliminated
 - Provided it has no implicit side effects that are visible (output, exceptions, etc.)
 - If b or c are function calls, they have to be assumed to have unknown side effects unless the compiler can prove otherwise

Aliases

- A variable or memory location may have multiple names or *aliases*
 - Call-by-reference parameters
 - Variables whose address is taken (&x)
 - Expressions that dereference pointers (p.x, *p)
 - Expressions involving subscripts (a[i])
 - Variables in nested scopes

Aliases vs Optimizations

- Example:

`p.x := 5; q.x := 7; a := p.x;`

- Does reaching definition analysis show that the definition of `p.x` reaches `a`?
- (Or: do `p` and `q` refer to the same variable/object?)
- (Or: *can* `p` and `q` refer to the same thing?)

Aliases vs Optimizations

- Example

```
void f(int *p, int *q) {  
    *p = 1; *q = 2;  
    return *p;  
}
```

- How do we account for the possibility that p and q might refer to the same thing?
- Safe approximation: since it's possible, assume it is true (but rules out a lot)
 - C programmers can use “restrict” to indicate no other pointer is an alias for this one

Types and Aliases (1)

- In Java, ML, MiniJava, and others, if two variables have incompatible types they cannot be names for the same location
 - Also helps that programmer cannot create arbitrary pointers to storage in these languages

Types and Aliases (2)

- **Strategy:** Divide memory locations into *alias classes* based on type information (every type, array, record field is a class)
- **Implication:** need to propagate type information from the semantics pass to optimizer
 - Not normally true of a minimally typed IR
- Items in different alias classes cannot refer to each other

Aliases and Flow Analysis

- **Idea:** Base alias classes on points where a value is created
 - Every new/malloc and each local or global variable whose address is taken is an alias class
 - Pointers can refer to values in multiple alias classes (so each memory reference is to a set of alias classes)
 - Use to calculate “may alias” information (e.g., p “may alias” q at program point s)

Using “may-alias” information

- Treat each alias class as a “variable” in dataflow analysis problems
- Example: framework for available expressions
 - Given statement $s: M[a] := b$,
 $gen[s] = \{ \}$
 $kill[s] = \{ M[x] \mid \text{a may alias } x \text{ at } s \}$

May-Alias Analysis

- Without alias analysis, #2 kills $M[t]$ since x and t might be related
- If analysis determines that “ x may-alias t ” is false, $M[t]$ is still available at #3; can eliminate the common subexpression and use copy propagation
- Code
 - 1: $u := M[t]$
 - 2: $M[x] := r$
 - 3: $w := M[t]$
 - 4: $b := u + w$

Loops

Loops

Much of the execution time of programs is spent here

∴ worth considerable effort to make loops go faster

∴ want to figure out how to recognize loops and figure out how to “improve” them

What Is a Loop?


- In source code, a loop is the set of statements in the body of a **for/while** construct
- But, in a language that permits free use of **GOTOs**, how do we recognize a loop?
- In a control-flow-graph (node = basic-block, arc = flow-of-control), how do we recognize a loop?

Example: Any Loops in this Code?


```
    i = 0
    goto L8
L7: i++
L8: if (i < N) goto L9
    s = 0
    j = 0
    goto L5
L4: j++
L5: N--
    if(j >= N) goto L3
    if (a[j+1] >= a[j]) goto L2
    t = a[j+1]
    a[j+1] = a[j]
    a[j] = t
    s = 1
L2: goto L4
L3: if(s != ) goto L1 else goto L9
L1: goto L7
L9: return
```

Anyone recognize or
guess the algorithm?

Any Loops in this Flowgraph?



Loop in a Flowgraph: Intuition



- Cluster of nodes, such that:
- There's one node called the "header"
- I can reach all nodes in the cluster from the header
- I can get back to the header from all nodes in the cluster
- Only once entrance - via the header
- One or more exits


What Is a Loop?

- In a control flow graph, a loop is a set of nodes S such that:
 - S includes a *header node* h
 - From any node in S there is a path of directed edges leading to h
 - There is a path from h to any node in S
 - There is no edge from any node outside S to any node in S other than h

Entries and Exits

- In a loop
 - An *entry node* is one with some predecessor outside the loop
 - An *exit node* is one that has a successor outside the loop
- **Corollary:** A loop may have multiple exit nodes, but only one entry node

Loop Terminology



Reducible Flow Graphs

- In a reducible flow graph, any two loops are either nested or disjoint
- Roughly, to discover if a flow graph is reducible, repeatedly delete edges and collapse together pairs of nodes (x,y) where x is the only predecessor of y
- If the graph can be reduced to a single node it is reducible
 - Caution: this is the “powerpoint” version of the definition – see a good compiler book for the careful details

Reducible Flow Graphs in Practice

- Common control-flow constructs yield reducible flow graphs
 - if-then[-else], while, do, for, break(!)
- A C function without goto will always be reducible
- Many dataflow analysis algorithms are very efficient on reducible graphs, but...
- We don't need to assume reducible control-flow graphs to handle loops

Finding Loops in Flow Graphs

- We use *dominators* for this
- Recall
 - Every control flow graph has a unique start node s_0
 - Node x dominates node y if every path from s_0 to y must go through x
 - A node x dominates itself

Calculating Dominator Sets

- $D[n]$ is the set of nodes that dominate n
 - $D[s_0] = \{ s_0 \}$
 - $D[n] = \{ n \} \cup (\cap_{p \in \text{pred}[n]} D[p])$
- Set up an iterative analysis as usual to solve this
 - Except initially each $D[n]$ must be all nodes in the graph – updates make these sets smaller if changed

Immediate Dominators

- Every node n has a single *immediate dominator* $\text{idom}(n)$
 - $\text{idom}(n)$ dominates n
 - $\text{idom}(n)$ differs from n – i.e., strictly dominates
 - $\text{idom}(n)$ does not dominate any other strict dominator of n
 - i.e., strictly dominates and is nearest dominator
- **Fact (er, theorem):** If a dominates n and b dominates n , then either a dominates b or b dominates a
 - $\therefore \text{idom}(n)$ is unique

Dominator Tree

- A *dominator tree* is constructed from a flowgraph by drawing an edge from every node n to $\text{idom}(n)$
 - This will be a tree. Why?

Back Edges & Loops

- A flow graph edge from a node n to a node h that dominates n is a *back edge*
- For every back edge there is a corresponding subgraph of the flow graph that is a loop

Natural Loops

- If h dominates n and $n \rightarrow h$ is a back edge, then the *natural loop* of that back edge is the set of nodes x such that
 - h dominates x
 - There is a path from x to n not containing h
- h is the *header* of this loop
- Standard loop optimizations can cope with loops whether they are natural or not

Inner Loops

- Inner loops are more important for optimization because most execution time is expected to be spent there
- If two loops share a header, it is hard to tell which one is “inner”
 - Common way to handle this is to merge natural loops with the same header

Inner (nested) loops

- Suppose
 - A and B are loops with headers a and b
 - $a \neq b$
 - b is in A
- Then
 - The nodes of B are a proper subset of A
 - B is nested in A, or B is the *inner loop*

Loop-Nest Tree

- Given a flow graph G
 1. Compute the dominators of G
 2. Construct the dominator tree
 3. Find the natural loops (thus all loop-header nodes)
 4. For each loop header h , merge all natural loops of h into a single loop: $\text{loop}[h]$
 5. Construct a tree of loop headers s.t. h_1 is above h_2 if h_2 is in $\text{loop}[h_1]$

Loop-Nest Tree Details

- Leaves of this tree are the innermost loops
- Need to put all non-loop nodes somewhere
 - Convention: lump these into the root of the loop-nest tree

Loop Preheader

- Often we need a place to park code right before the beginning of a loop
- Easy if there is a single node preceding the loop header h
 - But this isn't the case in general
- So insert a *preheader* node p
 - Include an edge $p \rightarrow h$
 - Change all edges $x \rightarrow h$ to be $x \rightarrow p$

Loop-Invariant Computations

- **Idea:** If $x := a1 \text{ op } a2$ always does the same thing each time around the loop, we'd like to *hoist* it and do it once outside the loop
- But can't always tell if $a1$ and $a2$ will have the same value
 - Need a conservative (safe) approximation

Loop-Invariant Computations

- d: $x := a_1 \text{ op } a_2$ is *loop-invariant* if for each a_i
 - a_i is a constant, or
 - All the definitions of a_i that reach d are outside the loop, or
 - Only one definition of a_i reaches d, and that definition is loop invariant
- Use this to build an iterative algorithm
 - **Base cases:** constants and operands defined outside the loop
 - **Then:** repeatedly find definitions with loop-invariant operands

Hoisting

- Assume that $d: x := a1 \text{ op } a2$ is loop invariant. We can hoist it to the loop preheader if:
 - d dominates all loop exits where x is live-out, and
 - There is only one definition of x in the loop, and
 - x is not live-out of the loop preheader
- Need to modify this if $a1 \text{ op } a2$ could have side effects or raise an exception

Hoisting: Possible?

- Example 1

L0: $t := 0$

L1: $i := i + 1$

d: $t := a \text{ op } b$

$M[i] := t$

if $i < n$ goto L1

L2: $x := t$

- Example 2

L0: $t := 0$

L1: if $i \geq n$ goto L2

$i := i + 1$

d: $t := a \text{ op } b$

$M[i] := t$

goto L1

L2: $x := t$

Hoisting: Possible?

- Example 3

L0: $t := 0$

L1: $i := i + 1$

d: $t := a \text{ op } b$

$M[i] := t$

$t := 0$

$M[j] := t$

if $i < n$ goto L1

L2: $x := t$

- Example 4

L0: $t := 0$

L1: $M[j] := t$

$i := i + 1$

d: $t := a \text{ op } b$

$M[i] := t$

if $i < n$ goto L1

L2: $x := t$

Induction Variables

- Suppose inside a loop
 - Variable i is incremented or decremented
 - Variable j is set to $i*c+d$ where c and d are loop-invariant
- Then we can calculate j 's value without using i
 - Whenever i is incremented by a , increment j by $c*a$

Example

- Original
 - s := 0
 - i := 0
 - L1: if $i \geq n$ goto L2
 - j := $i * 4$
 - k := j + a
 - x := M[k]
 - s := s + x
 - i := i + 1
 - goto L1
 - L2:
- To optimize, do...
 - Induction-variable analysis to discover i and j are related induction variables
 - Strength reduction to replace $*4$ with an addition
 - Induction-variable elimination to replace $i \geq n$
 - Assorted copy propagation

Result

- Original

```
s := 0
i := 0
L1: if i ≥ n goto L2
    j := i*4
    k := j+a
    x := M[k]
    s := s+x
    i := i+1
    goto L1
L2:
```

- Transformed

```
s := 0
k' = a
b = n*4
c = a+b
L1: if k' ≥ c goto L2
    x := M[k']
    s := s+x
    k' := k'+4
    goto L1
L2:
```

Details are somewhat messy – see your favorite compiler book

Basic and Derived Induction Variables

- Variable i is a *basic induction variable* in loop L with header h if the only definitions of i in L have the form $i := i \pm c$ where c is loop invariant
- Variable k is a *derived induction variable* in L if:
 - There is only one definition of k in L of the form $k := j * c$ or $k := j + d$ where j is an induction variable and c, d are loop-invariant, *and*
 - if j is a derived variable in the family of i , then:
 - The only definition of j that reaches k is the one in the loop, *and*
 - there is no definition of i on any path between the definition of j and the definition of k

Optimizing Induction Variables

- **Strength reduction:** if a derived induction variable is defined with $j := i * c$, try to replace it with an addition inside the loop
- **Elimination:** after strength reduction some induction variables are not used or are only compared to loop-invariant variables; delete them
- **Rewrite comparisons:** If a variable is used only in comparisons against loop-invariant variables and in its own definition, modify the comparison to use a related induction variable

Loop Unrolling

- If the body of a loop is small, much of the time is spent in the “increment and test” code
- **Idea:** reduce overhead by *unrolling* – put two or more copies of the loop body inside the loop

Loop Unrolling

- **Basic idea:** Given loop L with header node h and back edges $s_i \rightarrow h$
 1. Copy the nodes to make loop L' with header h' and back edges $s_i' \rightarrow h'$
 2. Change all back edges in L from $s_i \rightarrow h$ to $s_i \rightarrow h'$
 3. Change all back edges in L' from $s_i' \rightarrow h'$ to $s_i' \rightarrow h$

Unrolling Algorithm Results

- Before

L1: $x := M[i]$

$s := s + x$

$i := i + 4$

if $i < n$ goto L1 else L2

L2:

- After

L1: $x := M[i]$

$s := s + x$

$i := i + 4$

if $i < n$ goto L1' else L2

L1': $x := M[i]$

$s := s + x$

$i := i + 4$

if $i < n$ goto L1 else L2

L2:

Hmmmm....

- Not so great – just code bloat
- But: use induction variables and various loop transformations to clean up

After Some Optimizations

- Before

```
L1: x := M[i]
    s := s + x
    i := i + 4
    if i < n goto L1' else L2
L1':  x := M[i]
    s := s + x
    i := i + 4
    if i < n goto L1 else L2
L2:
```

- After

```
L1: x := M[i]
    s := s + x
    x := M[i+4]
    s := s + x
    i := i + 8
    if i < n goto L1 else L2
L2:
```

Still Broken...

- But in a different, better(?) way
- Good code, but only correct if original number of loop iterations was even
- Fix: add an epilogue to handle the “odd” leftover iteration

Fixed

- Before

```
L1: x := M[i]
    S := S + X
    x := M[i+4]
    S := S + X
    i := i + 8
    if i < n goto L1 else L2
L2:
```

- After

```
    if i < n - 8 goto L1 else L2
L1: x := M[i]
    S := S + X
    x := M[i+4]
    S := S + X
    i := i + 8
    if i < n - 8 goto L1 else L2
L2: x := M[i]
    S := S + X
    i := i + 4
    if i < n goto L2 else L3
L3:
```


Postscript

- This example only unrolls the loop by a factor of 2
- More typically, unroll by a factor of K
 - Then need an epilogue that is a loop like the original that iterates up to $K-1$ times

Memory Hierarchies

- One of the great triumphs of computer design
- Effect is a large, fast memory
- Reality is a series of progressively larger, slower, cheaper stores, with frequently accessed data automatically staged to faster storage (cache, main storage, disk)
- Programmer/compiler typically treats it as one large store. (but not always the best idea)
- Hardware maintains cache coherency – most of the time

Intel Haswell Caches



L1 = 64 KB per core

L2 = 256 KB per core

L3 = 2-8 MB shared

Just How Slow is Operand Access?

- Instruction ~5 per cycle
- Register 1 cycle
- L1 CACHE ~4 cycles
- L2 CACHE ~10 cycles
- L3 CACHE (unshared line) ~40 cycles
- DRAM ~100 ns

Memory Issues

- Byte load/store is often slower than whole (physical) word load/store
 - Unaligned access is often extremely slow
- **Temporal locality**: accesses to recently accessed data will usually find it in the (fast) cache
- **Spatial locality**: accesses to data near recently used data will usually be fast
 - “near” = in the same cache block
- But – alternating accesses to blocks that map to the same cache block will cause thrashing
- CPU speed increases have out-paced increases in memory access times
- Memory access now often determines overall execution speed
- “Instruction count” is not the only performance metric for optimization

Data Alignment

- Data objects (structs) often are similar in size to a cache block (≈ 64 bytes)
 - \therefore Better if objects don't span blocks
- Some strategies:
 - Allocate objects sequentially; bump to next block boundary if useful
 - Allocate objects of same common size in separate pools (all size-2, size-4, etc.)
- Tradeoff: speed for some wasted space

Instruction Alignment

- Align frequently executed basic blocks on cache boundaries (or avoid spanning cache blocks)
- Branch targets (particularly loops) may be faster if they start on a cache line boundary
 - Often see multi-byte nops in optimized code as padding to align loop headers
 - How much depends on architecture (current intel 16 bytes, current AMD 32 bytes)
- Try to move infrequent code (startup, exceptions) away from hot code
- Optimizing compiler may perform basic-block ordering

Loop Interchange

- Watch for bad cache patterns in inner loops; rearrange if possible
- Example

```
for (i = 0; i < m; i++)  
  for (j = 0; j < n; j++)  
    for (k = 0; k < p; k++)  
      a[i,k,j] = b[i,j-1,k] + b[i,j,k] + b[i,j+1,k]
```

 - $b[i,j+1,k]$ is reused in the next two iterations, but will have been flushed from the cache by the k loop

Loop Interchange

- Solution for this example: interchange j and k loops

```
for (i = 0; i < m; i++)  
  for (k = 0; k < p; k++)  
    for (j = 0; j < n; j++)  
      a[i,k,j] = b[i,j-1,k] + b[i,j,k] + b[i,j+1,k]
```

 - Now $b[i,j+1,k]$ will be used three times on each cache load
 - Safe here because loop iterations are independent

Loop Interchange

- Need to construct a data-dependency graph showing information flow between loop iterations
- For example, iteration (j,k) depends on iteration (j',k') if (j',k') computes values used in (j,k) or stores values overwritten by (j,k)
 - If there is a dependency and loops are interchanged, we could get different results – so can't do it

Blocking

- Consider matrix multiply

```
for (i = 0; i < n; i++)
  for (j = 0; j < n; j++) {
    c[i,j] = 0.0;
    for (k = 0; k < n; k++)
      c[i,j] = c[i,j] + a[i,k]*b[k,j]
  }
```
- If a, b fit in the cache together, great!
- If they don't, then every $b[k,j]$ reference will be a cache miss
- Loop interchange ($i \leftrightarrow j$) won't help; then every $a[i,k]$ reference would be a miss

Blocking

- **Solution:** reuse rows of A and columns of B while they are still in the cache
- Assume the cache can hold $2 \times c \times n$ matrix elements ($1 < c < n$)
- Calculate $c \times c$ blocks of C using c rows of A and c columns of B

Blocking

- Calculating $c \times c$ blocks of C
for ($i = i_0; i < i_0 + c; i++$)
 for ($j = j_0; j < j_0 + c; j++$) {
 $c[i,j] = 0.0;$
 for ($k = 0; k < n; k++$)
 $c[i,j] = c[i,j] + a[i,k] * b[k,j]$
 }

Blocking

- Then nest this inside loops that calculate successive $c \times c$ blocks


```
for (i0 = 0; i0 < n; i0+=c)
  for (j0 = 0; j0 < n; j0+=c)
    for (i = i0; i < i0+c; i++)
      for (j = j0; j < j0+c; j++) {
        c[i,j] = 0.0;
        for (k = 0; k < n; k++)
          c[i,j] = c[i,j] + a[i,k]*b[k,j]
      }
```

SSA

Def-Use (DU) Chains

- **Common dataflow analysis problem:** Find all sites where a variable is used, or find the definition site of a variable used in an expression
- **Traditional solution:** def-use chains – additional data structure on top of the dataflow graph
 - Link each statement defining a variable to all statements that use it
 - Link each use of a variable to its definition

Def-Use (DU) Chains



In this example, two DU chains intersect

DU-Chain Drawbacks

- **Expensive:** if a typical variable has N uses and M definitions, the total cost *per-variable* is $O(N * M)$, i.e., $O(n^2)$
 - Would be nice if cost were proportional to the size of the program
- Unrelated uses of the same variable are mixed together
 - Complicates analysis – variable looks live across all uses even if unrelated

SSA: Static Single Assignment

- IR where each variable has only one definition in the program text
 - This is a single *static* definition, but that definition can be in a loop that is executed dynamically many times
- Makes many analyses (and associated optimizations) more efficient
- Separates values from memory storage locations
- Complementary to CFG/DFG – better for some things, but cannot do everything

SSA in Basic Blocks

Idea: for each original variable x , create a new variable x_n at the n^{th} definition of the original x . Subsequent uses of x use x_n until the next definition point.

- Original


- $a := x + y$
- $b := a - 1$
- $a := y + b$
- $b := x * 4$
- $a := a + b$

- SSA

- $a_1 := x + y$
- $b_1 := a_1 - 1$
- $a_2 := y + b_1$
- $b_2 := x * 4$
- $a_3 := a_2 + b_2$


Merge Points

- The issue is how to handle merge points



Merge Points


- The issue is how to handle merge points




- Solution:** introduce a Φ -function
 $a_3 := \Phi(a_1, a_2)$
- Meaning:** a_3 is assigned either a_1 or a_2 depending on which control path is used to reach the Φ -function

Another Example

Original



SSA




How Does Φ “Know” What to Pick?

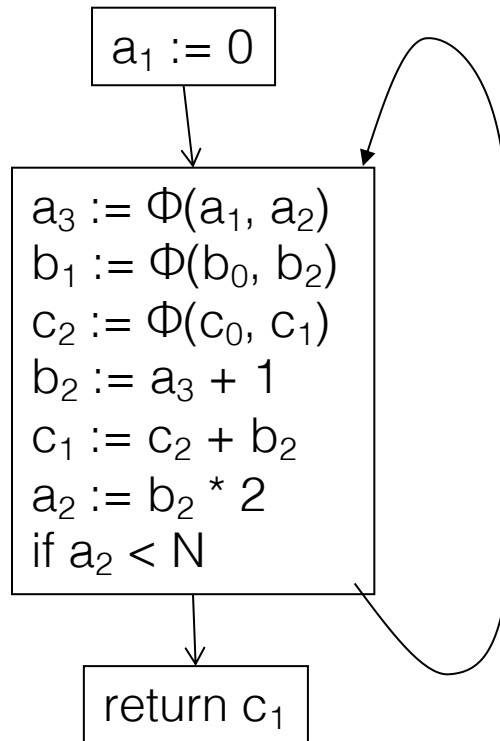
- It doesn't
- Φ -functions don't actually exist at runtime
 - When we're done using the SSA IR, we translate back out of SSA form, removing all Φ -functions
 - Basically by adding code to copy all SSA x_i values to the single, non-SSA, actual x
 - For analysis, all we typically need to know is the connection of uses to definitions – no need to “execute” anything

Example With a Loop

Original



SSA



Notes:

- Loop back edges are also merge points, so require Φ -functions
- a_0, b_0, c_0 are initial values of a, b, c on block entry
- b_1 is dead – can delete later
- c is live on entry – either input parameter or uninitialized

What does SSA “get” us?

- No need for DU or UD chains – implicit in SSA
- Compact representation
- SSA is “recent” (i.e., 80s)
- Prevalent in real compilers for { } languages

Converting To SSA Form

- Basic idea
 - First, add Φ -functions
 - Then, rename all definitions and uses of variables by adding subscripts

Inserting Φ -Functions

- Could simply add Φ -functions for every variable at every join point(!)
- Called “maximal SSA”
- But
 - Wastes *way* too much space and time
 - Not needed in many cases

Path-convergence criterion

- Insert a Φ -function for variable a at point z when:
 - There are blocks x and y , both containing definitions of a , and $x \neq y$
 - There are nonempty paths from x to z and from y to z
 - These paths have no common nodes other than z

Details

- The start node of the flow graph is considered to define every variable (even if “undefined”)
- Each Φ -function itself defines a variable, which may create the need for a new Φ -function
 - So we need to keep adding Φ -functions until things converge
- How can we do this efficiently?
Use a new concept: dominance frontiers

Dominators - Review

- **Definition:** a block x *dominates* a block y if and only if every path from the entry of the control-flow graph to y includes x
- So, by definition, x dominates x

Dominators and SSA

- One property of SSA is that definitions dominate uses; more specifically:
 - If $x := \Phi(\dots, x_i, \dots)$ is in block B , then the definition of x_i dominates the i^{th} predecessor of B
 - If x is used in a non- Φ statement in block B , then the definition of x dominates block B

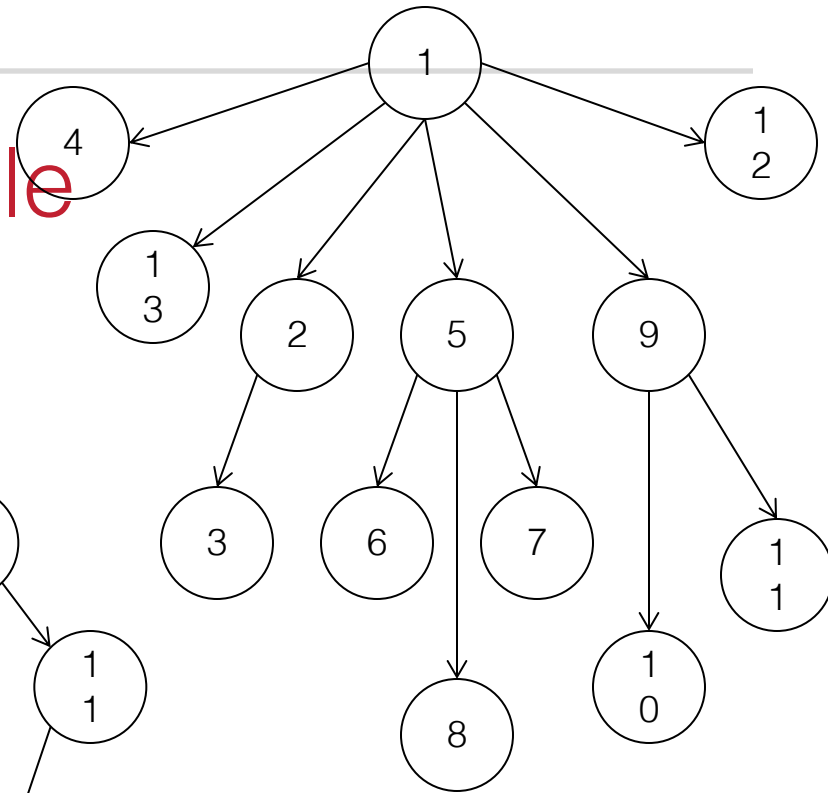

Dominance Frontier (1)

- To get a practical algorithm for placing Φ -functions, we need to avoid looking at all combinations of nodes leading from x to y
- Instead, use the dominator tree in the flow graph

Dominance Frontier (2)

- Definitions
 - x *strictly dominates* y if x dominates y and $x \neq y$
 - The *dominance frontier* of a node x is the set of all nodes w such that x dominates a predecessor of w , but x does not strictly dominate w
 - This means that x can be in *it's own* dominance frontier! That can happen if there is a back edge to x (i.e., x is the head of a loop)
- Essentially, the dominance frontier is the border between dominated and undominated nodes

Example

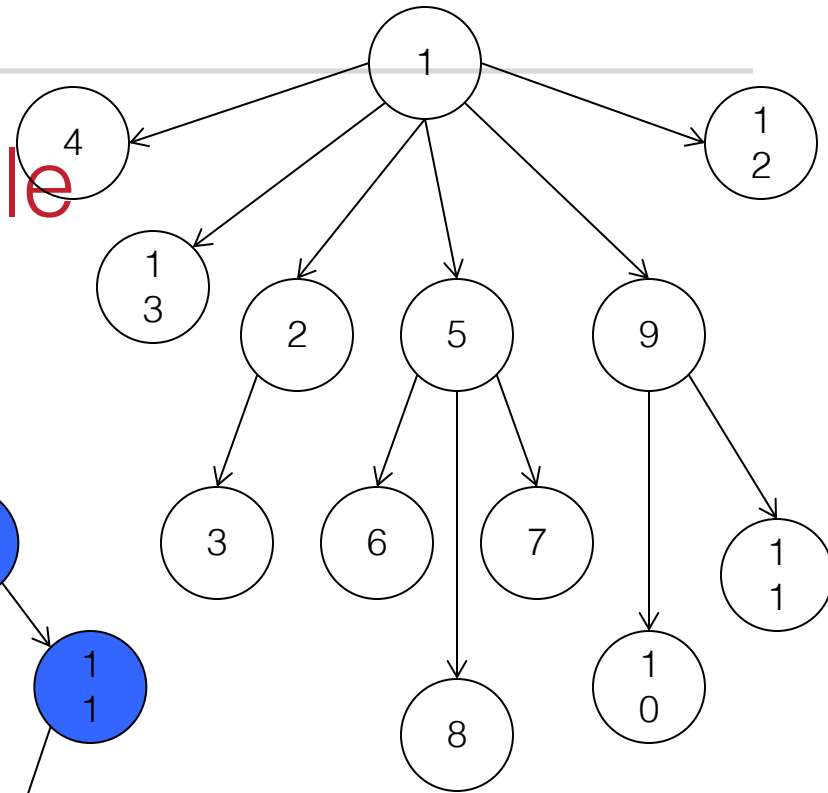



= x

= DomFrontier(x)


= StrictDom(x)

Example

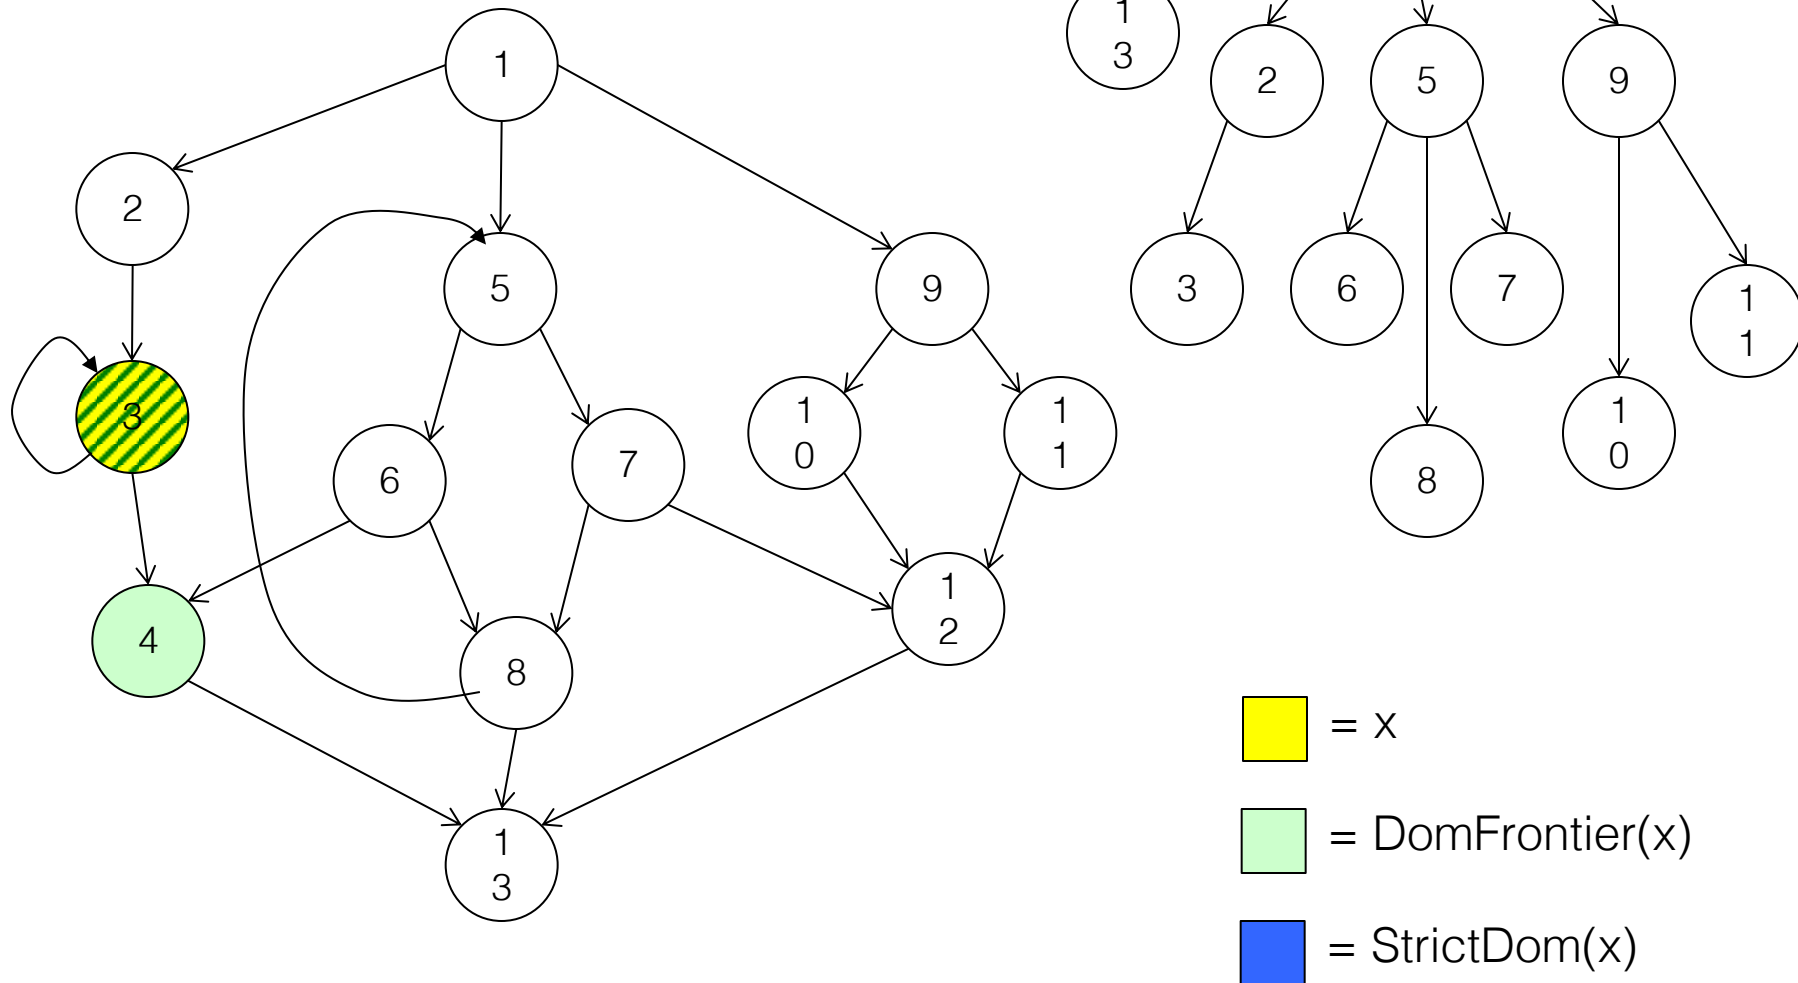


- = x
- = $\text{DomFrontier}(x)$
- = $\text{StrictDom}(x)$



Example



Example





Example



- = x
- = $\text{DomFrontier}(x)$
- = $\text{StrictDom}(x)$

Example





= x

= $\text{DomFrontier}(x)$

= $\text{StrictDom}(x)$

Example





= x

= DomFrontier(x)

= StrictDom(x)

Example




= x

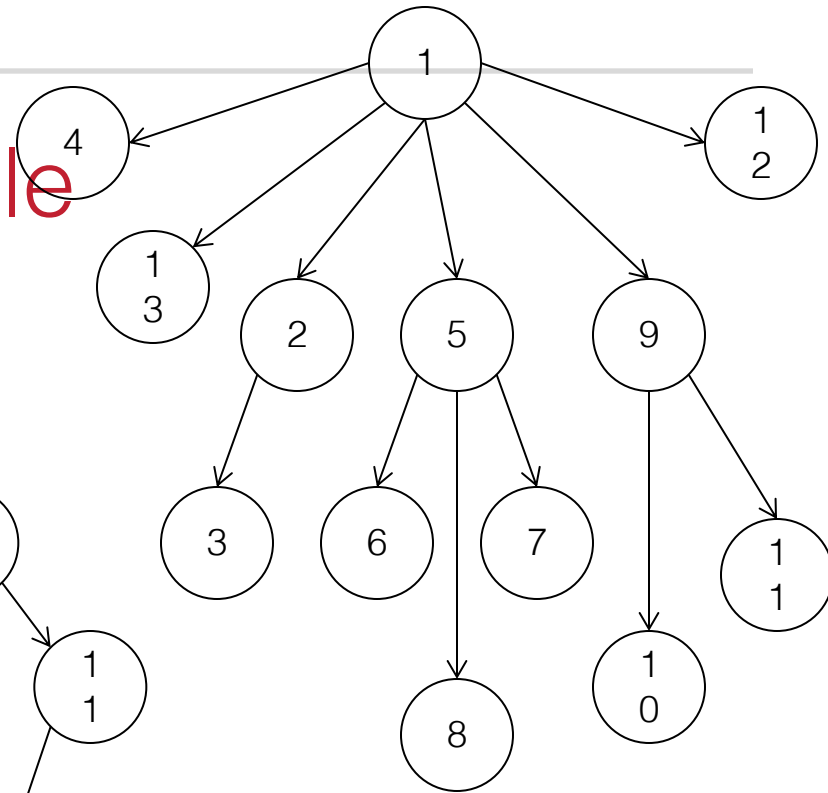

= DomFrontier(x)

= StrictDom(x)

Example



Example

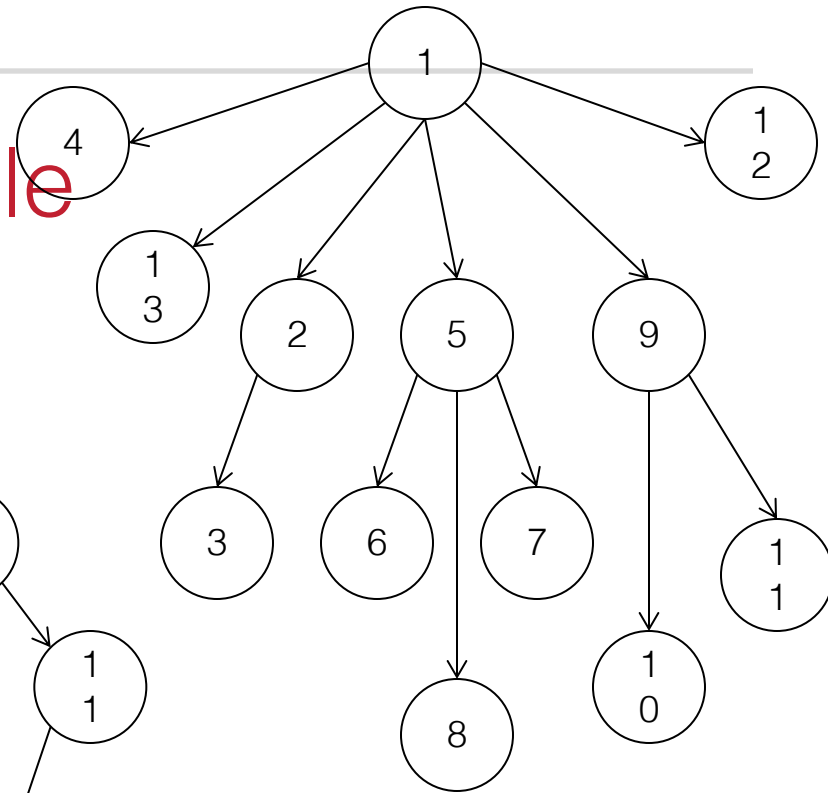



= x

= DomFrontier(x)

= StrictDom(x)

Example



= x

= DomFrontier(x)

= StrictDom(x)

Dominance Frontier Criterion for Placing Φ -Functions

- If a node x contains the definition of variable a , then every node in the dominance frontier of x needs a Φ -function for a
 - **Idea:** Everything dominated by x will see x 's definition of a . The dominance frontier represents the first nodes we could have reached via an alternative path, which *will* have an alternate reaching definition (recall we say the entry node defines everything)
 - Why is this right for loops? Hint: strict dominance...
 - Since the Φ -function itself is a definition, this placement rule needs to be iterated until it reaches a fixed-point
- **Theorem:** this algorithm places exactly the same set of Φ -functions as the path criterion given previously

Placing Φ -Functions: Details

- See the book for the full construction, but the basic steps are:
 1. Compute the dominance frontiers for each node in the flowgraph
 2. Insert just enough Φ -functions to satisfy the criterion. Use a worklist algorithm to avoid reexamining nodes unnecessarily
 3. Walk the dominator tree and rename the different definitions of each variable a to be a_1, a_2, a_3, \dots

Efficient Dominator Tree Computation

- Goal: SSA makes optimizing compilers faster since we can find definitions/uses without expensive bit-vector algorithms
- So, need to be able to compute SSA form quickly
- Computation of SSA from dominator trees are efficient, but...

Lengauer-Tarjan Algorithm

- Iterative set-based algorithm for finding dominator trees is slow in worst case
- Lengauer-Tarjan is near linear time
 - Uses depth-first spanning tree from start node of control flow graph
 - See books for details

SSA Optimizations

- Why go to the trouble of translating to SSA?
- The advantage of SSA is that it makes many optimizations and analyses simpler and more efficient
 - We'll give a couple of examples
- But first, what do we know? (i.e., what information is kept in the SSA graph?)

SSA Data Structures

- **Statement:** links to containing block, next and previous statements, variables defined, variables used.
- **Variable:** link to its (single) definition and (possibly multiple) use sites
- **Block:** List of contained statements, ordered list of predecessors, successor(s)

Dead-Code Elimination

- A variable is live \Leftrightarrow its list of uses is not empty(!)
 - That's it! Nothing further to compute
- Algorithm to delete dead code:
 - while there is some variable v with no uses
 - if the statement that defines v has no other side effects, then delete it
 - Need to remove this statement from the list of uses for its operand variables – which may cause those variables to become dead

Sparse Simple Constant Propagation

- If c is a constant in $v := c$, any use of v can be replaced by c
 - Then update every use of v to use constant c
- If the c_i 's in $v := \Phi(c_1, c_2, \dots, c_n)$ are all the same constant c , we can replace this with $v := c$
- Incorporate copy propagation, constant folding, and others in the same worklist algorithm

Simple Constant Propagation

W := list of all statements in SSA program

while W is not empty

 remove some statement S from W

 if S is $v := \Phi(c, c, \dots, c)$, replace S with $v := c$

 if S is $v := c$

 delete S from the program

 for each statement T that uses v

 substitute c for v in T

 add T to W

Converting Back from SSA

- Unfortunately, real machines do not include a Φ instruction
- So after analysis, optimization, and transformation, need to convert back to a “ Φ -less” form for execution

Translating Φ -functions

- The meaning of $x := \Phi(x_1, x_2, \dots, x_n)$ is “set $x := x_1$ if arriving on edge 1, set $x := x_2$ if arriving on edge 2, etc.”
- So, for each i , insert $x := x_i$ at the end of predecessor block i
- Rely on copy propagation and coalescing in register allocation to eliminate redundant copy instructions

SSA Wrapup

- More details needed to fully and efficiently implement SSA, but these are the main ideas
 - See recent compiler books (but not the new Dragon book!)
- Allows efficient implementation of many optimizations
- SSA is used in most modern optimizing compilers (llvm is based on it) and has been retrofitted into many older ones (gcc is a well-known example)
- Not a silver bullet – some optimizations still need non-SSA forms, but very effective for many

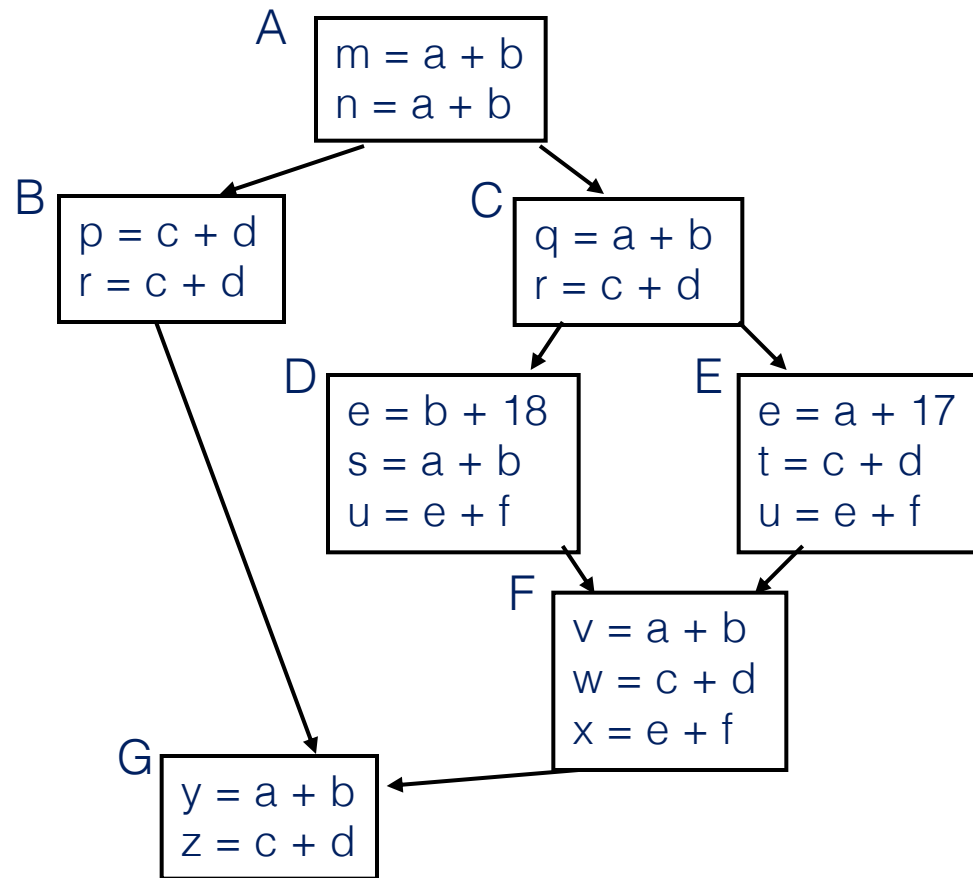
Code Replication

- Sometimes replicating code increases opportunities – modify the code to create larger regions with simple control flow
- Two examples
 - Cloning
 - Inline substitution


Cloning

- **Idea:** duplicate blocks with multiple predecessors
- **Tradeoff**
 - More local optimization possibilities – larger blocks, fewer branches
 - But: larger code size, may slow down if it interacts badly with cache

Original VN Example



Example with Cloning




Inline Substitution

- **Problem:** an optimizer has to treat a procedure call as if it (could have) modified all globally reachable data
 - Plus there is the basic expense of calling the procedure
- **Inline Substitution:** replace each call site with a copy of the called function body

Inline Substitution Issues

- Pro
 - More effective optimization – better local context and don't need to invalidate local assumptions
 - Eliminate overhead of normal function call
- Con
 - Potential code bloat
 - Need to manage recompilation when either caller or callee changes



[Meme credit: imgflip.com]